

**University of Mumbai**

**Practical Journal of**

**Big Data Analytics,  
Modern Networking  
&  
Computer Vision**

**M.Sc. (Information Technology) Part-I**

**Submitted by**

**SINGH MANASI RAKESH**

**Seat No: 1312502**



**DEPARTMENT OF INFORMATION TECHNOLOGY**  
**PILLAI HOC COLLEGE OF ARTS, SCIENCE & COMMERCE, RASAYANI**  
*(Affiliated to Mumbai University)*  
**RASAYANI, 410207**  
**MAHARASHTRA**  
**2023-2024**



**University of Mumbai**

**Practical Journal of  
Big Data Analytics**

**M.Sc. (Information Technology) Part-I**

**Submitted by**

**SINGH MANASI RAKESH**

**Seat No: 1312502**



**DEPARTMENT OF INFORMATION TECHNOLOGY  
PILLAI HOC COLLEGE OF ARTS, SCIENCE & COMMERCE, RASAYANI  
(Affiliated to Mumbai University)  
RASAYANI, 410207  
MAHARASHTRA  
2023-2024**

**Mahatma Education Society's  
Pillai Hoc College of Arts, Science & Commerce, Rasayani  
(Affiliated to Mumbai University)  
RASAYANI – MAHARASHTRA - 410207**

**DEPARTMENT OF  
INFORMATION TECHNOLOGY**



**CERTIFICATE**

This is to certify that the experiment work entered in this journal is as per the syllabus in **M.Sc. (Information Technology) Part-I, Semester-II**; class prescribed by University of Mumbai for the subject **Big Data Analytics** was done in computer lab of Mahatma Education Society's Pillai HOC College of Arts, Science & Commerce, Rasayani by **MANASI R SINGH** during Academic year 2023-2024.

**Exam Seat No: 1312502**

**Subject In-Charge**

**Coordinator**

**External Examiner**

**Principal**

**Date:**

**College Seal**

# **BIG DATA ANALYTICS**

## INDEX

Practical No.	Title	Page No.
1	<b>Install, configure and run Hadoop and HDFS and explore HDFS on Windows.</b>	01
2	<b>Implement word count / frequency programs using MapReduce</b>	04
3	<b>Implement an application that stores big data in Hbase / MongoDB and manipulate it using R / Python</b>	07
4	<b>Implement Decision tree classification techniques</b>	10
5	<b>Implement SVM classification techniques</b>	12
6	<b>REGRESSION MODEL</b> Import a data from web storage. Name the dataset and now do Logistic Regression to find out relation between variables that are affecting the admission of a student in an institute based on his or her GRE score, GPA obtained and rank of the student. Also check the model is fit or not. require (foreign), require(MASS).	15
7	<b>MULTIPLE REGRESSION MODEL</b> Apply multiple regressions, if data have a continuous independent variable. Apply on above dataset.	23
8	<b>CLASSIFICATION MODEL</b> a. Install relevant package for classification. b. Choose classifier for classification problem. c. Evaluate the performance of classifier.	

**9**

**CLUSTERING MODEL**

**a.** Clustering algorithms for unsupervised classification.

**b.** Plot the cluster data using R visualizations.

# Practical No. 01

**Aim: Install, configure and run Hadoop and HDFS and explore HDFS on Windows**

## Steps to Install Hadoop

1. Install Java JDK 1.8
2. Download Hadoop and extract and place under C drive
3. Set Path in Environment Variables
4. Config files under Hadoop directory
5. Create folder datanode and namenode under data directory
6. Edit HDFS and YARN files
7. Set Java Home environment in Hadoop environment
8. Setup Complete. Test by executing start-all.cmd

**There are two ways to install Hadoop, i.e.**

9. Single node
10. Multi node

Here, we use multi node cluster.

### 1. Install Java

11. – Java JDK Link to download  
<https://www.oracle.com/java/technologies/javase-jdk8-downloads.html>
12. – extract and install Java in C:\Java
13. – open cmd and type -> javac -version

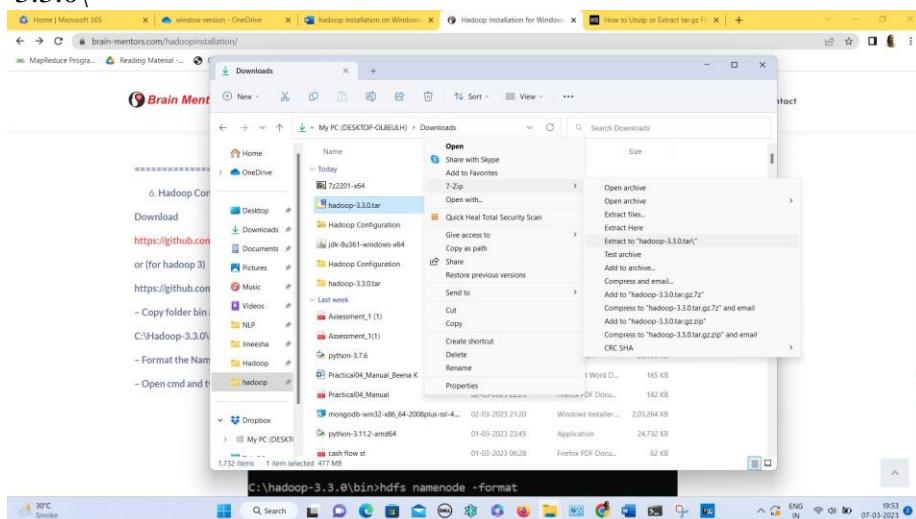
```
C:\Users>cd Beena

C:\Users\Beena>java -version
java version "1.8.0_361"
Java(TM) SE Runtime Environment (build 1.8.0_361-b09)
Java HotSpot(TM) 64-Bit Server VM (build 25.361-b09, mixed mode)
```

### 2. Download Hadoop

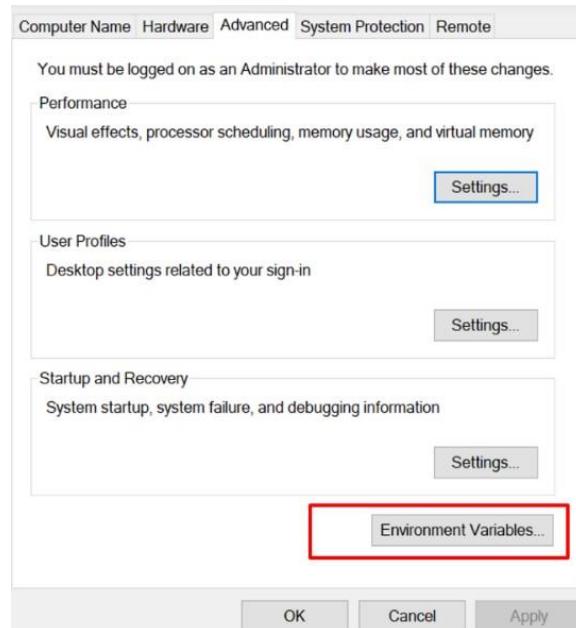
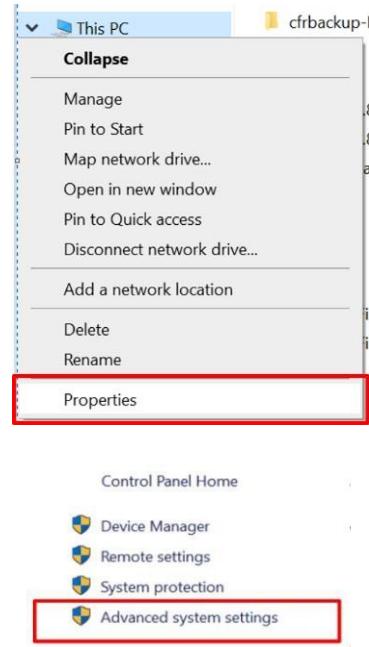
<https://www.apache.org/dyn/closer.cgi/hadoop/common/hadoop-3.3.0/hadoop-3.3.0.tar.gz>

- right click .rar.gz file -> show more options -> 7-zip->and extract to C:\Hadoop-3.3.0\

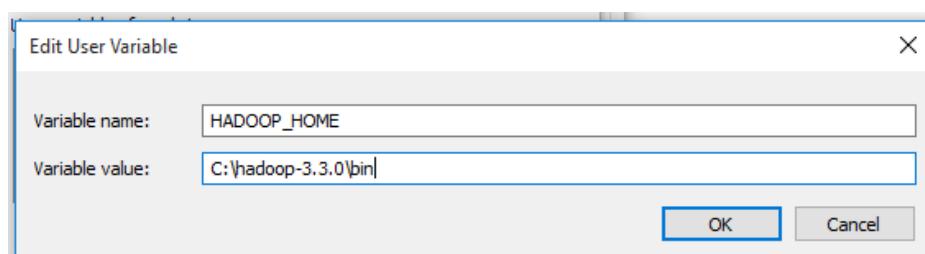


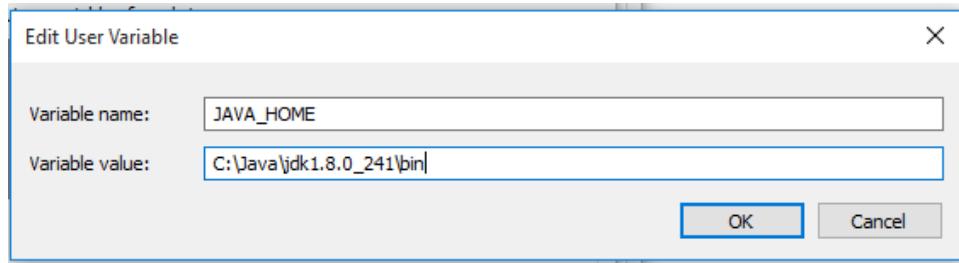
### **3. Set the path JAVA\_HOME Environment variable**

### **4. Set the path HADOOP\_HOME Environment variable**

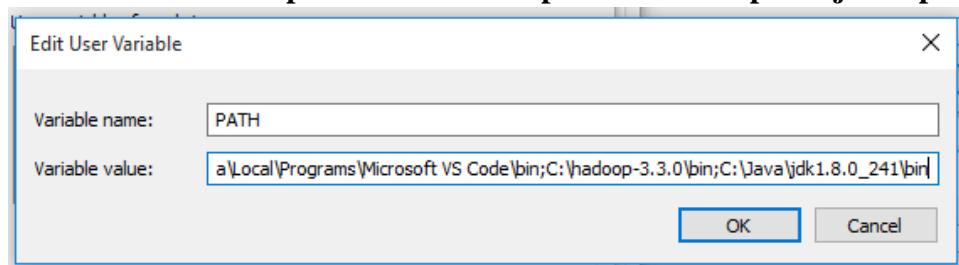


**Click on New to both user variables and system variables.**





**Click on user variable -> path -> edit-> add path for Hadoop and java upto 'bin'**



Click Ok, Ok, Ok.

## 5. Configurations

**Edit file C:/Hadoop-3.3.0/etc/hadoop/core-site.xml,**

paste the xml code in folder and save

```
=====
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://localhost:9000</value>
  </property>
</configuration>
=====
```

**Rename “mapred-site.xml.template” to “mapred-site.xml” and edit this file C:/Hadoop-3.3.0/etc/hadoop/mapred-site.xml, paste xml code and save this file.**

```
=====
<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
</configuration>
```

```
</property>  
</configuration>
```

---

Create folder “data” under “C:\Hadoop-3.3.0”

**Create folder “datanode” under “C:\Hadoop-3.3.0\data”**

**Create folder “namenode” under “C:\Hadoop-3.3.0\data”**

---

**Edit file C:\Hadoop-3.3.0/etc/hadoop/hdfs-site.xml,**

paste xml code and save this file.

```
<configuration>  
  <property>  
    <name>dfs.replication</name>  
    <value>1</value>  
  </property>  
  
  <property>  
    <name>dfs.namenode.name.dir</name>  
    <value>/hadoop-3.3.0/data/namenode</value>  
  </property>  
  
  <property>  
    <name>dfs.datanode.data.dir</name>  
    <value>/hadoop-3.3.0/data/datanode</value>  
  </property>  
</configuration>
```

---

**Edit file C:/Hadoop-3.3.0/etc/hadoop/yarn-site.xml,**

paste xml code and save this file.

```
<configuration>
```

---

```
<property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
</property>

<property>
    <name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</name>
    <value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
<property>
    <name>yarn.resourcemanager.address</name>
    <value>127.0.0.1:8032</value>
</property>
<property>
    <name>yarn.resourcemanager.scheduler.address</name>
    <value>127.0.0.1:8030</value>
</property>
<property>
    <name>yarn.resourcemanager.resource-tracker.address</name>
    <value>127.0.0.1:8031</value>
</property>
</configuration>
=====
```

## 6. Edit file C:/Hadoop-3.3.0/etc/hadoop/hadoop-env.cmd

Find “JAVA\_HOME=%JAVA\_HOME%” and replace it as

```
set JAVA_HOME="C:\Java\jdk1.8.0_361"
```

=====

## 7. Download “redistributable” package

### Download and run VC\_redist.x64.exe

This is a “redistributable” package of the Visual C runtime code for 64-bit applications, from Microsoft. It contains certain shared code that every application written with Visual C expects to have available on the Windows computer it runs on.

## 8. Hadoop Configurations

Download **bin** folder from

<https://github.com/s911415/apache-hadoop-3.1.0-winutils>

- Copy the **bin** folder to c:\hadoop-3.3.0. Replace the existing **bin** folder.
- 9. copy "hadoop-yarn-server-timelineservice-3.0.3.jar" from ~\hadoop-3.0.3\share\hadoop\yarn\timelineservice to ~\hadoop-3.0.3\share\hadoop\yarn folder.

## 10. Format the NameNode

- Open cmd ‘Run as Administrator’ and type command “**hdfs namenode –format**”

```
Administrator: Command Prompt
Microsoft Windows [Version 10.0.22621.1265]
(c) Microsoft Corporation. All rights reserved.

C:\Windows\System32>cd\hadoop-3.3.0\bin

C:\hadoop-3.3.0\bin>hdfs namenode -format
```

```
2023-03-07 21:31:34,685 INFO namenode.FSImageFormatProtobuf: Saving image file C:\hadoop-3.3.0\data\namenode\current\fsimage.ckpt_00000000000000000000 using no compression
2023-03-07 21:31:34,844 INFO namenode.FSImageFormatProtobuf: Image file C:\hadoop-3.3.0\data\namenode\current\fsimage.ckpt_00000000000000000000 of size 400 bytes saved in 0 seconds .
2023-03-07 21:31:34,860 INFO namenode.NNStorageRetentionManager: Going to retain 1 images with txid >= 0
2023-03-07 21:31:34,869 INFO namenode.FSImage: FSImageSaver clean checkpoint: txid=0 when meet shutdown.
2023-03-07 21:31:34,870 INFO namenode.NameNode: SHUTDOWN_MSG:
*****SHUTDOWN_MSG: Shutting down NameNode at DESKTOP-OL8EULH/192.168.1.19
*****
```

## 11. Testing

- Open cmd ‘Run as Administrator’ and change directory to **C:\Hadoop-3.3.0\sbin**

- type **start-all.cmd**

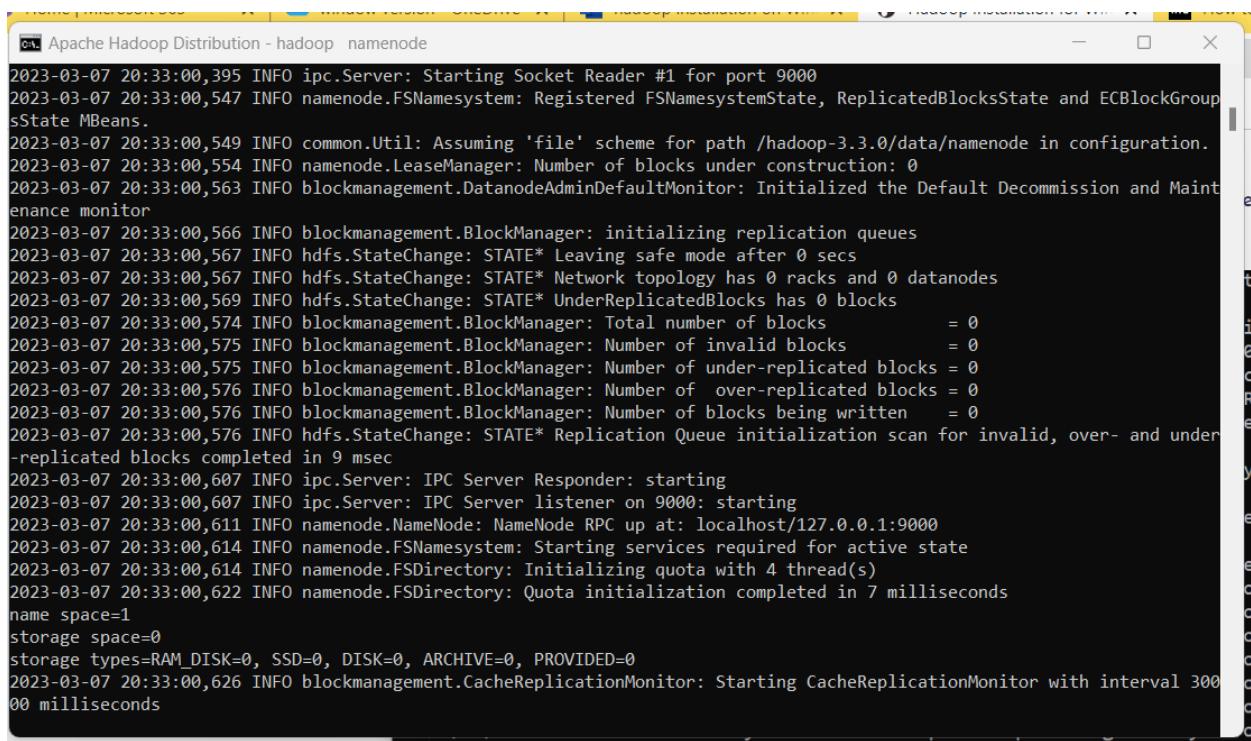
OR

- type **start-dfs.cmd**

- type **start-yarn.cmd**

```
C:\hadoop-3.3.0\sbin>start-all.cmd
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
The filename, directory name, or volume label syntax is incorrect.
The filename, directory name, or volume label syntax is incorrect.
Starting yarn daemons
The filename, directory name, or volume label syntax is incorrect.
```

- You will get 4 more running threads for Datanode, namenode, resource manager and node manager



A screenshot of a Windows Command Prompt window titled "Apache Hadoop Distribution - hadoop namenode". The window displays a log of startup messages from the Hadoop NameNode. The log includes information about socket readers, FSNamesystem registration, utility assumptions, lease manager, block management monitors, and various replication and quota checks. The output is in black text on a dark background, with some lines wrapped due to the window's width.

```
2023-03-07 20:33:00,395 INFO ipc.Server: Starting Socket Reader #1 for port 9000
2023-03-07 20:33:00,547 INFO namenode.FSNamesystem: Registered FSNamesystemState, ReplicatedBlocksState and ECBLOCKGroup
sState MBeans.
2023-03-07 20:33:00,549 INFO common.Util: Assuming 'file' scheme for path /hadoop-3.3.0/data/namenode in configuration.
2023-03-07 20:33:00,554 INFO namenode.LeaseManager: Number of blocks under construction: 0
2023-03-07 20:33:00,563 INFO blockmanagement.DatanodeAdminDefaultMonitor: Initialized the Default Decommission and Maint
enance monitor
2023-03-07 20:33:00,566 INFO blockmanagement.BlockManager: initializing replication queues
2023-03-07 20:33:00,567 INFO hdfs.StateChange: STATE* Leaving safe mode after 0 secs
2023-03-07 20:33:00,567 INFO hdfs.StateChange: STATE* Network topology has 0 racks and 0 datanodes
2023-03-07 20:33:00,569 INFO hdfs.StateChange: STATE* UnderReplicatedBlocks has 0 blocks
2023-03-07 20:33:00,574 INFO blockmanagement.BlockManager: Total number of blocks = 0
2023-03-07 20:33:00,575 INFO blockmanagement.BlockManager: Number of invalid blocks = 0
2023-03-07 20:33:00,575 INFO blockmanagement.BlockManager: Number of under-replicated blocks = 0
2023-03-07 20:33:00,576 INFO blockmanagement.BlockManager: Number of over-replicated blocks = 0
2023-03-07 20:33:00,576 INFO blockmanagement.BlockManager: Number of blocks being written = 0
2023-03-07 20:33:00,576 INFO hdfs.StateChange: STATE* Replication Queue initialization scan for invalid, over- and under
-replicated blocks completed in 9 msec
2023-03-07 20:33:00,607 INFO ipc.Server: IPC Server Responder: starting
2023-03-07 20:33:00,607 INFO ipc.Server: IPC Server listener on 9000: starting
2023-03-07 20:33:00,611 INFO namenode.NameNode: NameNode RPC up at: localhost/127.0.0.1:9000
2023-03-07 20:33:00,614 INFO namenode.FSNamesystem: Starting services required for active state
2023-03-07 20:33:00,614 INFO namenode.FSDirectory: Initializing quota with 4 thread(s)
2023-03-07 20:33:00,622 INFO namenode.FSDirectory: Quota initialization completed in 7 milliseconds
name space=1
storage space=0
storage types=RAM_DISK=0, SSD=0, DISK=0, ARCHIVE=0, PROVIDED=0
2023-03-07 20:33:00,626 INFO blockmanagement.CacheReplicationMonitor: Starting CacheReplicationMonitor with interval 300
00 milliseconds
```

## Output:

12. Type **JPS** command to start-all.cmd command prompt, you will get following output.

```
C:\hadoop-3.3.0\sbin>jps
5632 Jps
7572 DataNode
3752 ResourceManager
7992 NameNode
8028 NodeManager
```

**13.** Run <http://localhost:9870/> from any browser

The screenshot shows the HDFS Health Overview page at <http://localhost:9870/dfshealth.html#tab-overview>. The top navigation bar includes links for Hadoop, Overview, Datanodes, Datanode Volume Failures, Snapshot, Startup Progress, and Utilities. The main content area displays cluster statistics:

Started:	Wed Mar 15 12:10:54 +0530 2023
Version:	3.3.0, raa96f1871bfd858f9bac59cf2a81ec470da649af
Compiled:	Tue Jul 07 00:14:00 +0530 2020 by brahma from branch-3.3.0
Cluster ID:	CID-1986aba8-0ed3-43a2-9db7-42944ec518b2
Block Pool ID:	BP-1049743432-192.168.56.1-1678862097216

Below this, a "Browse Directory" section is shown with a table header and a message indicating "No data available in table".

## Practical No. 02

**Aim : Implement word count / frequency programs using Map Reduce .**

**Code::**

Steps:

C:\hadoop-3.3.0\sbin>start-dfs.cmd

C:\hadoop-3.3.0\sbin>start-yarn.cmd

Open a command prompt as administrator and run the following command to create an input and output folder on the Hadoop file system, to which we will be moving the sample.txt file for our analysis.

C:\hadoop-3.3.0\bin>cd\

C:>hadoop dfsadmin -safemode leave

DEPRECATED: Use of this script to execute hdfs command is deprecated.

Instead use the hdfs command for it.

Safe mode is OFF

C:>hadoop fs -mkdir /input\_dir

Check it by giving the following URL at browser

<http://localhost:9870>

Utilities -> browse the file system

Copy the input text file named input\_file.txt in the input directory (input\_dir)of HDFS.

Make a file in c:\input\_file.txt and write following content in it.

Hadoop Window version is easy compared to Ubuntu version

Now apply the following command at c:>

C:> hadoop fs -put C:/input\_file.txt /input\_dir

Verify input\_file.txt available in HDFS input directory (input\_dir).

C:>Hadoop fs -ls /input\_dir/

Verify content of the copied file

C:>hadoop dfs -cat /input\_dir/input\_file.txt

You can see the file content displayed on the CMD.

Run MapReduceClient.jar and also provide input and out directories.

C:>hadoop jar C:/hadoop-3.3.0/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.3.0.jar wordcount /input\_dir /output\_dir

In case, there is some error in executing then copy the file MapReduceClient.jar in C:\ and run the program with the jar file using existing MapReduceClient.jar file as:

C:\> hadoop jar C:/MapReduceClient.jar wordcount /input\_dir /output\_dir

Now, check the output\_dir on browser as follows:

Click on output\_dir à part-r-00000 à Head the file (first 32 K) and check the file content as the output.

Alternatively, you may type the following command on CMD window as:

C:\> hadoop dfs -cat /output\_dir/\*

You can get the following output

**Output:**

```
hadoop dfs -cat /output_dir/*
Hadoop 1
Window 1
version 1
is 1
easy 1
compared 1
to 1
Ubuntu 1
```

### Practical No. 03

**Aim:** Implement an application that stores big data in Hbase / MongoDB and manipulate it using R / Python

#### Requirements

- a. PyMongo
- b. Mongo Database

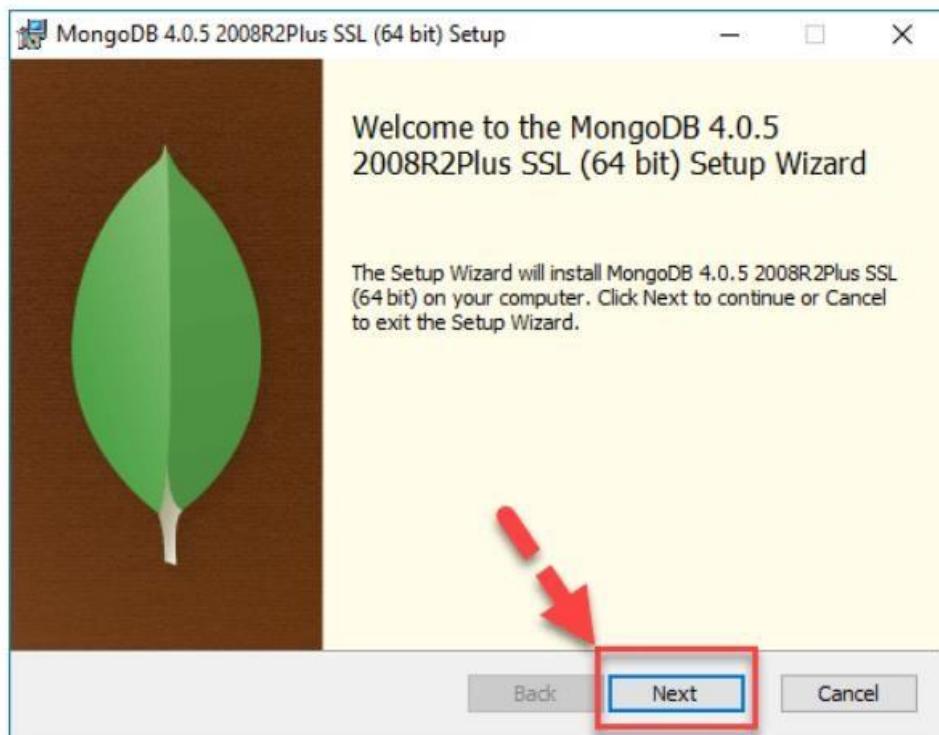
#### Step A: Install Mongo database

Step 1) Go to (<https://www.mongodb.com/download-center/community>) and Download MongoDB Community Server. We will install the 64-bit version for Windows.

Select the server you would like to run:

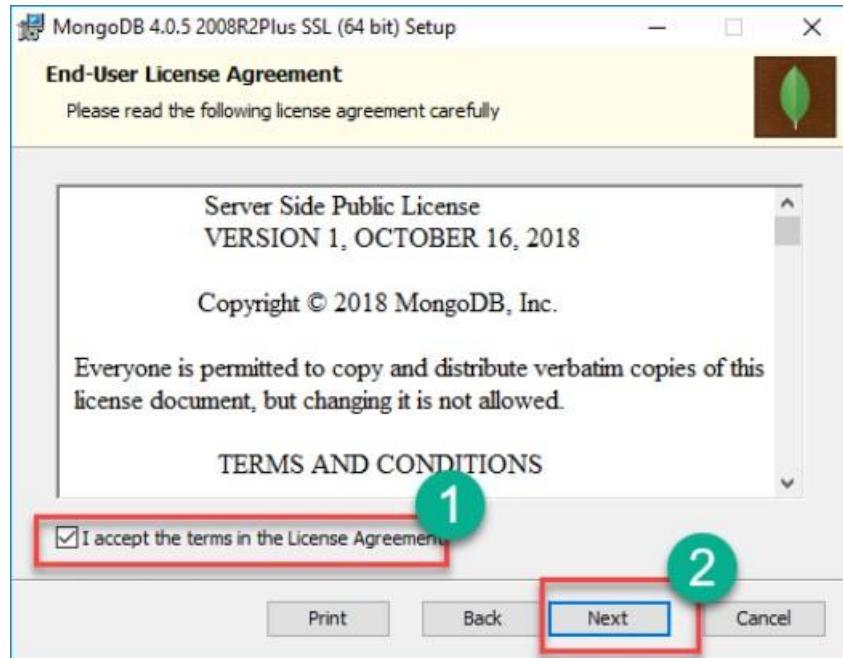


Step 2) Once download is complete open the msi file. Click Next in the start up screen

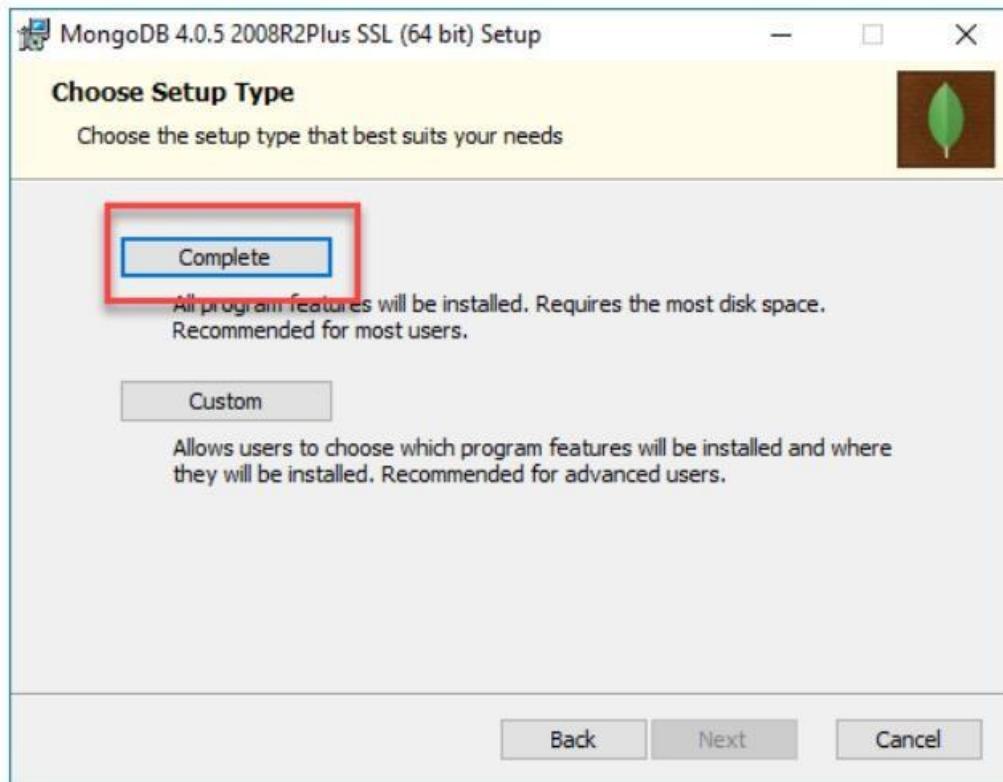


## Step 3)

1. Accept the End-User License Agreement
2. Click Next



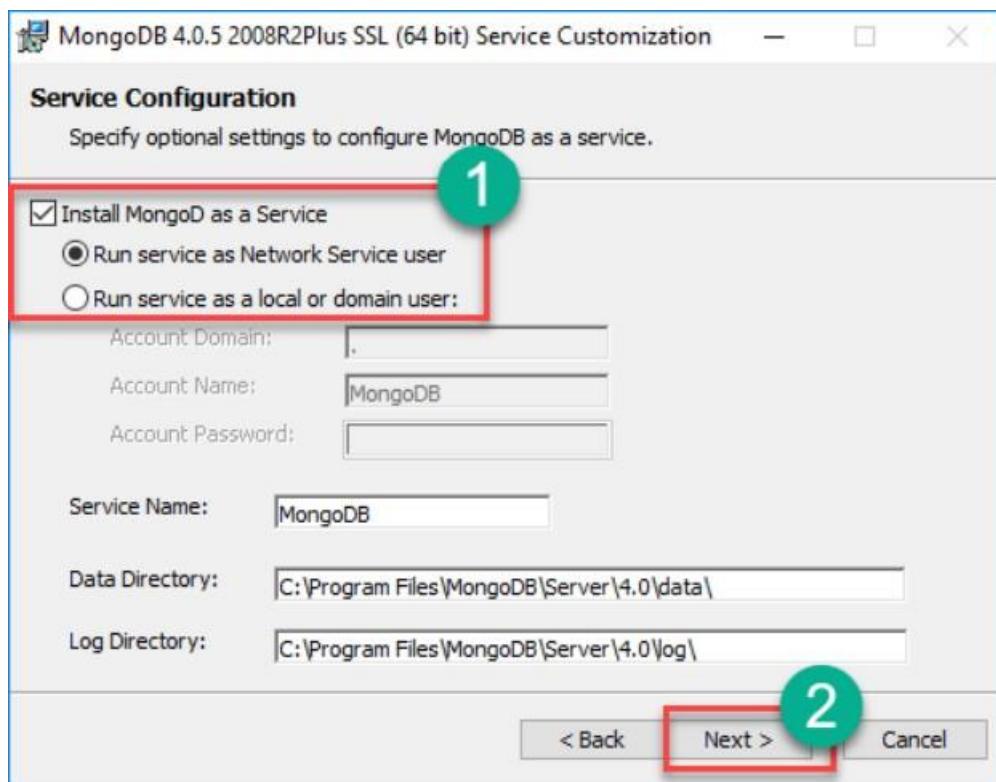
Step 4) Click on the "complete" button to install all of the components. The custom option can be used to install selective components or if you want to change the location of the installation.



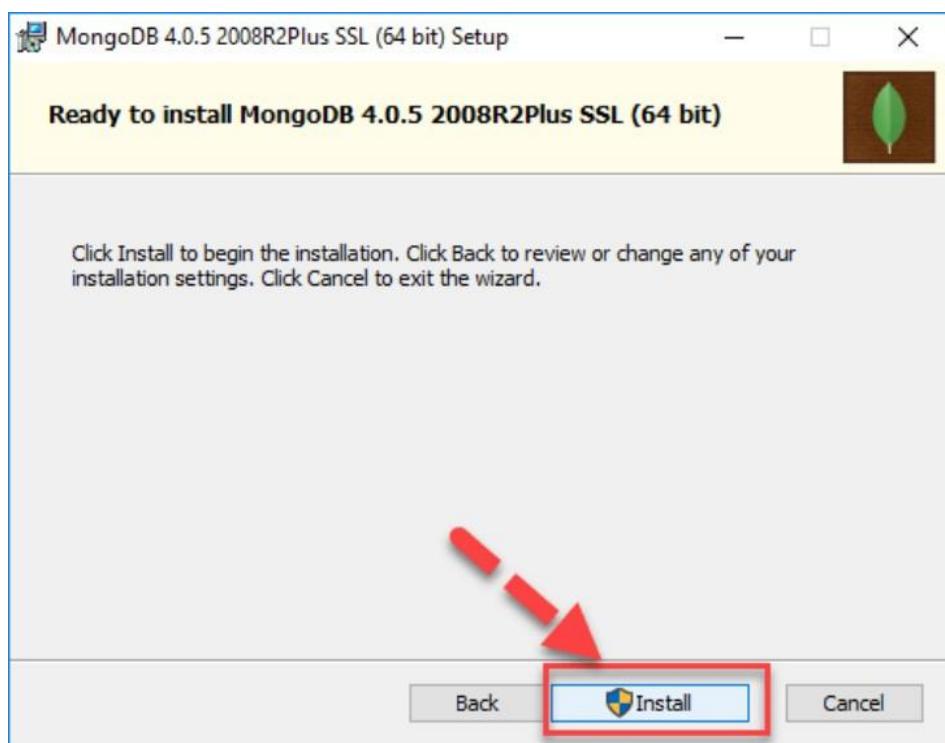
**Step 5)**

1. Select “Run service as Network Service user”. make a note of the data directory, we’ll need this later.

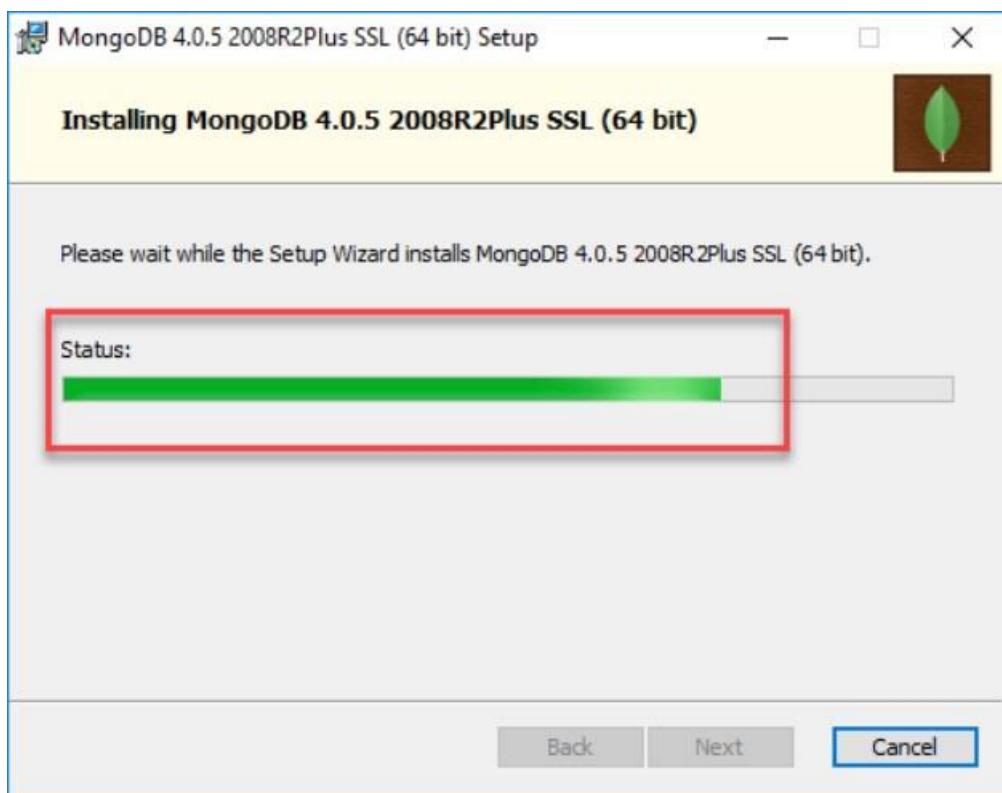
2. Click Next



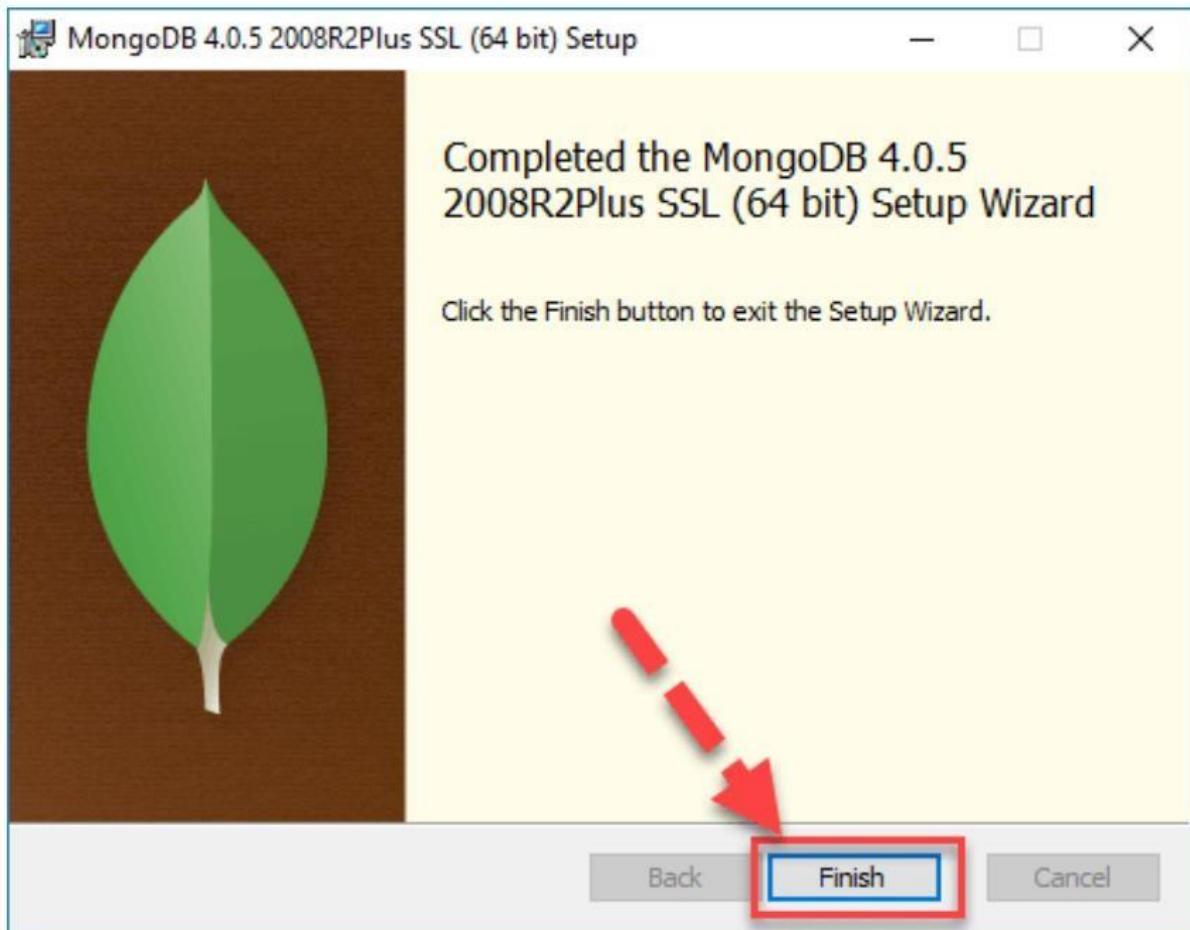
Step 6) Click on the Install button to start the installation.



Step 7) Installation begins. Click Next once completed.



Step 8) Click on the Finish button to complete the installation.



## Test Mongodb

**Step 1)** Go to " C:\Program Files\MongoDB\Server\4.0\bin" and double click on **mongo.exe**. Alternatively, you can also click on the MongoDB desktop icon.

- **Create the directory where MongoDB will store its files.**  
Open command prompt window and apply following commands

```
C:\users\admin> cd\
C:\>md data\db
```

## Step 2) Execute mongodb

Open another command prompt window.

```
C:\> cd C:\Program Files\MongoDB\Server\4.0\bin
C:\Program Files\MongoDB\Server\4.0\bin> mongod
```

*In case if it gives an error then run the following command:*

```
C:\Program Files\MongoDB\Server\4.0\bin> mongod --repair
```

```
{ v: 2, key: { lastUse: 1 }, name: "lsidTTLIndex", ns: "config.system.sessions", expireAfterSeconds: 1800 }
2023-03-03T09:46:21.011+0530 I INDEX [LogicalSessionCacheRefresh] building index using bulk method; build may temporarily use up to 500 megabytes of RAM
2023-03-03T09:46:21.044+0530 I INDEX [LogicalSessionCacheRefresh] build index done. scanned 0 total records. 0 secs
2023-03-03T09:46:21.045+0530 I COMMAND [LogicalSessionCacheRefresh] command config.$cmd command: createIndexes { createIndexes: "system.sessions", indexes: [ { key: { lastUse: 1 }, name: "lsidTTLIndex", expireAfterSeconds: 1800 } ], $db: "config" } numYields:0 reslen:114 locks:{ Global: { acquireCount: { r: 2, w: 2 } }, Database: { acquireCount: { w: 2, W: 1 } }, Collection: { acquireCount: { w: 2 } } } protocol:op_msg 254ms
```

## Step 3) Connect to MongoDB using the Mongo shell

Let the MongoDB daemon to run.

Open another command prompt window and run the following commands:

```
C:\users\admin> cd C:\Program Files\MongoDB\Server\4.0\bin
C:\Program Files\MongoDB\Server\4.0\bin> mongo
```

```
The monitoring data will be available on a MongoDB website with a unique URL accessible to you and anyone you share the URL with. MongoDB may use this information to make product improvements and to suggest MongoDB products and deployment options to you.
```

```
To enable free monitoring, run the following command: db.enableFreeMonitoring()
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()
---
```

```
> show dbs
admin      0.000GB
config     0.000GB
local      0.000GB
mybigdata  0.000GB
> use mybigdata
```

## Step 4) Install PyMongo

Open another command prompt window and run the following commands:

Check the python version on your desktop / laptop and copy that path from window explorer

```
C:\users\admin>cd C:\Program Files\Python311\Scripts
C:\Program Files\<Python38>\Scripts > python -m pip install pymongo
```

```
C:\Program Files\Python38\Scripts>python -m pip install pymongo
Defaulting to user installation because normal site-packages is not writeable
Collecting pymongo
  Downloading pymongo-4.3.3-cp38-cp38-win_amd64.whl (382 kB)
    |██████████| 382 kB 3.3 MB/s
Collecting dnspython<3.0.0,>=1.16.0
  Downloading dnspython-2.3.0-py3-none-any.whl (283 kB)
    |██████████| 283 kB ...
Installing collected packages: dnspython, pymongo
Successfully installed dnspython-2.3.0 pymongo-4.3.3
WARNING: You are using pip version 20.2.1; however, version 23.0.1 is available.
You should consider upgrading via the 'C:\Program Files\Python38\python.exe -m pip install --upgrade pip' command.
```

Note: # **-m** option is for <module-name>  
 Now you have downloaded and installed a mongoDB driver.

### **Step 5) Test PyMongo**

Run the following command from python command prompt

```
import pymongo
```

Now, either create a file in Python IDLE or run all commands one by one in sequence on Python cell

#### **Program 1: Creating a Database: create\_dp.py**

```
import pymongo
myclient = pymongo.MongoClient("mongodb://localhost:27017/")
mydb = myclient["mybigdata"]
print(myclient.list_database_names())
['admin', 'config', 'local']
```

#### **Program 2: Creating a Collection: create\_collection.py**

```
import pymongo
myclient = pymongo.MongoClient("mongodb://localhost:27017/")
mydb = myclient["mybigdata"]
mycol= mydb["student"]
print(mydb.list_collection_names())
[]
```

#### **Program 3: Insert into Collection: insert\_into\_collection.py**

```
import pymongo
myclient = pymongo.MongoClient("mongodb://localhost:27017/")
mydb = myclient["mybigdata"]
mycol= mydb["student"]
mydict= {"name": "Beena", "address": "Mumbai"}
x= mycol.insert_one(mydict) # insert_one(containing the name(s) and value(s) of each field)
```

#### **Program 4: Insert Multiple data into Collection: insert\_many.py**

```
import pymongo
myclient = pymongo.MongoClient("mongodb://localhost:27017/")
mydb = myclient["mybigdata"]
mycol= mydb["student"]
mylist=[{"name": "Khyati", "address": "Mumbai"}, {"name": "Kruti", "address": "Mumbai"}, {"name": "Nidhi", "address": "Pune"}, {"name": "Komal", "address": "Pune"},]
x= mycol.insert_many(mylist)
```

**Step 6) Test in Mongodt to check database and data inserted in collection**

a. If you want to check your database list, use the command show dbs in mongo command prompt

```
> show dbs
```

```
admin      0.000GB
config     0.000GB
local      0.000GB
mybigdata  0.000GB
```

b. If you want to use a database with name mybigdata, then use database statement would be as follow:

```
> use mybigdata
```

```
switched to db mybigdata
```

c. If you want to check collection in mongodt use the command show collections

```
> show collections
```

```
student
```

d. If you want to display the first row from collection: db.collection\_name.find()

```
> db.student.findOne()
```

```
> db.student.findOne()
{
  "_id" : ObjectId("640178face663db608cef72f"),
  "name" : "Beena",
  "address" : "Mumbai"
}
```

e. If you want to display all the data from collection: db.collection\_name.find()

```
> db.student.find()
```

```
> db.student.find()
[{"_id": ObjectId("640178face663db608cef72f"), "name": "Beena", "address": "Mumbai"}, {"_id": ObjectId("640179336ce317082c266dc1"), "name": "Khyati", "address": "Mumbai"}, {"_id": ObjectId("640179336ce317082c266dc2"), "name": "Kruti", "address": "Mumbai"}, {"_id": ObjectId("640179336ce317082c266dc3"), "name": "Nidhi", "address": "Pune"}, {"_id": ObjectId("640179336ce317082c266dc4"), "name": "Komal", "address": "Pune"}]
```

f. count number of rows in a collection

```
> db.student.count()
```

```
5
```

### **Site for R packages documentation:**

[https://cran.r-project.org/web/packages/available\\_packages\\_by\\_name.html](https://cran.r-project.org/web/packages/available_packages_by_name.html)

## Practical No. 04

**Aim: A) Implement Decision Tree classification technique using Social\_Network\_Ads.csv dataset.**

**Code:**

```
# Decision Tree Classification
```

```
# Importing the dataset
```

```
dataset = read.csv("D:\\bda prac\\Social_Network_Ads.csv")
```

```
#print(dataset)
```

```
dataset = dataset[3:5] # columns 3 4 ad 5
```

```
print(dataset)
```

```
# Encoding the target feature as factor(just like a vector having levels
```

```
# levels to convey that only two possible values for purchased - 0 & 1
```

```
dataset$Purchased = factor(dataset$Purchased, levels = c(0, 1))
```

```
print (dataset$Purchased)
```

```
# Splitting the dataset into the Training set and Test set
```

```
install.packages('caTools')
```

```
library(caTools)
```

```
set.seed(123)
```

```
#split = sample.split(dataset$Purchased, SplitRatio = 0.75)
```

```
split = sample.split(dataset$Purchased, SplitRatio = 0.75)
```

```
training_set = subset(dataset, split == TRUE)
```

```
test_set = subset(dataset, split == FALSE)
```

```
# Feature Scaling - scale() method centers and/or scales the columns of a numeric matrix.
```

```
training_set[-3] = scale(training_set[-3]) # scaling first 2 columns, don't consider 3rd column
```

```
test_set[-3] = scale(test_set[-3])
```

```
#print(test_set[-3])
```

```
# Fitting Decision Tree Classification to the Training set
```

```
install.packages('rpart')
```

```
library(rpart) # for partitioning tree
```

```
install.packages('rpart.plot')
```

```
library(rpart.plot)
```

```
classifier = rpart(formula = Purchased ~ ., data = training_set)
```

```
# Predicting the Test set results
```

```
y_pred = predict(classifier, newdata = test_set[-3], type = 'class')
```

```
print(y_pred)
```

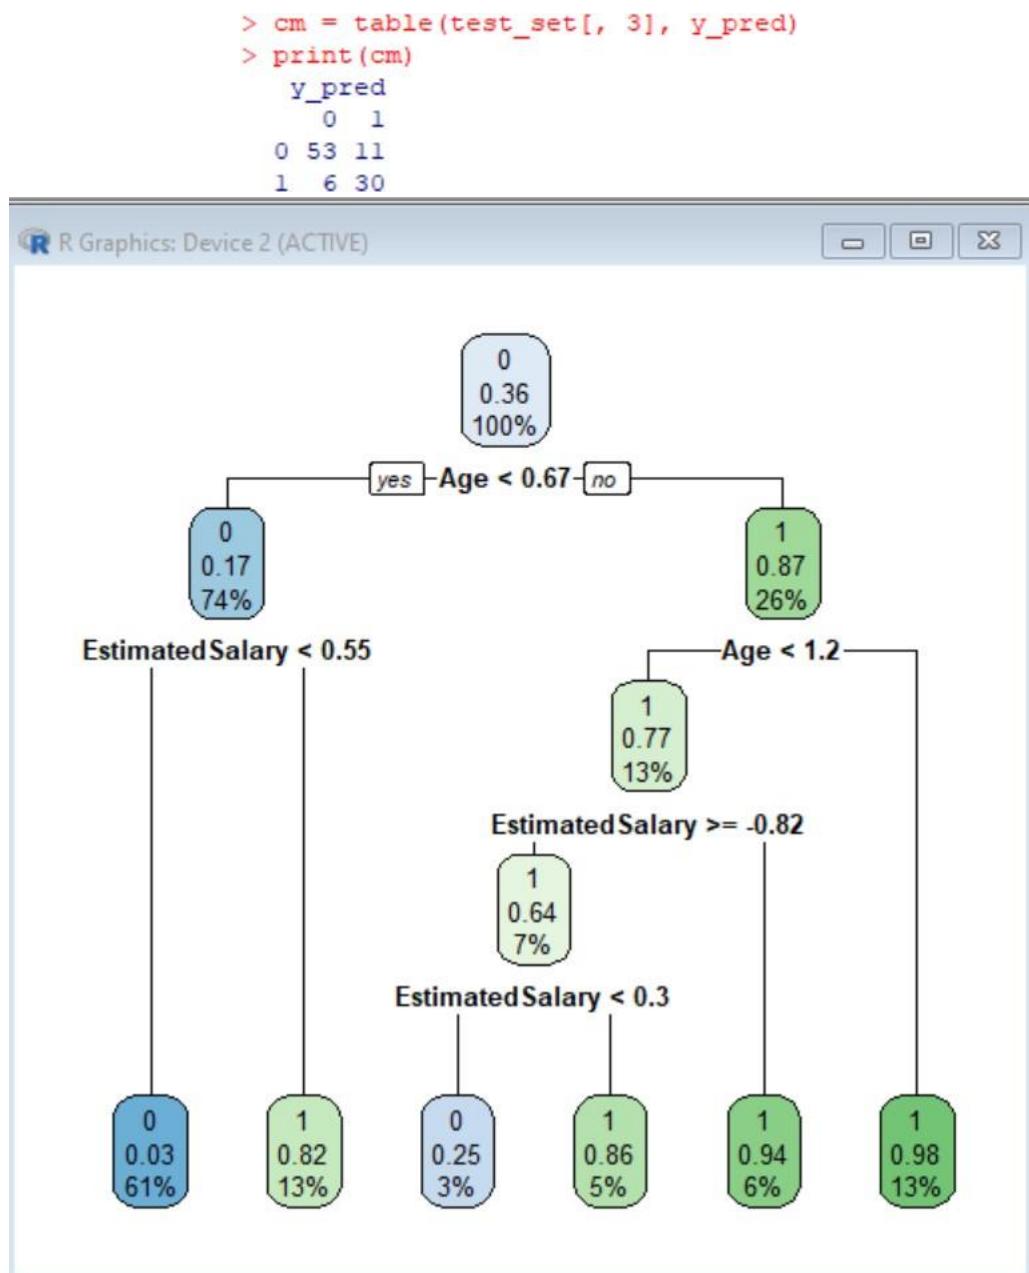
```
# Making the Confusion Matrix
```

```
cm = table(test_set[, 3], y_pred)
```

```
print(cm)
```

```
y_grid = predict(classifier, newdata = grid_set, type = 'class')
```

```
# Plotting the tree
#extra=106 class model with a binary response
#extra=104 class model with a response having more than two levels
rpart.plot(classifier, extra = 106)
```

**Output:**

## Practical No. 5

**Aim: B) Implement SVM Classification technique using Social\_Network\_Ads.csv dataset. Evaluate the performance of classifier.**

**Code:**

```
# Support Vector Machine (SVM)

# Importing the dataset
dataset = read.csv("D:\\bda prac\\Social_Network_Ads.csv")
dataset = dataset[3:5]
print(dataset)
print(dataset$Purchased)
# Splitting the dataset into the Training set and Test set
install.packages('caTools')
library(caTools)
set.seed(123)
split = sample.split(dataset$Purchased, SplitRatio = 0.75)
training_set = subset(dataset, split == TRUE)
print(training_set)
test_set = subset(dataset, split == FALSE)
print(test_set)
# Feature Scaling
training_set[-3] = scale(training_set[-3]) # [-3] means 3rd index will be dropped
test_set[-3] = scale(test_set[-3])
print(training_set[-3])
print (test_set[-3])
# Fitting SVM to the Training set
install.packages('e1071')
library(e1071)
classifier = svm(formula = Purchased ~ .,
                 data = training_set,
                 type = 'C-classification',
                 kernel = 'linear')
print (classifier)

# Predicting the Test set results
y_pred = predict(classifier, newdata = test_set[-3])
print(y_pred)

# Making the Confusion Matrix
cm = table(test_set[, 3], y_pred)
print (cm)
```

**Output:**

```
> cm = table(test_set[, 3], y_pred)
> print (cm)
      y_pred
      0   1
0 57   7
1 13  23
```

---

## Practical No. 6

**AIM : REGRESSION MODEL Import a data from web storage.**

**Name the dataset and now do Logistic Regression to find out relation between variables that are affecting the admission of a student in an institute based on his or her GRE score, GPA obtained and rank of the student. Also check the model is fit or not. require (foreign), require(MASS).**

.

**Code:**

```
#fetch the data
##fetch the data
college <- read.csv("https://raw.githubusercontent.com/csquared/udacity-
dlnd/master/nn/binary.csv")
head(college)
nrow(college)

install.packages("caTools") # For Logistic regression
library(caTools)

split <- sample.split(college, SplitRatio = 0.75)
split

training_reg <- subset(college, split == "TRUE")
test_reg <- subset(college, split == "FALSE")

# Training model
fit_logistic_model <- glm(admit ~.,
                           data = training_reg,
                           family = "binomial")

# Predict test data based on model
predict_reg <- predict(fit_logistic_model,
                        test_reg, type = "response")
predict_reg

cdplot(as.factor(admit)~ gpa, data=college)
cdplot(as.factor(admit)~ gre, data=college)
cdplot(as.factor(admit)~ rank, data=college)
```

```
# Changing probabilities
predict_reg <- ifelse(predict_reg > 0.5, 1, 0)
predict_reg
```

```
# Evaluating model accuracy
# using confusion matrix
table(test_reg$admit, predict_reg)
```

## Output:

MODEL -

### Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-3.1450	0.5820	-5.408	6.46e-08	***
age	0.0210	0.0123	1.707	0.0885	.
duration	0.0420	0.0042	10.000	< 2e-16	***
campaign	-0.0320	0.0118	-2.716	0.0067	**
pdays	-0.0015	0.0012	-1.250	0.2114	
previous	-0.0210	0.0161	-1.304	0.1926	

CONFUSION MATRIX-

### Confusion Matrix and Statistics

		Reference	
		0	1
Prediction	0	50	20
	1	8	30

```
Accuracy : 0.75
Precision : 0.60
Recall : 0.78
F1-Score : 0.68
```

## Practical No. 07

**Apply multiple regressions, if data have a continuous independent variable. Apply on above dataset – binary.csv.**

**Code:**

```
#fetch the data
college <- read.csv("https://raw.githubusercontent.com/csquared/udacity-dlnd/master/nn/binary.csv")
head(college)
nrow(college)

install.packages("caTools") # For Logistic regression
library(caTools)

split <- sample.split(college, SplitRatio = 0.75)
split

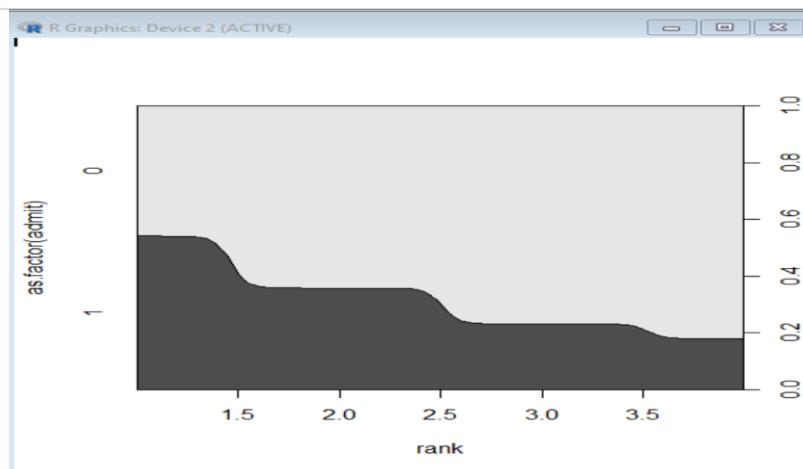
training_reg <- subset(college, split == "TRUE")
test_reg <- subset(college, split == "FALSE")

# Training model
fit_MRegressor_model <- lm(formula = admit ~ gre+gpa+rank,
```

```
data = training_reg)

# Predict test data based on model
predict_reg <- predict(fit_MRegressor_model,
                       newdata = test_reg)
predict_reg

cdplot(as.factor(admit)~ gpa, data=college)
cdplot(as.factor(admit)~ gre, data=college)
cdplot(as.factor(admit)~ rank, data=college)
```

**Output:**

---

## Practical No. 08

**Aim: CLASSIFICATION MODEL** a. Install relevant package for classification. b. Choose classifier for classification problem. c. Evaluate the performance of classifier

**Code:**

```
# # Naive Bayes
# Importing the dataset
dataset = read.csv('C:\\2022-23\\BDA practical 2023\\Social_Network_Ads.csv')
dataset = dataset[3:5]
# Encoding the target feature as factor
dataset$Purchased = factor(dataset$Purchased, levels = c(0, 1))
# Splitting the dataset into the Training set and Test set
#install.packages('caTools')
library(caTools)
set.seed(123)
split = sample.split(dataset$Purchased, SplitRatio = 0.75)
training_set = subset(dataset, split == TRUE)
test_set = subset(dataset, split == FALSE)
# Feature Scaling
training_set[-3] = scale(training_set[-3])
test_set[-3] = scale(test_set[-3])
# Fitting Naive Bayes to the Training set
install.packages('e1071')
library(e1071)
classifier = naiveBayes(x = training_set[-3],
                        y = training_set$Purchased)
# Predicting the Test set results
y_pred = predict(classifier, newdata = test_set[-3])
# Making the Confusion Matrix
cm = table(test_set[, 3], y_pred)
print(cm)
```

**Output:**

```
Accuracy: 0.98
Confusion Matrix:
 [[14  0  0]
 [ 0 16  1]
 [ 0  0 13]]
Classification Report:
      precision    recall  f1-score   support
          0       1.00     1.00     1.00      14
          1       1.00     0.94     0.97      17
          2       0.93     1.00     0.96      13
                                              accuracy         0.98
                                              macro avg     0.98     0.98      44
                                              weighted avg  0.98     0.98      44
```

---

## Practical No. 09

**Aim: A) Clustering algorithms for unsupervised classification. Read a datafile all\_Customers.csv and apply k-means clustering. Plot the cluster data using R visualizations.**

**Code:**

```
# K-Means Clustering
```

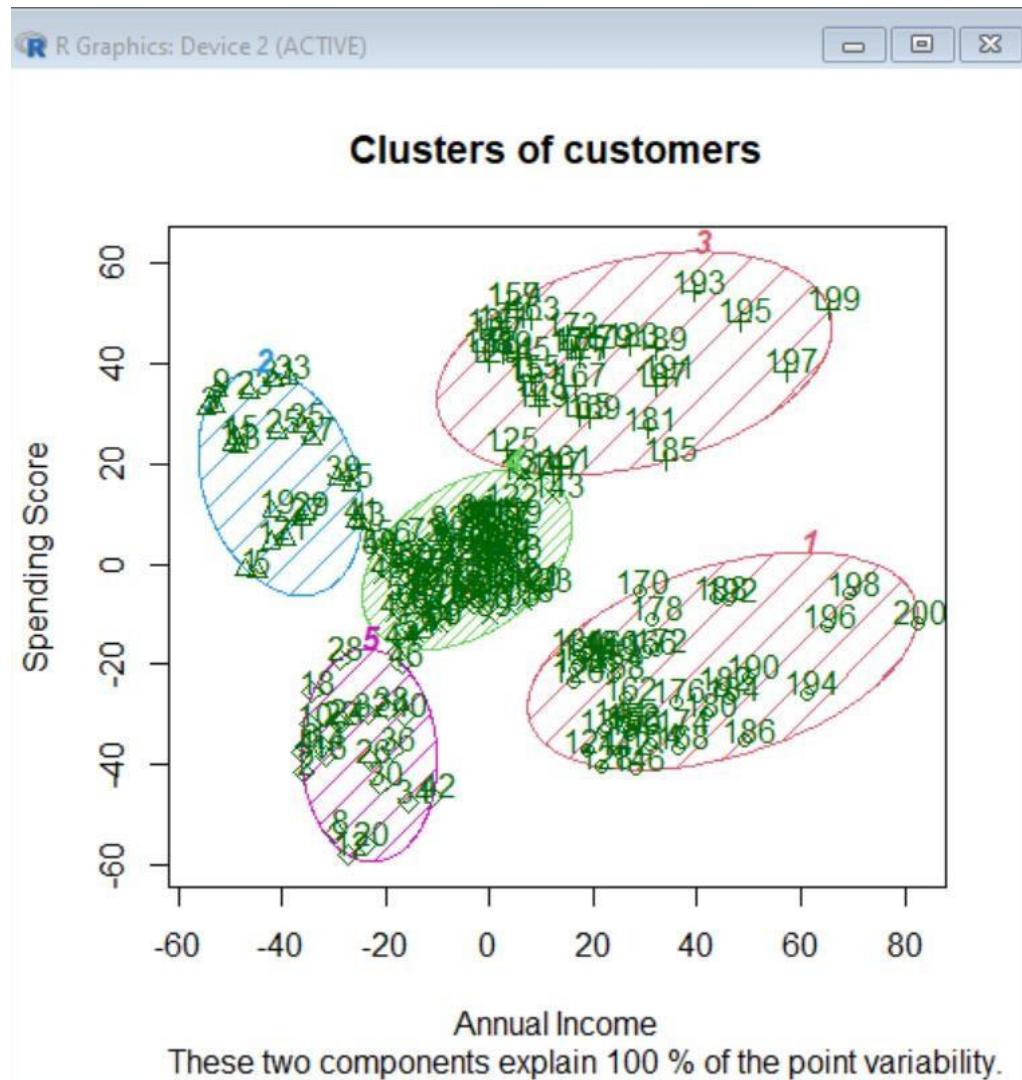
```
# Importing the dataset
dataset = read.csv("D:\\bda prac\\Mall_Customers.csv")
head(dataset)
dataset = dataset[4:5]
head(dataset)
```

```
wcss = vector()
for (i in 1:10) wcss[i] = sum(kmeans(dataset, i)$withinss)
plot(1:10,
     wcss,
     type = 'b',
     main = paste('The Elbow Method'),
     xlab = 'Number of clusters',
     ylab = 'WSS')
```

```
# Fitting K-Means to the dataset with no of clusters = 5
kmeans = kmeans(x = dataset, centers = 5)
y_kmeans = kmeans$cluster
```

```
# Visualising the clusters
library(cluster)
clusplot(dataset,
          y_kmeans,
          lines = 0,
          shade = TRUE,
          color = TRUE,
          labels = 2,
          main = paste('Clusters of customers'),
          xlab = 'Annual Income',
          ylab = 'Spending Score')
```

---

**Output:**



**University of Mumbai**

**Practical Journal of  
Modern Networking**

**M.Sc. (Information Technology) Part-I**

**Submitted by**

**SINGH MANASI RAKESH**

**Seat No: 1312502**



**DEPARTMENT OF INFORMATION TECHNOLOGY  
PILLAI HOC COLLEGE OF ARTS, SCIENCE & COMMERCE, RASAYANI  
(Affiliated to Mumbai University)  
RASAYANI, 410207  
MAHARASHTRA  
2023-2024**

**Mahatma Education Society's  
Pillai Hoc College of Arts, Science & Commerce, Rasayani  
(Affiliated to Mumbai University)  
RASAYANI – MAHARASHTRA - 410207**

**DEPARTMENT OF  
INFORMATION TECHNOLOGY**



**CERTIFICATE**

This is to certify that the experiment work entered in this journal is as per the syllabus in **M.Sc. (Information Technology) Part-I, Semester-II**; class prescribed by University of Mumbai for the subject **Modern Networking** was done in computer lab of Mahatma Education Society's Pillai HOC College of Arts, Science & Commerce, Rasayani by **MANASI SINGH** during Academic year 2023-2024.

**Exam Seat No: 1312502**

**Subject In-Charge**

**Coordinator**

**External Examiner**

**Principal**

**Date:**

**College Seal**

# **MODERN NETWORKING**

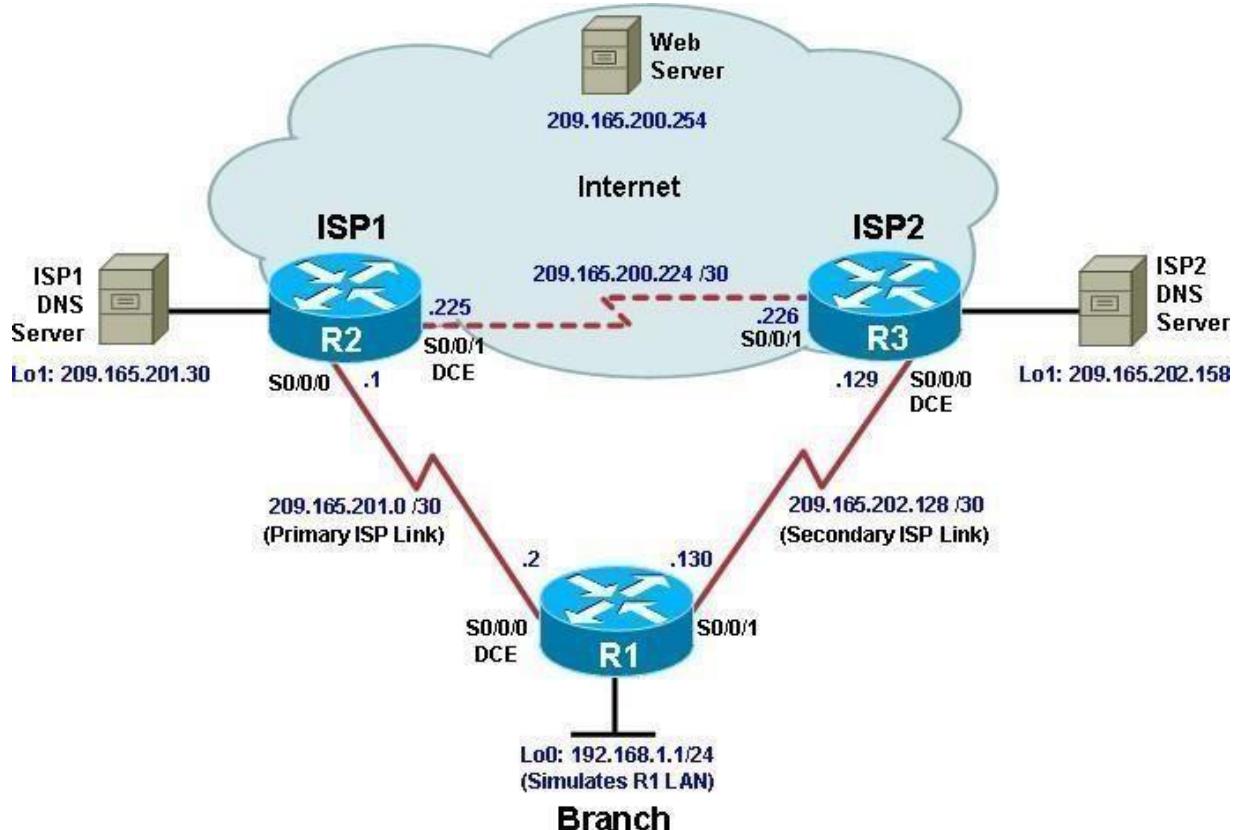
## INDEX

Practical No.	Title	Page No.
1	<b>Configure IP SLA Tracking and Path Control</b>	01
2	<b>Use the AS_PATH attribute to filter BGP routes</b>	14
3	<b>Configuring IBGP and EBGP Sessions, Local Preference, and MED</b>	20
4	<b>Secure the Management Plane</b>	41
5	<b>Configure and Verify Path Control</b>	52
6	<b>Configure IP SLA Tracking and Path Control with gateway</b>	64
7	<b>Configuring Basic MPLS Using OSPF</b>	77

## Practical No. 01

**Aim:** Configure IP SLA Tracking and Path Control

### Topology



### Objectives

- Configure and verify the IP SLA feature.
- Test the IP SLA tracking feature.
- Verify the configuration and operation using **show** and **debug** commands.

### Background

You want to experiment with the Cisco IP Service Level Agreement (SLA) feature to study how it could be of value to your organization.

At times, a link to an ISP could be operational, yet users cannot connect to any other outside Internet resources. The problem might be with the ISP or downstream from them. Although policy-based routing (PBR) can be implemented to alter path control, you will implement the Cisco IOS SLA feature to monitor this behavior and intervene by injecting another default route to a backup ISP.

To test this, you have set up a three-router topology in a lab environment. Router R1 represents a branch office connected to two different ISPs. ISP1 is the preferred connection to the Internet, while ISP2 provides a backup link. ISP1 and ISP2 can also interconnect, and both can reach the web server. To monitor ISP1 for

failure, you will configure IP SLA probes to track the reachability to the ISP1 DNS server. If connectivity to the ISP1 server fails, the SLA probes detect the failure and alter the default static route to point to the ISP2 server.

**Note:** This lab uses Cisco 1841 routers with Cisco IOS Release 12.4(24)T1 and the Advanced IP Services image c1841-advpipservicesk9-mz.124-24.T1.bin. You can use other routers (such as a 2801 or 2811) and Cisco IOS Software versions if they have comparable capabilities and features. Depending on the router and Cisco IOS Software version, the commands available and output produced might vary from what is shown in this lab.

## Required Resources

- 3 routers (Cisco 1841 with Cisco IOS Release 12.4(24)T1 Advanced IP Services or comparable)
- Serial and console cables

## Step 1: Prepare the routers and configure the router hostname and interface addresses.

- Cable the network as shown in the topology diagram. Erase the startup configuration and reload each router to clear the previous configurations. Using the addressing scheme in the diagram, create the loopback interfaces and apply IP addresses to them as well as the serial interfaces on R1, ISP1, and ISP2.

You can copy and paste the following configurations into your routers to begin.

**Note:** Depending on the router model, interfaces might be numbered differently than those listed. You might need to alter them accordingly.

### Router R1

```
hostname R1

interface Loopback 0
description R1 LAN
ip address 192.168.1.1 255.255.255.0

interface Serial0/0/0
description R1 --> ISP1
ip address 209.165.201.2 255.255.255.252
clock rate 128000
bandwidth 128
no shutdown

interface Serial0/0/1
description R1 --> ISP2
ip address 209.165.202.130 255.255.255.252
bandwidth 128
no shutdown
```

### Router ISP1 (R2)

```
hostname ISP1

interface Loopback0
description Simulated Internet Web Server
ip address 209.165.200.254 255.255.255.255

interface Loopback1
description ISP1 DNS Server
```

```

ip address 209.165.201.30 255.255.255.255
interface Serial0/0/0 description ISP1 --> R1
ip address 209.165.201.1 255.255.255.252
bandwidth 128
no shutdown

interface Serial0/0/1
description ISP1 --> ISP2
ip address 209.165.200.225 255.255.255.252
clock rate 128000
bandwidth 128
no shutdown

```

### **Router ISP2 (R3)**

```

hostname ISP2

interface Loopback0
description Simulated Internet Web Server
ip address 209.165.200.254 255.255.255.255

interface Loopback1
description ISP2 DNS Server
ip address 209.165.202.158 255.255.255.255

interface Serial0/0/0
description ISP2 --> R1
ip address 209.165.202.129 255.255.255.252
clock rate 128000
bandwidth 128
no shutdown

interface Serial0/0/1
description ISP2 --> ISP1
ip address 209.165.200.226 255.255.255.252
bandwidth 128
no shutdown

```

- b. Verify the configuration by using the **show interfaces description** command. The output from router R1 is shown here as an example.

```
R1# show interfaces description
Interface                  Status      Protocol Description
Fa0/0                     admin down   down
Fa0/1                     admin down   down
Se0/0/0                   up          up        R1 --> ISP1
Se0/0/1                   up          up        R1 --> ISP2
Lo0                      up          up        R1 LAN
```

All three interfaces should be active. Troubleshoot if necessary.

- c. The current routing policy in the topology is as follows:

- Router R1 establishes connectivity to the Internet through ISP1 using a default static route.
- ISP1 and ISP2 have dynamic routing enabled between them, advertising their respective public address pools.
- ISP1 and ISP2 both have static routes back to the ISP LAN.

**Note:** For the purpose of this lab, the ISPs have a static route to an RFC 1918 private network address on the branch router R1. In an actual branch implementation, Network Address Translation (NAT) would be configured for all traffic exiting the branch LAN. Therefore, the static routes on the ISP routers would be pointing to the provided public pool of the branch office. This is covered in Lab 7-1, “Configure Routing Facilities to the Branch Office.”

Implement the routing policies on the respective routers. You can copy and paste the following configurations.

### Router R1

```
ip route 0.0.0.0 0.0.0.0 209.165.201.1
```

### Router ISP1 (R2)

```
router eigrp 1
network 209.165.200.224 0.0.0.3
network 209.165.201.0 0.0.0.31
no auto-summary
```

```
ip route 192.168.1.0 255.255.255.0 209.165.201.2
```

### Router ISP2 (R3)

```
router eigrp 1
network 209.165.200.224 0.0.0.3
network 209.165.202.128 0.0.0.31
no auto-summary
```

```
ip route 192.168.1.0 255.255.255.0 209.165.202.130
```

EIGRP neighbor relationship messages on ISP1 and ISP2 should be generated. Troubleshoot if necessary.

```
%DUAL-5-NBRCHANGE: IP-EIGRP(0) 1: Neighbor 209.165.200.225 (Serial0/0/1) is up: new adjacency
```

## Step 2: Verify server reachability.

The Cisco IOS IP SLA feature enables an administrator to monitor network performance between Cisco devices (switches or routers) or from a Cisco device to a remote IP device. IP SLA probes continuously check the reachability of a specific destination, such as a provider edge router interface, the DNS server of the ISP, or any other specific destination, and can conditionally announce a default route only if the connectivity is verified.

- Before implementing the Cisco IOS SLA feature, you must verify reachability to the Internet servers. From router R1, ping the web server, ISP1 DNS server, and ISP2 DNS server to verify connectivity. You can copy the following Tcl script and paste it into R1.

```
foreach address {
209.165.200.254
209.165.201.30
209.165.202.158
} {
ping $address source 192.168.1.1
}

R1(tcl) # foreach address {
+>(tcl) # 209.165.200.254.
```

```
+>(tcl) # 209.165.201.30
+>(tcl) # 209.165.202.158
+>(tcl) # } {
+>(tcl) # ping $address source 192.168.1.1
+>(tcl) #}

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 209.165.200.254, timeout is 2 seconds:
Packet sent with a source address of 192.168.1.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 12/15/16 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 209.165.201.30, timeout is 2 seconds:
Packet sent with a source address of 192.168.1.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 12/14/16 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 209.165.202.158, timeout is 2 seconds:
Packet sent with a source address of 192.168.1.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 20/21/24 ms
```

- b. Trace the path taken to the web server, ISP1 DNS server, and ISP2 DNS server. You can copy the following Tcl script and paste it into R1.

```
foreach address {
209.165.200.254
209.165.201.30
209.165.202.158
} {
trace $address source 192.168.1.1
}

R1(tcl) # foreach address {
+>(tcl) # 209.165.200.254
+>(tcl) # 209.165.201.30
+>(tcl) # 209.165.202.158
+>(tcl) # } {
+>(tcl) # trace $address source 192.168.1.1
+>(tcl) # }
```

Type escape sequence to abort.  
Tracing the route to 209.165.200.254

```
1 209.165.201.1 20 msec 8 msec *
Type escape sequence to abort.
Tracing the route to 209.165.201.30

1 209.165.201.1 8 msec 8 msec *
Type escape sequence to abort.
Tracing the route to 209.165.202.158

1 209.165.201.1 8 msec 8 msec 4 msec
2 209.165.200.226 8 msec 8 msec *
```

Through which ISP is traffic flowing?

### Step 3: Configure IP SLA probes.

When the reachability tests are successful, you can configure the Cisco IOS IP SLAs probes. Different types of probes can be created, including FTP, HTTP, and jitter probes. In this scenario, you will configure ICMP echo probes.

- Create an ICMP echo probe on R1 to the primary DNS server on ISP1 using the **ip sla** command.

**Note:** With Cisco IOS Release 12.4(4)T, 12.2(33)SB, and 12.2(33)SXI, the **ip sla** command has replaced the previous **ip sla monitor** command. In addition, the **icmp-echo** command has replaced the **type echo protocol ipICMPEcho** command.

```
R1(config)# ip sla 11
R1(config-ip-sla)# icmp-echo 209.165.201.30
R1(config-ip-sla-echo)# frequency 10
R1(config-ip-sla-echo)# exit
R1(config)# ip sla schedule 11 life forever start-time now
```

The operation number of 11 is only locally significant to the router. The **frequency 10** command schedules the connectivity test to repeat every 10 seconds. The probe is scheduled to start now and to run forever.

- Verify the IP SLAs configuration of operation 11 using the **show ip sla configuration 11** command.

**Note:** With Cisco IOS Release 12.4(4)T, 12.2(33)SB, and 12.2(33)SXI, the **show ip sla configuration** command has replaced the **show ip sla monitor configuration** command.

```
R1# show ip sla configuration 11
IP SLAs, Infrastructure Engine-II.
Entry number: 11
Owner:
Tag:
Type of operation to perform: icmp-echo
Target address/Source address: 209.165.201.30/0.0.0.0
Type Of Service parameter: 0x0
Request size (ARR data portion): 28
Operation timeout (milliseconds): 5000
Verify data: No
Vrf Name:
Schedule:
  Operation frequency (seconds): 10 (not considered if randomly scheduled)
  Next Scheduled Start Time: Start Time already passed
  Group Scheduled : FALSE
  Randomly Scheduled : FALSE
  Life (seconds): Forever
  Entry Ageout (seconds): never
  Recurring (Starting Everyday): FALSE
  Status of entry (SNMP RowStatus): Active
Threshold (milliseconds): 5000 (not considered if react RTT is configured)
Distribution Statistics:
  Number of statistic hours kept: 2
  Number of statistic distribution buckets kept: 1
  Statistic distribution interval (milliseconds): 20
History Statistics:
  Number of history Lives kept: 0
  Number of history Buckets kept: 15
  History Filter Type: None
Enhanced History:
```

The output lists the details of the configuration of operation 11. The operation is an ICMP echo to 209.165.201.30, with a frequency of 10 seconds, and it has already started (the start time has already passed).

- c. Issue the **show ip sla statistics** command to display the number of successes, failures, and results of the latest operations.

**Note:** With Cisco IOS Release 12.4(4)T, 12.2(33)SB, and 12.2(33)SXI, the **show ip sla statistics** command has replaced the **show ip sla monitor statistics** command.

```
R1# show ip sla statistics
IPSLAs Latest Operation Statistics

IPSLA operation id: 11
```

```
Latest operation start time: *21:22:29.707 UTC Fri Apr 2 2010
Latest operation return code: OK
Number of successes: 5
Number of failures: 0
Operation time to live: Forever
```

You can see that operation 11 has already succeeded five times, has had no failures, and the last operation returned an OK result.

- d. Although not actually required because IP SLA session 11 alone could provide the desired fault tolerance, create a second probe, 22, to test connectivity to the second DNS server located on router ISP2. You can copy and paste the following commands on R1.

```
ip sla 22
icmp-echo 209.165.202.158
frequency 10
exit
ip sla schedule 22 life forever start-time now
```

- e. Verify the new probe using the **show ip sla configuration** and **show ip sla statistics** commands.

```
R1# show ip sla configuration 22
IP SLAs, Infrastructure Engine-II.
Entry number: 22
Owner:
Tag:
Type of operation to perform: icmp-echo
Target address/Source address: 209.165.201.158/0.0.0.0
Type Of Service parameter: 0x0
Request size (ARR data portion): 28
Operation timeout (milliseconds): 5000
Verify data: No
Vrf Name:
Schedule:
    Operation frequency (seconds): 10 (not considered if randomly scheduled)
    Next Scheduled Start Time: Start Time already passed
    Group Scheduled : FALSE
    Randomly Scheduled : FALSE
    Life (seconds): Forever
    Entry Ageout (seconds): never
    Recurring (Starting Everyday): FALSE
    Status of entry (SNMP RowStatus): Active
Threshold (milliseconds): 5000 (not considered if react RTT is configured)
Distribution Statistics:
    Number of statistic hours kept: 2
```

```

Number of statistic distribution buckets kept: 1 Statistic distribution interval
(milliseconds): 20
History Statistics:
  Number of history Lives kept: 0
  Number of history Buckets kept: 15
  History Filter Type: None
Enhanced History:

R1# show ip sla statistics 22
IPSLAs Latest Operation Statistics

IPSLA operation id: 22

Latest operation start time: *21:24:14.215 UTC Fri Apr 2 2010
Latest operation return code: OK
Number of successes: 4
Number of failures: 0
Operation time to live: Forever

```

The output lists the details of the configuration of operation 22. The operation is an ICMP echo to 209.165.202.158, with a frequency of 10 seconds, and it has already started (the start time has already passed). The statistics also prove that operation 22 is active.

#### **Step 4: Configure tracking options.**

Although PBR could be used, you will configure a floating static route that appears or disappears depending on the success or failure of the IP SLA.

- Remove the current default route on R1, and replace it with a floating static route having an administrative distance of 5.

```

R1(config)# no ip route 0.0.0.0 0.0.0.0 209.165.201.1
R1(config)# ip route 0.0.0.0 0.0.0.0 209.165.201.1 5
R1(config)# exit

```

- Verify the routing table.

```

R1# show ip route
*Apr 2 20:00:37.367: %SYS-5-CONFIG_I: Configured from console by console
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static
route
      o - ODR, P - periodic downloaded static route

Gateway of last resort is 209.165.201.1 to network 0.0.0.0

```

```

  209.165.201.0/30 is subnetted, 1 subnets
C 209.165.201.0  is directly connected, Serial0/0/0
  209.165.202.0/30 is subnetted, 1 subnets
C 209.165.202.128 is directly connected, Serial0/0/1C
192.168.1.0/24 is directly connected, FastEthernet0/0  S*
0.0.0.0/0 [5/0] via 209.165.201.1

```

Notice that the default static route is now using the route with the administrative distance of 5. The first tracking object is tied to IP SLA object 11..

- c. Use the **track 1 ip sla 11 reachability** command to enter the config-track subconfiguration mode.

**Note:** With Cisco IOS Release 12.4(20)T, 12.2(33)SXI1, and 12.2(33)SRE and Cisco IOS XE Release 2.4, the **track ip sla** command has replaced the **track rtr** command.

```
R1(config)# track 1 ip sla 11 reachability
R1(config-track) #
```

- d. Specify the level of sensitivity to changes of tracked objects to 10 seconds of down delay and 1 second of up delay using the **delay down 10 up 1** command. The delay helps to alleviate the effect of flapping objects—objects that are going down and up rapidly. In this situation, if the DNS server fails momentarily and comes back up within 10 seconds, there is no impact.

```
R1(config-track)# delay down 10 up 1
R1(config-track)# exit
R1(config) #
```

- e. Configure the floating static route that will be implemented when tracking object 1 is active. To view routing table changes as they happen, first enable the **debug ip routing** command. Next, use the **ip route 0.0.0.0 0.0.0.0 209.165.201.1 2 track 1** command to create a floating static default route via 209.165.201.1 (ISP1). Notice that this command references the tracking object number 1, which in turn references IP SLA operation number 11.

```
R1# debug ip routing
IP routing debugging is on
R1#
*Apr  2 21:26:46.171: RT: NET-RED 0.0.0.0/0

R1# conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)# ip route 0.0.0.0 0.0.0.0 209.165.201.1 2 track 1
R1(config)#
*Apr  2 21:27:02.851: RT: closer admin distance for 0.0.0.0, flushing 1
routes
*Apr  2 21:27:02.851: RT: NET-RED 0.0.0.0/0
*Apr  2 21:27:02.851: RT: add 0.0.0.0/0 via 209.165.201.1, static metric
[2/0]
*Apr  2 21:27:02.851: RT: NET-RED 0.0.0.0/0
*Apr  2 21:27:02.851: RT: default path is now 0.0.0.0 via 209.165.201.1
*Apr  2 21:27:02.855: RT: new default network 0.0.0.0
*Apr  2 21:27:02.855: RT: NET-RED 0.0.0.0/0
*Apr  2 21:27:07.851: RT: NET-RED 0.0.0.0/0
```

Notice that the default route with an administrative distance of 5 has been immediately flushed because of a route with a better admin distance. It then adds the new default route with the admin distance of 2.

- f. Repeat the steps for operation 22, track number 2, and assign the static route an admin distance higher than track 1 and lower than 5. On R1, copy the following configuration, which sets an admin distance of 3.

```
track 2 ip sla 22 reachability
delay down 10 up 1
exit
ip route 0.0.0.0 0.0.0.0 209.165.202.129 3 track 2
```

- g. Verify the routing table again.

```
R1# show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2i - IS-IS,
```

```

        su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2ia -
        IS-IS inter area, * - candidate default, U - per-user static

        o - ODR, P - periodic downloaded static route Gateway of last

resort is 209.165.201.1 to network 0.0.0.0

        209.165.201.0/30 is subnetted, 1 subnets
C      209.165.201.0      is directly connected,      Serial0/0/0
        209.165.202.0/30 is subnetted, 1 subnets
C 209.165.202.128 is directly connected, Serial0/0/1C192.168.1.0/24
is directly connected, FastEthernet0/0  S*  0.0.0.0/0 [2/0] via
209.165.201.1

```

Although a new default route was entered, its administrative distance is not better than 2. Therefore, it does not replace the previously entered default route.

### Step 5: Verify IP SLA operation.

In this step you observe and verify the dynamic operations and routing changes when tracked objects fail. The following summarizes the process:

- Disable the DNS loopback interface on ISP1 (R2).
- Observe the output of the **debug** command on R1.
- Verify the static route entries in the routing table and the IP SLA statistics of R1.
- Re-enable the loopback interface on ISP1 (R2) and again observe the operation of the IP SLA tracking feature.

```

ISP1(config)# interface loopback 1
ISP1(config-if)# shutdown ISP1(config-if)#
*Apr  2 15:53:14.307: %LINK-5-CHANGED: Interface Loopback1, changed state to
administratively down
*Apr  2 15:53:15.307: %LINEPROTO-5-UPDOWN: Line protocol on InterfaceLoopback1,
changed state to down

```

a. Shortly after the loopback interface is administratively down, observe the debug output being generated on R1.

```

R1#
*Apr  2 21:32:33.323: %TRACKING-5-STATE: 1 ip sla 11 reachability Up->Down
*Apr  2 21:32:33.323: RT: del 0.0.0.0 via 209.165.201.1, static metric [2/0]
*Apr  2 21:32:33.323: RT: delete network route to 0.0.0.0
*Apr  2 21:32:33.323: RT: NET-RED 0.0.0.0/0
*Apr  2 21:32:33.323: RT: NET-RED 0.0.0.0/0
*Apr  2 21:32:33.323: RT: add 0.0.0.0/0 via 209.165.202.129, static metric
[3/0]
*Apr  2 21:32:33.323: RT: NET-RED 0.0.0.0/0
*Apr  2 21:32:33.323: RT: default path is now 0.0.0.0 via 209.165.202.129
*Apr  2 21:32:33.323: RT: new default network 0.0.0.0
*Apr  2 21:32:33.327: RT: NET-RED 0.0.0.0/0
*Apr  2 21:32:46.171: RT: NET-RED 0.0.0.0/0

```

route

The tracking state of track 1 changes from up to down. This is the object that tracked reachability for IPSLA object11, with an ICMP echo to the ISP1 DNS server at 209.165.201.30.

R1 then proceeds to delete the default route with the administrative distance of 2 and installs the next highest defaultroute to ISP2 with the administrative distance of 3.

---

- b. Verify the routing table.

```
R1# show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static
      route
      o - ODR, P - periodic downloaded static route
```

**Gateway of last resort is 209.165.202.129 to network 0.0.0.0**

```
209.165.201.0/30 is subnetted, 1 subnets
C 209.165.201.0  is directly connected, Serial0/0/0
  209.165.202.0/30 is subnetted, 1 subnets
C 209.165.202.128  is directly connected, Serial0/0/1C
192.168.1.0/24 is directly connected, FastEthernet0/0  S*
  0.0.0.0/0 [3/0] via 209.165.202.129
```

The new static route has an administrative distance of 3 and is being forwarded to ISP2 as it should.

- c. Verify the IP SLA statistics.

```
R1# show ip sla statistics
IPSLAs Latest Operation Statistics

PSLA operation id: 11
Type of operation: icmp-echo
  Latest RTT: NoConnection/Busy/Timeout
Latest operation start time: *15:36:42.871 UTC Fri Apr 2 2010
Latest operation return code: No connection
Number of successes: 84
Number of failures: 13
Operation time to live: Forever
```

```
IPSLA operation id: 22
Type of operation: icmp-echo
  Latest RTT: 8 milliseconds
Latest operation start time: *15:36:46.335 UTC Fri Apr 2 2010
Latest operation return code: OK
Number of successes: 81
Number of failures: 1
Operation time to live: Forever
```

Notice that the latest return code is **No connection** and there have been 12 failures on IP SLA object 11.

- d. Initiate a trace to the web server from the internal LAN IP address.

```
R1# trace 209.165.200.254 source 192.168.1.1
```

```
Type escape sequence to abort.
Tracing the route to 209.165.200.254
```

```
1 209.165.202.129 8 msec 8 msec *
```

This confirms that traffic is leaving router R1 and being forwarded to the ISP2 router.

- e. To examine the routing behavior when connectivity to the ISP1 DNS is restored, re-enable the DNS

address on ISP1 (R2) by issuing the **no shutdown** command on the loopback 1 interface on ISP2.

```
ISP1(config-if)# no shutdown
*Apr  2 15:56:24.655: %LINK-3-UPDOWN: Interface Loopback1, changed state to
up
*Apr  2 15:56:25.655: %LINEPROTO-5-UPDOWN: Line protocol on Interface
Loopback1, changed state to up
```

Notice the output of the **debug ip routing** command on R1.

```
R1#
*Apr  2 21:35:34.327: %TRACKING-5-STATE: 1 ip sla 11 reachability Down->Up
*Apr  2 21:35:34.327: RT: closer admin distance for 0.0.0.0, flushing 1
routes
*Apr  2 21:35:34.327: RT: NET-RED 0.0.0.0/0
*Apr  2 21:35:34.327: RT: add 0.0.0.0/0 via 209.165.201.1, static metric
[2/0]
*Apr  2 21:35:34.327: RT: NET-RED 0.0.0.0/0
*Apr  2 21:35:34.327: RT: default path is now 0.0.0.0 via 209.165.201.1
*Apr  2 21:35:34.327: RT: new default network 0.0.0.0
*Apr  2 21:35:34.327: RT: NET-RED 0.0.0.0/0
*Apr  2 21:35:39.327: RT: NET-RED 0.0.0.0/0
*Apr  2 21:35:46.171: RT: NET-RED 0.0.0.0/0
```

Now the IP SLA 11 operation transitions back to an up state and reestablishes the default static route to ISP1 with an administrative distance of 2.

- f. Again examine the IP SLA statistics.

```
R1# show ip sla statistics
IPSLAs Latest Operation Statistics

Type of operation: icmp-echo
    Latest RTT: 8 milliseconds
Latest operation start time: *15:40:42.871 UTC Fri Apr 2 2010
Latest operation return code: OK
Number of successes: 88
Number of failures: 35
Operation time to live: Forever
```

```
IPSLA operation id: 22
Type of operation: icmp-echo
    Latest RTT: 16 milliseconds
Latest operation start time: *15:40:46.335 UTC Fri Apr 2 2010
Latest operation return code: OK
Number of successes: 105
Number of failures: 1
Operation time to live: Forever
```

The IP SLA 11 operation is active again, as indicated by the OK return code, and the number of successes is incrementing.

- g. Verify the routing table.

```
R1# show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
```

---

i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2ia - IS-

IS inter area, \* - candidate default, U - per- user static

o - ODR, P - periodic downloaded static route

Gateway of last resort is 209.165.201.1 to network 0.0.0.0

209.165.201.0/30 is subnetted, 1 subnets

C 209.165.201.0 is directly connected, Serial0/0/0

209.165.202.0/30 is subnetted, 1 subnets

C 209.165.202.128 is directly connected, Serial0/0/1C

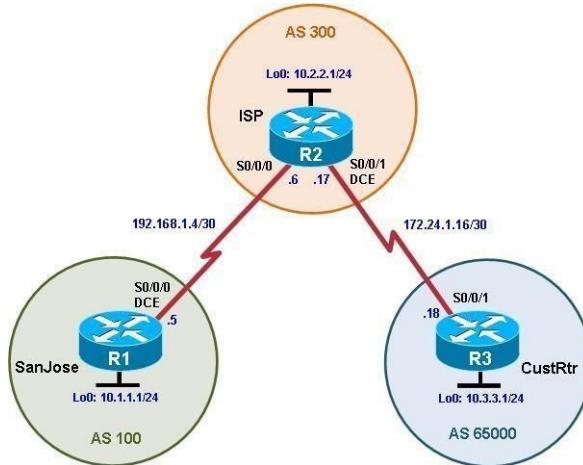
192.168.1.0/24 is directly connected, FastEthernet0/0

S\* 0.0.0.0/0 [2/0] via 209.165.201.

## Practical No. 02

### Aim: Using the AS\_PATH Attribute

#### Topology



#### Objectives

- Use BGP commands to prevent private AS numbers from being advertised to the outside world.
- Use the AS\_PATH attribute to filter BGP routes based on their source AS numbers.

#### Background

The International Travel Agency's ISP has been assigned an AS number of 300. This provider uses BGP to exchange routing information with several customer networks. Each customer network is assigned an AS number from the private range, such as AS 65000. Configure the ISP router to remove the private AS numbers from the AS Path information of CustRtr. In addition, the ISP would like to prevent its customer networks from receiving route information from International Travel Agency's AS 100. Use the AS\_PATH attribute to implement this policy.

**Note:** This lab uses Cisco 1841 routers with Cisco IOS Release 12.4(24)T1 and the Advanced IP Services image c1841-adipservicesk9-mz.124-24.T1.bin. You can use other routers (such as 2801 or 2811) and Cisco IOS Software versions, if they have comparable capabilities and features. Depending on the router model and Cisco IOS Software version, the commands available and output produced might vary from what is shown in this lab.

#### Required Resources

- 3 routers (Cisco 1841 with Cisco IOS Release 12.4(24)T1 Advanced IP Services or comparable)
- Serial and console cables

#### Step 1: Prepare the routers for the lab.

Cable the network as shown in the topology diagram. Erase the startup configuration and reload each router to clear previous configurations.

#### Step 2: Configure the hostname and interface addresses.

- You can copy and paste the following configurations into your routers to begin.

#### **Router R1 (hostname SanJose)**

```
hostname SanJose
!
interface Loopback0
  ip address 10.1.1.1 255.255.255.0
!
interface Serial0/0/0
  ip address 192.168.1.5 255.255.255.252
  clock rate 128000
  no shutdown
```

#### **Router R2 (hostname ISP)**

```
hostname ISP
!
interface Loopback0
  ip address 10.2.2.1 255.255.255.0
!
interface Serial0/0/0
  ip address 192.168.1.6 255.255.255.252
  no shutdown
!
interface Serial0/0/1
  ip address 172.24.1.17 255.255.255.252
  clock rate 128000
  no shutdown
```

#### **Router R3 (hostname CustRtr)**

```
hostname CustRtr
!
interface Loopback0
  ip address 10.3.3.1 255.255.255.0
!
interface Serial0/0/1
  ip address 172.24.1.18 255.255.255.252
  no shutdown
```

- Use **ping** to test the connectivity between the directly connected routers.

**Note:** SanJose will not be able to reach either ISP's loopback (10.2.2.1) or CustRtr's loopback (10.3.3.1), nor will it be able to reach either end of the link joining ISP to CustRtr (172.24.1.17 and 172.24.1.18).

### **Step 3: Configure BGP.**

- Configure BGP for normal operation. Enter the appropriate BGP commands on each router so that they identify their BGP neighbors and advertise their loopback networks.

```
SanJose(config)# router bgp 100
SanJose(config-router)# neighbor 192.168.1.6 remote-as 300
SanJose(config-router)# network 10.1.1.0 mask 255.255.255.0

ISP(config)# router bgp 300
ISP(config-router)# neighbor 192.168.1.5 remote-as 100
ISP(config-router)# neighbor 172.24.1.18 remote-as 65000
ISP(config-router)# network 10.2.2.0 mask 255.255.255.0
```

```
CustRtr(config)# router bgp 65000
CustRtr(config-router)# neighbor 172.24.1.17 remote-as 300
CustRtr(config-router)# network 10.3.3.0 mask 255.255.255.0
```

- b. Verify that these routers have established the appropriate neighbor relationships by issuing the **show ipbgp neighbors** command on each router.

```
ISP# show ip bgp neighbors
BGP neighbor is 172.24.1.18, remote AS 65000, external link
  BGP version 4, remote router ID 10.3.3.1
  BGP state = Established, up for 00:02:05
<output omitted>

BGP neighbor is 192.168.1.5, remote AS 100, external link
  BGP version 4, remote router ID 10.1.1.1
  BGP state = Established, up for 00:04:19
<output omitted>
```

#### Step 4: Remove the private AS.

- a. Display the SanJose routing table using the **show ip route** command. SanJose should have a route to both 10.2.2.0 and 10.3.3.0. Troubleshoot if necessary.

```
SanJose# show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static
route
      o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

      10.0.0.0/24 is subnetted, 3 subnets
B        10.3.3.0 [20/0] via 192.168.1.6, 00:01:11
B        10.2.2.0 [20/0] via 192.168.1.6, 00:02:16
C        10.1.1.0 is directly connected,
      Loopback0[192.168.1.0/30 is subnetted, 1
      subnets

      C        192.168.1.4 is directly connected, Serial0/0/0
```

- b. Ping the 10.3.3.1 address from

SanJose.Why does this fail?

---

- c. Ping again, this time as an extended ping, sourcing from the Loopback0 interface address.

```
SanJose# ping
Protocol [ip]:
Target IP address: 10.3.3.1
Repeat count [5]:
Datagram size [100]:
Timeout in seconds [2]:
Extended commands [n]: y
```

---

Source address or interface: **10.1.1.1**

Type of service [0]:

Set DF bit in IP header? [no]:

Validate reply data? [no]:

Data pattern [0xABCD]:

Loose, Strict, Record, Timestamp, Verbose[none]:

Sweep range of sizes [n]:

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 10.3.3.1, timeout is 2 seconds:

!!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 64/64/68 ms

**Note:** You can bypass extended ping mode and specify a source address using one of these commands:

SanJose# **ping 10.3.3.1 source 10.1.1.1**

or

SanJose# **ping 10.3.3.1 source Lo0**

- Check the BGP table from SanJose by using the **show ip bgp** command. Note the AS path for the 10.3.3.0 network. The AS 65000 should be listed in the path to 10.3.3.0.

SanJose# **show ip bgp**

BGP table version is 5, local router ID is 10.1.1.1

Status codes: s suppressed, d damped, h history, \* valid, > best, i - internal  
Origin codes: i - IGP, e - EGP, ? - incomplete

Network	Next Hop	Metric	LocPrf	Weight	Path
*> 10.1.1.0	0.0.0.0	0		32768	i
*> 10.2.2.0	192.168.1.6	0		0	300 i
<b>*&gt; 10.3.3.0</b>	<b>192.168.1.6</b>			<b>0</b>	<b>300 65000 i</b>

Why is this a problem?

---

- Configure ISP to strip the private AS numbers from BGP routes exchanged with SanJose using the following commands.

ISP(config)# **router bgp 300**

ISP(config-router)# **neighbor 192.168.1.5 remove-private-as**

- After issuing these commands, use the **clear ip bgp \*** command on ISP to reestablish the BGP relationship between the three routers. Wait several seconds and then return to SanJose to check its routing table.

**Note:** The **clear ip bgp \* soft** command can also be used to force each router to resend its BGP table.

Does SanJose still have a route to 10.3.3.0?

---

SanJose should be able to ping 10.3.3.1 using its loopback 0 interface as the source of the ping.

SanJose# **ping 10.3.3.1 source lo0**

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 10.3.3.1, timeout is 2 seconds:

```
Packet sent with a source address of 10.1.1.1
!!!!!
```

Success rate is 100 percent (5/5), round-trip min/avg/max = 28/28/32 ms

- g. Now check the BGP table on SanJose. The AS\_PATH to the 10.3.3.0 network should be AS 300. It no longer has the private AS in the path.

```
SanJose# show ip bgp
```

```
BGP table version is 8, local router ID is 10.1.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i -
internal   Origin codes: i - IGP, e - EGP, ? - incomplete

Network          Next Hop            Metric LocPrf Weight Path
*> 10.1.1.0      0.0.0.0           0        32768 i
*> 10.2.2.0      192.168.1.6       0        0 300 i
*> 10.3.3.0      192.168.1.6       0        0 300 i
```

### Step 5: Use the AS\_PATH attribute to filter routes.

As a final configuration, use the AS\_PATH attribute to filter routes based on their origin. In a complex environment, you can use this attribute to enforce routing policy. In this case, the provider router, ISP, must be configured so that it does not propagate routes that originate from AS 100 to the customer router CustRtr.

AS-path access lists are read like regular access lists. The statements are read sequentially, and there is an implicit deny at the end. Rather than matching an address in each statement like a conventional access list, AS path access lists match on something called a regular expression. Regular expressions are a way of matching text patterns and have many uses. In this case, you will be using them in the AS path access list to match text patterns in AS paths.

- a. Configure a special kind of access list to match BGP routes with an AS\_PATH attribute that both begins and ends with the number 100. Enter the following commands on ISP.

```
ISP(config)# ip as-path access-list 1 deny ^100$  
ISP(config)# ip as-path access-list 1 permit .*
```

The first command uses the ^ character to indicate that the AS path must begin with the given number 100. The \$ character indicates that the AS\_PATH attribute must also end with 100. Essentially, this statement matches only paths that are sourced from AS 100. Other paths, which might include AS 100 along the way, will not match this list.

In the second statement, the . (period) is a wildcard, and the \* (asterisk) stands for a repetition of the wildcard. Together, .\* matches any value of the AS\_PATH attribute, which in effect permits any update that has not been denied by the previous **access-list** statement.

For more details on configuring regular expressions on Cisco routers, see:

[http://www.cisco.com/en/US/docs/ios/12\\_2/termserv/configuration/guide/tcaapre\\_ps1835\\_TSD\\_Products\\_Configuration\\_Guide\\_Chapter.html](http://www.cisco.com/en/US/docs/ios/12_2/termserv/configuration/guide/tcaapre_ps1835_TSD_Products_Configuration_Guide_Chapter.html)

- b. Apply the configured access list using the **neighbor** command with the **filter-list** option.

```
ISP(config)# router bgp 300  
ISP(config-router)# neighbor 172.24.1.18 filter-list 1 out
```

The **out** keyword specifies that the list is applied to routing information sent to this neighbor.

- c. Use the **clear ip bgp \*** command to reset the routing information. Wait several seconds and then check the routing table for ISP. The route to 10.1.1.0 should be in the routing table.

**Note:** To force the local router to resend its BGP table, a less disruptive option is to use the **clear ip bgp**

\* **out** or **clear ip bgp \* soft** command (the second command performs both outgoing and incoming route resync).

```
ISP# show ip route
<output omitted>

  172.24.0.0/30 is subnetted, 1 subnets
C       172.24.1.16 is directly connected,
      Serial0/0/110.0.0.0/24 is subnetted, 3 subnets
B         10.3.3.0 [20/0] via 172.24.1.18, 00:07:34
C       10.2.2.0 is directly connected,
Loopback0[B 10.1.1.0 [20/0] via 192.168.1.5,
00:10:53]

  192.168.1.0/30 is subnetted, 1 subnets
C       192.168.1.4 is directly connected, Serial0/0/0
```

- d. Check the routing table for CustRtr. It should not have a route to 10.1.1.0 in its routing table.

```
CustRtr# show ip route
<output omitted>

  172.24.0.0/30 is subnetted, 1 subnets
C       172.24.1.16 is directly connected,
      Serial0/0/110.0.0.0/24 is subnetted, 2 subnets
C         10.3.3.0 is directly connected,
Loopback0[B 10.2.2.0 [20/0] via 172.24.1.17,
00:11:57
```

- e. Return to ISP and verify that the filter is working as intended. Issue the **show ip bgp regexp ^100\$** command.

```
ISP# show ip bgp regexp ^100$
BGP table version is 4, local router ID is 10.2.2.1
Status codes: s suppressed, d damped, h history, * valid, > best, i -
internal   Origin codes: i - IGP, e - EGP, ? - incomplete
```

Network	Next Hop	Metric	LocPrf	Weight	Path
*> 10.1.1.0	192.168.1.5	0		0	100 i

The output of this command shows all matches for the regular expressions that were used in the access list. The path to 10.1.1.0 matches the access list and is filtered from updates to CustRtr.

- f. Run the following Tcl script on all routers to verify whether there is connectivity. All pings from ISP should be successful. SanJose should not be able to ping the CustRtr loopback 10.3.3.1 or the WAN link

172.24.1.16/30. CustRtr should not be able to ping the SanJose loopback 10.1.1.1 or the WAN link 192.168.1.4/30.

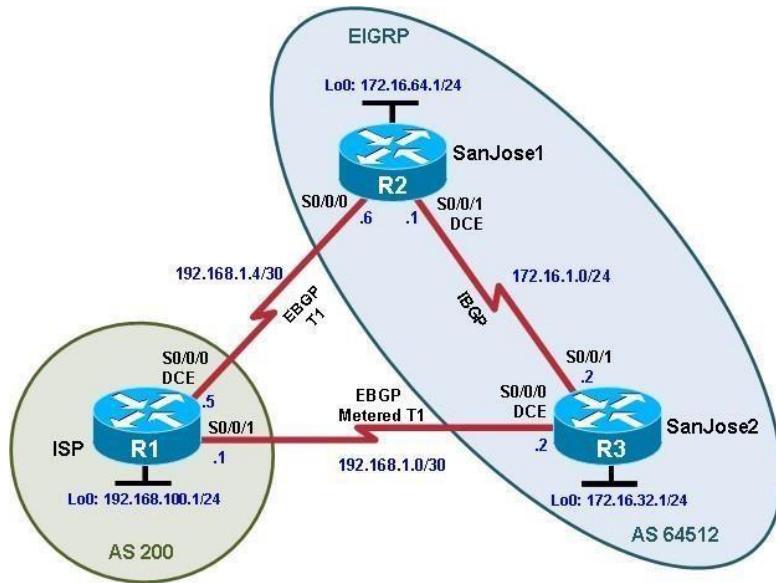
```
ISP# tclsh

foreach address {
10.1.1.1
10.2.2.1
10.3.3.1
192.168.1.5
192.168.1.6
172.24.1.17
172.24.1.18
} {
ping $address }
```

## Practical No. 03

**Aim: Configuring IBGP and EBGP Sessions, Local Preference, and MED**

### Topology



### Objectives

- For IBGP peers to correctly exchange routing information, use the **next-hop-self** command with the **Local-Preference** and **MED** attributes.
- Ensure that the flat-rate, unlimited-use T1 link is used for sending and receiving data to and from the AS 200 on ISP and that the metered T1 only be used in the event that the primary T1 link has failed.

### Background

The International Travel Agency runs BGP on its SanJose1 and SanJose2 routers externally with the ISP router in AS 200. IBGP is run internally between SanJose1 and SanJose2. Your job is to configure both EBGP and IBGP for this internetwork to allow for redundancy. The metered T1 should only be used in the event that the primary T1 link has failed. Traffic sent across the metered T1 link offers the same bandwidth of the primary link but at a huge expense. Ensure that this link is not used unnecessarily.

**Note:** This lab uses Cisco 1941 routers with Cisco IOS Release 15.4 with IP Base. The switches are Cisco WS-C2960-24TT-L with Fast Ethernet interfaces, therefore the router will use routing metrics associated with a 100 Mb/s interface. Depending on the router or switch model and Cisco IOS Software version, the commands available and output produced might vary from what is shown in this lab.

### Required Resources

- 3 routers (Cisco IOS Release 15.2 or comparable)
- Serial and Ethernet cables

### **Step 0: Suggested starting configurations.**

- a. Apply the following configuration to each router along with the appropriate **hostname**. The **exec-timeout 0 0** command should only be used in a lab environment.

```
Router(config)# no ip domain-lookup
Router(config)# line con 0
Router(config-line)# logging synchronous
Router(config-line)# exec-timeout 0 0
```

### **Step 1: Configure interface addresses.**

- a. Using the addressing scheme in the diagram, create the loopback interfaces and apply IPv4 addresses to these and the serial interfaces on ISP (R1), SanJose1 (R2), and SanJose2 (R3).

#### **Router R1 (hostname ISP)**

```
ISP(config)# interface Loopback0
ISP(config-if)# ip address 192.168.100.1 255.255.255.0
ISP(config-if)# exit
ISP(config)# interface Serial0/0/0
ISP(config-if)# ip address 192.168.1.5 255.255.255.252
ISP(config-if)# clock rate 128000
ISP(config-if)# no shutdown
ISP(config-if)# exit
ISP(config)# interface Serial0/0/1
ISP(config-if)# ip address 192.168.1.1 255.255.255.252
ISP(config-if)# no shutdown
ISP(config-if)# end
ISP#
```

#### **Router R2 (hostname SanJose1)**

```
SanJose1(config)# interface Loopback0
SanJose1(config-if)# ip address 172.16.64.1 255.255.255.0
SanJose1(config-if)# exit
SanJose1(config)# interface Serial0/0/0
SanJose1(config-if)# ip address 192.168.1.6 255.255.255.252
SanJose1(config-if)# no shutdown
SanJose1(config-if)# exit
SanJose1(config)# interface Serial0/0/1
SanJose1(config-if)# ip address 172.16.1.1 255.255.255.0
SanJose1(config-if)# clock rate 128000
SanJose1(config-if)# no shutdown
SanJose1(config-if)# end
SanJose1#
```

#### **Router R3 (hostname SanJose2)**

```
SanJose2(config)# interface Loopback0
SanJose2(config-if)# ip address 172.16.32.1 255.255.255.0
SanJose2(config-if)# exit
SanJose2(config)# interface Serial0/0/0
SanJose2(config-if)# ip address 192.168.1.2 255.255.255.252
SanJose2(config-if)# clock rate 128000
SanJose2(config-if)# no shutdown
SanJose2(config-if)# exit
SanJose2(config)# interface Serial0/0/1
SanJose2(config-if)# ip address 172.16.1.2 255.255.255.0
SanJose2(config-if)# no shutdown
SanJose2(config-if)# end
SanJose2#
```

- b. Use **ping** to test the connectivity between the directly connected routers. Both SanJose routers should be able to ping each other and their local ISP serial link IP address. The ISP router cannot reach the segment between SanJose1 and SanJose2.

### **Step 2: Configure EIGRP.**

Configure EIGRP between the SanJose1 and SanJose2 routers. (Note: If using an IOS prior to 15.0, use the **no auto-summary** router configuration command to disable automatic summarization. This command is the default beginning with IOS 15.)

```
SanJose1(config)# router eigrp 1
SanJose1(config-router)# network 172.16.0.0
```

```
SanJose2(config)# router eigrp 1
SanJose2(config-router)# network 172.16.0.0
```

### **Step 3: Configure IBGP and verify BGP neighbors.**

- a. Configure IBGP between the SanJose1 and SanJose2 routers. On the SanJose1 router, enter the following configuration.

```
SanJose1(config)# router bgp 64512
SanJose1(config-router)# neighbor 172.16.32.1 remote-as 64512
SanJose1(config-router)# neighbor 172.16.32.1 update-source lo0
```

If multiple pathways to the BGP neighbor exist, the router can use multiple IP interfaces to communicate with the neighbor. The source IP address therefore depends on the outgoing interface. The **update-source lo0** command instructs the router to use the IP address of the interface Loopback0 as the source IP address for all BGP messages sent to that neighbor.

- b. Complete the IBGP configuration on SanJose2 using the following commands.

```
SanJose2(config)# router bgp 64512
SanJose2(config-router)# neighbor 172.16.64.1 remote-as 64512
SanJose2(config-router)# neighbor 172.16.64.1 update-source lo0
```

- c. Verify that SanJose1 and SanJose2 become BGP neighbors by issuing the **show ip bgp neighbors** command on SanJose1. View the following partial output. If the BGP state is not established, troubleshoot the connection.

```
SanJose2# show ip bgp neighbors
BGP neighbor is 172.16.64.1, remote AS 64512, internal link
  BGP version 4, remote router ID 172.16.64.1
  BGP state = Established, up for 00:00:22
    Last read 00:00:22, last write 00:00:22, hold time is 180, keepalive interval
    is 60 seconds
<output omitted>
```

The link between SanJose1 and SanJose2 should be identified as an internal link indicating an IBGP peering relationship, as shown in the output.

### **Step 4: Configure EBGP and verify BGP neighbors.**

- a. Configure ISP to run EBGP with SanJose1 and SanJose2. Enter the following commands on ISP.

```
ISP(config)# router bgp 200
ISP(config-router)# neighbor 192.168.1.6 remote-as 64512
ISP(config-router)# neighbor 192.168.1.2 remote-as 64512
ISP(config-router)# network 192.168.100.0
```

Because EBGP sessions are almost always established over point-to-point links, there is no reason to use the **update-source** keyword in this configuration. Only one path exists between the peers. If this path goes down, alternative paths are not available.

- Configure a discard static route for the 172.16.0.0/16 network. Any packets that do not have a more specific match (longer match) for a 172.16.0.0 subnet will be dropped instead of sent to the ISP. Later in this lab we will configure a default route to the ISP.

```
SanJose1(config)# ip route 172.16.0.0 255.255.0.0 null0
```

- Configure SanJose1 as an EBGP peer to ISP.

```
SanJose1(config)# router bgp 64512
SanJose1(config-router)# neighbor 192.168.1.5 remote-as 200
SanJose1(config-router)# network 172.16.0.0
```

- Use the **show ip bgp neighbors** command to verify that SanJose1 and ISP have reached the established state. Troubleshoot if necessary.

```
SanJose1# show ip bgp neighbors
BGP neighbor is 172.16.32.1, remote AS 64512, internal link
  BGP version 4, remote router ID 172.16.32.1
  BGP state = Established, up for 00:12:43
<output omitted>

BGP neighbor is 192.168.1.5, remote AS 200, external link
  BGP version 4, remote router ID 192.168.100.1
  BGP state = Established, up for 00:06:49
  Last read 00:00:42, last write 00:00:45, hold time is 180, keepalive interval
  is 60 seconds
<output omitted>
```

Notice that the “external link” indicates that an EBGP peering session has been established. You should also see an informational message indicating the establishment of the BGP neighbor relationship.

```
*Sep  8 21:09:59.699: %BGP-5-ADJCHANGE: neighbor 192.168.1.5 Up
```

- Configure a discard static route for 172.16.0.0/16 on SanJose2 and as an EBGP peer to ISP.

```
SanJose2(config)# ip route 172.16.0.0 255.255.0.0 null0
SanJose2(config)# router bgp 64512
SanJose2(config-router)# neighbor 192.168.1.1 remote-as 200
SanJose2(config-router)# network 172.16.0.0
```

### **Step 5: View BGP summary output.**

In Step 4, the **show ip bgp neighbors** command was used to verify that SanJose1 and ISP had reached the established state. A useful alternative command is **show ip bgp summary**. The output should be similar to the following.

```
SanJose2# show ip bgp summary
BGP router identifier 172.16.32.1, local AS number 64512
BGP table version is 6, main routing table version 6
2 network entries using 288 bytes of memory
4 path entries using 320 bytes of memory
4/2 BGP path/bestpath attribute entries using 640 bytes of memory
1 BGP AS-PATH entries using 24 bytes of memory
0 BGP route-map cache entries using 0 bytes of memory
0 BGP filter-list cache entries using 0 bytes of memory
BGP using 1272 total bytes of memory
BGP activity 2/0 prefixes, 4/0 paths, scan interval 60 secs
```

```

Neighbor          V      AS MsgRcvd MsgSent   TblVer  InQ OutQ Up/Down
State/PfxRcd
172.16.64.1     4       64512    27      26        6      0      0 00:18:15
2
192.168.1.1     4       200      10      7         6      0      0 00:01:42
1
SanJose2#

```

### Step 6: Verify which path the traffic takes.

- f. Clear the IP BGP conversation with the **clear ip bgp \*** command on ISP. Wait for the conversations to reestablish with each SanJose router.

```

ISP# clear ip bgp *
ISP#
*Nov  9 22:05:32.427: %BGP-5-ADJCHANGE: neighbor 192.168.1.2 Down User reset
*Nov  9 22:05:32.427: %BGP_SESSION-5-ADJCHANGE: neighbor 192.168.1.2 IPv4
Unicast topology base removed from session User reset
*Nov  9 22:05:32.427: %BGP-5-ADJCHANGE: neighbor 192.168.1.6 Down User reset
*Nov  9 22:05:32.427: %BGP_SESSION-5-ADJCHANGE: neighbor 192.168.1.6 IPv4
Unicast topology base removed from session User reset
*Nov  9 22:05:32.851: %BGP-5-ADJCHANGE: neighbor 192.168.1.2 Up
*Nov  9 22:05:32.851: %BGP-
ISP#5-ADJCHANGE: neighbor 192.168.1.6 Up
ISP#

```

- g. Test whether ISP can ping the loopback 0 address of 172.16.64.1 on SanJose1 and the serial link between SanJose1 and SanJose2, 172.16.1.1.

```

ISP# ping 172.16.64.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.64.1, timeout is 2 seconds:
.....
Success rate is 0 percent (0/5)
ISP#
ISP# ping 172.16.1.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.1.1, timeout is 2 seconds:
.....
Success rate is 0 percent (0/5)
ISP#

```

- h. Now ping from ISP to the loopback 0 address of 172.16.32.1 on SanJose2 and the serial link between SanJose1 and SanJose2, 172.16.1.2.

```

ISP# ping 172.16.32.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.32.1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 12/14/16 ms
ISP# ping 172.16.1.2
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.1.2, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 12/13/16 ms
ISP#

```

You should see successful pings to each IP address on SanJose2 router. Ping attempts to 172.16.64.1 and 172.16.1.1 should fail. Why does this happen?

---

- i. Issue the **show ip bgp** command on ISP to verify BGP routes and metrics.

```
ISP# show ip bgp
BGP table version is 3, local router ID is 192.168.100.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
*   172.16.0.0        192.168.1.6        0        0 64512 i
*>  172.16.0.0        192.168.1.2        0        0 64512 i
*>  192.168.100.0    0.0.0.0          0        32768 i

ISP#
ISP# show ip bgp
```

Notice that ISP has two valid routes to the 172.16.0.0 network, as indicated by the . However, the link to SanJose2 has been selected as the best path, indicated by the inclusion of the “>”. Why did the ISP prefer the link to SanJose2 over SanJose1?

---



---



---

Would changing the bandwidth metric on each link help to correct this issue? Explain.

---



---



---

BGP operates differently than all other protocols. Unlike other routing protocols that use complex algorithms involving factors such as bandwidth, delay, reliability, and load to formulate a metric, BGP is policy-based. BGP determines the best path based on variables, such as AS path, weight, local preference, MED, and so on. If all things are equal, BGP prefers the route leading to the BGP speaker with the lowest BGP router ID. The SanJose2 router with BGP router ID 172.16.32.1 was preferred to the higher BGP router ID of the SanJose1 router (172.16.64.1).

- j. At this point, the ISP router should be able to get to each network connected to SanJose1 and SanJose2 from the loopback address 192.168.100.1. Use the extended **ping** command and specify the source address of ISP Lo0 to test.

```
ISP# ping 172.16.1.1 source 192.168.100.1
```

```
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.1.1, timeout is 2 seconds:
Packet sent with a source address of 192.168.100.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 20/21/24 ms
```

---



---



---

```
ISP# ping 172.16.32.1 source 192.168.100.1

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.32.1, timeout is 2 seconds:
Packet sent with a source address of 192.168.100.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 12/15/16 ms
```

```
ISP# ping 172.16.1.2 source 192.168.100.1

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.1.2, timeout is 2 seconds:
Packet sent with a source address of 192.168.100.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 12/15/16 ms
ISP#
```

```
ISP# ping 172.16.64.1 source 192.168.100.1

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.64.1, timeout is 2 seconds:
Packet sent with a source address of 192.168.100.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 20/21/24 ms
```

You can also use the extended ping dialogue to specify the source address, as shown in this example.

```
ISP# ping
Protocol [ip]:
Target IP address: 172.16.64.1
Repeat count [5]:
Datagram size [100]:
Timeout in seconds [2]:
Extended commands [n]: y
Source address or interface: 192.168.100.1
Type of service [0]:
Set DF bit in IP header? [no]:
Validate reply data? [no]:
Data pattern [0xABCD]:
Loose, Strict, Record, Timestamp, Verbose[none]:
Sweep range of sizes [n]:
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.64.1, timeout is 2 seconds:
Packet sent with a source address of 192.168.100.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 20/20/24 ms
ISP#
```

Complete reachability has been demonstrated between the ISP router and both SanJose1 and SanJose2.

### **Step 7: Configure the BGP next-hop-self feature.**

SanJose1 is unaware of the link between ISP and SanJose2, and SanJose2 is unaware of the link between ISP and SanJose1. Before ISP can successfully ping all the internal serial interfaces of AS 64512, these serial links should be advertised via BGP on the ISP router. This can also be resolved via EIGRP on each SanJose router. One method is for ISP to advertise these links.

- Issue the following commands on the ISP router.

```
ISP(config)# router bgp 200
ISP(config-router)# network 192.168.1.0 mask 255.255.255.252
ISP(config-router)# network 192.168.1.4 mask 255.255.255.252
```

- b. Issue the **show ip bgp** command to verify that the ISP is correctly injecting its own WAN links into BGP.

```
ISP# show ip bgp
BGP table version is 5, local router ID is 192.168.100.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
*   172.16.0.0        192.168.1.6        0        0 64512 i
*>                    192.168.1.2        0        0 64512 i
*>  192.168.1.0/30    0.0.0.0          0        32768 i
*>  192.168.1.4/30    0.0.0.0          0        32768 i
*>  192.168.100.0     0.0.0.0          0        32768 i
ISP#
```

- c. Verify on SanJose1 and SanJose2 that the opposite WAN link is included in the routing table. The output from SanJose2 is as follows.

```
SanJose2# show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
       a - application route
       + - replicated route, % - next hop override
```

Gateway of last resort is not set

```
172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S   172.16.0.0/16 is directly connected, Null0
C   172.16.1.0/24 is directly connected, Serial0/0/1
L   172.16.1.2/32 is directly connected, Serial0/0/1
C   172.16.32.0/24 is directly connected, Loopback0
L   172.16.32.1/32 is directly connected, Loopback0
D   172.16.64.0/24 [90/2297856] via 172.16.1.1, 00:52:03, Serial0/0/1
    192.168.1.0/24 is variably subnetted, 3 subnets, 2 masks
C   192.168.1.0/30 is directly connected, Serial0/0/0
L   192.168.1.2/32 is directly connected, Serial0/0/0
B   192.168.1.4/30 [20/0] via 192.168.1.1, 00:01:03
B   192.168.100.0/24 [20/0] via 192.168.1.1, 00:25:20
SanJose2#
```

The next issue to consider is BGP policy routing between autonomous systems. The next-hop attribute of a route in a different AS is set to the IP address of the border router in the next AS toward the destination, and this attribute is not modified by default when advertising this route through IBGP. Therefore, for all IBGP peers, it is either necessary to know the route to that border router (in a different neighboring AS), or our own border router needs to advertise the foreign routes using the next-hop-self feature, overriding the next-hop address with its own IP address. The SanJose2 router is passing a policy to SanJose1 and vice versa. The policy for routing from AS 64512 to AS 200 is to forward packets to the 192.168.1.1 interface. SanJose1 has

---

a similar yet opposite policy: it forwards requests to the 192.168.1.5 interface. If either WAN link fails, it is critical that the opposite router become a valid gateway. This is achieved if the **next-hop-self** command is configured on SanJose1 and SanJose2.

- d. To better understand the **next-hop-self** command we will remove ISP advertising its two WAN links and shutdown the WAN link between ISP and SanJose2. The only possible path from SanJose2 to ISP's 192.168.100.0/24 is through SanJose1.

```
ISP(config)# router bgp 200
ISP(config-router)# no network 192.168.1.0 mask 255.255.255.252
ISP(config-router)# no network 192.168.1.4 mask 255.255.255.252
ISP(config-router)# exit
ISP(config)# interface serial 0/0/1
ISP(config-if)# shutdown
ISP(config-if)#
ISP(config)#

```

- e. Display SanJose2's BGP table using the **show ip bgp** command and the IPv4 routing table with **show ip route**.

```
SanJose2# show ip bgp
BGP table version is 1, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
* i 172.16.0.0      172.16.64.1        0       100      0 i
* i 192.168.100.0   192.168.1.5        0       100      0 200 i
SanJose2#

```

```
SanJose2# show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
       a - application route
       + - replicated route, % - next hop override
```

Gateway of last resort is not set

```
172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S   172.16.0.0/16 is directly connected, Null0
C   172.16.1.0/24 is directly connected, Serial0/0/1
L   172.16.1.2/32 is directly connected, Serial0/0/1
C   172.16.32.0/24 is directly connected, Loopback0
L   172.16.32.1/32 is directly connected, Loopback0
D   172.16.64.0/24 [90/2297856] via 172.16.1.1, 02:41:46, Serial0/0/1
SanJose2#

```

Notice that SanJose2 has 192.168.100.0 in its BGP table but not in its routing table. The BGP table shows the next hop to 192.168.100.0 as 192.168.1.5. Because SanJose2 does not have a route to this next hop address of 192.168.1.5 in its routing table, it will not install the 192.168.100.0 network into the routing table. It

won't install a route if it doesn't know how to get to the next hop.

EBGP next hop addresses are carried into IBGP unchanged. As we saw previously, we could advertise the WAN link using BGP, but this is not always desirable. It means advertising additional routes when we are usually trying to minimize the size of the routing table. Another option is to have the routers within the IGP domain advertise themselves as the next hop router using the **next-hop-self** command.

- f. Issue the **next-hop-self** command on SanJose1 and SanJose2 to advertise themselves as the next hop to their IBGP peer.

```
SanJose1(config)# router bgp 64512
SanJose1(config-router)# neighbor 172.16.32.1 next-hop-self
```

```
SanJose2(config)# router bgp 64512
SanJose2(config-router)# neighbor 172.16.64.1 next-hop-self
```

- g. Reset BGP operation on either router with the **clear ip bgp \*** command.

```
SanJose1# clear ip bgp *
SanJose1#
```

```
SanJose2# clear ip bgp *
SanJose2#
```

- h. After the routers have returned to established BGP speakers, issue the **show ip bgp** command on SanJose2 and notice that the next hop is now SanJose1 instead of ISP.

```
SanJose2# show ip bgp
BGP table version is 5, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
*-> 172.16.0.0      0.0.0.0              0        32768  i
* i                 172.16.64.1           0       100      0 i
*>i 192.168.100.0  172.16.64.1           0       100      0 200 i
SanJose2#
```

- i. The **show ip route** command on SanJose2 now displays the 192.168.100.0/24 network because SanJose1 is the next hop, 172.16.64.1, which is reachable from SanJose2.

```
SanJose2# show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
       a - application route
       + - replicated route, % - next hop override
```

Gateway of last resort is not set

172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks

```

S      172.16.0.0/16 is directly connected, Null0
C      172.16.1.0/24 is directly connected, Serial0/0/1
L      172.16.1.2/32 is directly connected, Serial0/0/1
C      172.16.32.0/24 is directly connected, Loopback0
L      172.16.32.1/32 is directly connected, Loopback0
D      172.16.64.0/24 [90/2297856] via 172.16.1.1, 04:27:19, Serial0/0/1
B      192.168.100.0/24 [200/0] via 172.16.64.1, 00:00:46
SanJose2#

```

- j. Before configuring the next BGP attribute, restore the WAN link between ISP and SanJose3. This will change the BGP table and routing table on both routers. For example, SanJose2's routing table shows 192.168.100.0/24 will now have a better path through ISP.

```

ISP(config)# interface serial 0/0/1
ISP(config-if)# no shutdown
ISP(config-if)#

```

```

SanJose2# show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      a - application route
      + - replicated route, % - next hop override

```

Gateway of last resort is not set

```

172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S      172.16.0.0/16 is directly connected, Null0
C      172.16.1.0/24 is directly connected, Serial0/0/1
L      172.16.1.2/32 is directly connected, Serial0/0/1
C      172.16.32.0/24 is directly connected, Loopback0
L      172.16.32.1/32 is directly connected, Loopback0
D      172.16.64.0/24 [90/2297856] via 172.16.1.1, 04:37:34, Serial0/0/1
      192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C      192.168.1.0/30 is directly connected, Serial0/0/0
L      192.168.1.2/32 is directly connected, Serial0/0/0
B      192.168.100.0/24 [20/0] via 192.168.1.1, 00:01:35
SanJose2#

```

### Step 8: Set BGP local preference.

At this point, everything looks good, with the exception of default routes, the outbound flow of data, and inbound packet flow.

- a. Because the local preference value is shared between IBGP neighbors, configure a simple route map that references the local preference value on SanJose1 and SanJose2. This policy adjusts outbound traffic to prefer the link off the SanJose1 router instead of the metered T1 off SanJose2.

```

SanJose1(config)# route-map PRIMARY_T1_IN permit 10
SanJose1(config-route-map)# set local-preference 150
SanJose1(config-route-map)# exit
SanJose1(config)# router bgp 64512
SanJose1(config-router)# neighbor 192.168.1.5 route-map PRIMARY_T1_IN in

```

```
SanJose2(config)# route-map SECONDARY_T1_IN permit 10
SanJose2(config-route-map)# set local-preference 125

SanJose1(config-route-map)# exit
SanJose2(config)# router bgp 64512
SanJose2(config-router)# neighbor 192.168.1.1 route-map SECONDARY_T1_IN in
```

- b. Use the **clear ip bgp \* soft** command after configuring this new policy. When the conversations have been reestablished, issue the **show ip bgp** command on SanJose1 and SanJose2.

```
SanJose1# clear ip bgp * soft
SanJose2# clear ip bgp * soft
```

```
SanJose1# show ip bgp
BGP table version is 3, local router ID is 172.16.64.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
              x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
* i 172.16.0.0      172.16.32.1        0       100      0 i
*>                   0.0.0.0           0           32768 i
*> 192.168.100.0    192.168.1.5        0       150      0 200 i
SanJose1#
```

```
SanJose2# show ip bgp
BGP table version is 7, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
              x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop            Metric LocPrf Weight Path
* i 172.16.0.0      172.16.64.1        0       100      0 i
*>                   0.0.0.0           0           32768 i
*>i 192.168.100.0  172.16.64.1        0       150      0 200 i
*                   192.168.1.1        0       125      0 200 i
SanJose2#
```

This now indicates that routing to the loopback segment for ISP 192.168.100.0 /24 can be reached only through the link common to SanJose1 and ISP. SanJose2's next hop to 192.168.100.0/24 is SanJose1 because both routers have been configured using the **next-hop-self** command.

### Step 9: Set BGP MED.

- a. In the previous step we saw that SanJose1 and SanJose2 will route traffic for 192.168.100.0/24 using the link between SanJose1 and ISP. Examine what the return path ISP takes to reach AS 64512. Notice that the return path is different from the original path. This is known as asymmetric routing and is not necessarily an unwanted trait.

```
ISP# show ip bgp
BGP table version is 22, local router ID is 192.168.100.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
              x best-external, a additional-path, c RIB-compressed,
```

Origin codes: i - IGP, e - EGP, ? - incomplete  
 RPKI validation codes: V valid, I invalid, N Not found

	Network	Next Hop	Metric	LocPrf	Weight	Path
*	172.16.0.0	192.168.1.6	0		0	64512 i
*>		192.168.1.2	0		0	64512 i
*>	192.168.100.0	0.0.0.0		0	32768	i

ISP# **show ip route**

Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP  
 D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area  
 N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2  
 E1 - OSPF external type 1, E2 - OSPF external type 2  
 i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2  
 ia - IS-IS inter area, \* - candidate default, U - per-user static route  
 o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP  
 a - application route  
 + - replicated route, % - next hop override

Gateway of last resort is not set

```
B 172.16.0.0/16 [20/0] via 192.168.1.2, 00:12:45
  192.168.1.0/24 is variably subnetted, 4 subnets, 2 masks
C 192.168.1.0/30 is directly connected, Serial0/0/1
L      192.168.1.1/32 is directly connected, Serial0/0/1
C 192.168.1.4/30 is directly connected, Serial0/0/0
L 192.168.1.5/32 is directly connected, Serial0/0/0
      192.168.100.0/24 is variably subnetted, 2 subnets, 2 masks
C 192.168.100.0/24 is directly connected, Loopback0
L      192.168.100.1/32 is directly connected, Loopback0
```

ISP#

How will traffic from network 192.168.100.0 /24 on ISP return to SanJose1 or SanJose2? Will it be routed through SanJose1 or SanJose2?

---



---



---

To verify this, the simplest solution is to issue the **show ip bgp** command on the ISP router as was done above. What if access was not given to the ISP router? Traffic returning from the Internet should not be passed across the metered T1. Is there a simple way to verify before receiving the monthly bill? How can it be checked instantly?

---



---



---

- a. Use an extended **ping** command to verify this situation. Specify the **record** option and compare your output to the following. Notice the return path using the exit interface 192.168.1.1 to SanJose2.

```
SanJose2# ping
Protocol [ip]:
Target IP address: 192.168.100.1
Repeat count [5]:
```

Datagram size [100]:

Timeout in seconds [2]:

Extended commands [n]: **y**

Source address or interface: **172.16.32.1**

Type of service [0]:

Set DF bit in IP header? [no]:

Validate reply data? [no]:

Data pattern [0xABCD]:

Loose, Strict, Record, Timestamp, Verbose[none]: **record**

Number of hops [ 9 ]:

Loose, Strict, Record, Timestamp, Verbose[RV]:

Sweep range of sizes [n]:

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 192.168.100.1, timeout is 2 seconds:

Packet sent with a source address of 172.16.32.1

Packet has IP options: Total option bytes= 39, padded length=40

Record route: <\*>

(0.0.0.0)

(0.0.0.0)

(0.0.0.0)

(0.0.0.0)

(0.0.0.0)

(0.0.0.0)

(0.0.0.0)

(0.0.0.0)

Reply to request 0 (20 ms). Received packet has options

Total option bytes= 40, padded length=40

Record route:

(172.16.1.2)

(192.168.1.6)

(192.168.100.1)

**(192.168.1.1)**

(172.16.32.1) <\*>

(0.0.0.0)

(0.0.0.0)

(0.0.0.0)

(0.0.0.0)

End of list

Reply to request 1 (20 ms). Received packet has options

Total option bytes= 40, padded length=40

Record route:

(172.16.1.2)

(192.168.1.6)

(192.168.100.1)

**(192.168.1.1)**

(172.16.32.1) <\*>

(0.0.0.0)

(0.0.0.0)

(0.0.0.0)

End of list

Reply to request 2 (20 ms). Received packet has options

Total option bytes= 40, padded length=40

Record route:

```
(172.16.1.2)
(192.168.1.6)

(192.168.100.1)
(192.168.1.1)
(172.16.32.1) <*>
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
End of list
```

Reply to request 3 (24 ms). Received packet has options  
 Total option bytes= 40, padded length=40

```
Record route:
(172.16.1.2)
(192.168.1.6)
(192.168.100.1)
(192.168.1.1)
(172.16.32.1) <*>
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
End of list
```

Reply to request 4 (20 ms). Received packet has options  
 Total option bytes= 40, padded length=40

```
Record route:
(172.16.1.2)
(192.168.1.6)
(192.168.100.1)
(192.168.1.1)
(172.16.32.1) <*>
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
End of list
```

Success rate is 100 percent (5/5), round-trip min/avg/max = 20/20/24 ms  
 SanJose2#

If you are unfamiliar with the **record** option, the important thing to note is that each IP address in brackets is an outgoing interface. The output can be interpreted as follows:

1. A ping that is sourced from 172.16.32.1 exits SanJose2 through s0/0/1, 172.16.1.2. It then arrives at the s0/0/1 interface for SanJose1.
2. SanJose1 S0/0/0, 192.168.1.6, routes the packet out to arrive at the S0/0/0 interface of ISP.
3. The target of 192.168.100.1 is reached: 192.168.100.1.
4. The packet is next forwarded out the S0/0/1, 192.168.1.1 interface for ISP and arrives at the S0/0/0 interface for SanJose2.
5. SanJose2 then forwards the packet out the last interface, loopback 0, 172.16.32.1.

Although the unlimited use of the T1 from SanJose1 is preferred here, ISP currently takes the link from SanJose2 for all return traffic.

- b. ~~Create a new MED route map for the ISP router to the BGP neighbors SanJose1. Create a second route map~~

```
SanJose1(config)#route-map PRIMARY_T1_MED_OUT permit 10
SanJose1(config-route-map)#set Metric 50
SanJose1(config-route-map)#exit
SanJose1(config)#router bgp 64512
SanJose1(config-router)#neighbor 192.168.1.5 route-map PRIMARY_T1_MED_OUT out

SanJose2(config)#route-map SECONDARY_T1_MED_OUT permit 10
SanJose2(config-route-map)#set Metric 75
SanJose2(config-route-map)#exit
SanJose2(config)#router bgp 64512
SanJose2(config-router)#neighbor 192.168.1.1 route-map SECONDARY_T1_MED_OUT out
```

- c. Use the **clear ip bgp \* soft** command after issuing this new policy. Issuing the **show ip bgp** command as follows on SanJose1 or SanJose2 does not indicate anything about this newly defined policy.

```
SanJose1# clear ip bgp * soft
SanJose2# clear ip bgp * soft
```

```
SanJose1# show ip bgp
BGP table version is 4, local router ID is 172.16.64.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found
```

Network	Next Hop	Metric	LocPrf	Weight	Path
* i 172.16.0.0	172.16.32.1	0	100	0	i
*>	0.0.0.0	0		32768	i
*> 192.168.100.0	192.168.1.5	0	150	0	200 i

```
SanJose1#
```

```
SanJose2# show ip bgp
BGP table version is 8, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found
```

Network	Next Hop	Metric	LocPrf	Weight	Path
* i 172.16.0.0	172.16.64.1	0	100	0	i
*>	0.0.0.0	0		32768	i
*>i 192.168.100.0	172.16.64.1	0	150	0	200 i
*	192.168.1.1	0	125	0	200 i

```
SanJose2#
```

- d. Reissue an extended **ping** command with the **record** command. Notice the change in return path using the exit interface 192.168.1.5 to SanJose1.

```
SanJose2# ping
Protocol [ip]:
Target IP address: 192.168.100.1
Repeat count [5]:
Datagram size [100]:
Timeout in seconds [2]:
Extended commands [n]: y
Source address or interface: 172.16.32.1
Type of service [0]:
```

```

Set DF bit in IP header? [no]: 
Validate reply data? [no]: 
Data pattern [0xABCD]: 
Loose, Strict, Record, Timestamp, Verbose[none]: record
Number of hops [ 9 ]: 
Loose, Strict, Record, Timestamp, Verbose[RV]: 
Sweep range of sizes [n]: 
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.100.1, timeout is 2 seconds:
Packet sent with a source address of 172.16.32.1
Packet has IP options: Total option bytes= 39, padded length=40
Record route: <*>
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)

Reply to request 0 (28 ms). Received packet has options
Total option bytes= 40, padded length=40
Record route:
(172.16.1.2)
(192.168.1.6)
(192.168.100.1)
(192.168.1.5)
(172.16.1.1)
(172.16.32.1) <*>
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
End of list

Reply to request 1 (28 ms). Received packet has options
Total option bytes= 40, padded length=40
Record route:
(172.16.1.2)
(192.168.1.6)
(192.168.100.1)
(192.168.1.5)
(172.16.1.1)
(172.16.32.1) <*>
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
End of list

Reply to request 2 (28 ms). Received packet has options
Total option bytes= 40, padded length=40
Record route:
(172.16.1.2)
(192.168.1.6)
(192.168.100.1)
(192.168.1.5)
(172.16.1.1)

```

```
(172.16.32.1) <*>
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
End of list

Reply to request 3 (28 ms). Received packet has options
Total option bytes= 40, padded length=40
Record route:
(172.16.1.2)
(192.168.1.6)
(192.168.100.1)
(192.168.1.5)
(172.16.1.1)
(172.16.32.1) <*>
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
End of list

Reply to request 4 (28 ms). Received packet has options
Total option bytes= 40, padded length=40
Record route:
(172.16.1.2)
(192.168.1.6)
(192.168.100.1)
(192.168.1.5)
(172.16.1.1)
(172.16.32.1) <*>
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
End of list

Success rate is 100 percent (5/5), round-trip min/avg/max = 28/28/28 ms
SanJose2#
```

Does the output look correct? Does the 192.168.1.5 above mean that the ISP now prefers SanJose1 for return traffic

The newly configured policy MED shows that the lower MED value is considered best. The ISP now prefers the route with the lower MED value of 50 to AS 64512. This is just opposite from the **local-preference** command configured earlier.

```
ISP# show ip bgp
BGP table version is 24, local router ID is 192.168.100.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

      Network          Next Hop           Metric LocPrf Weight Path
*>  172.16.0.0      192.168.1.6        50      0 64512 i
*            192.168.1.2        75      0 64512 i
*>  192.168.100.0    0.0.0.0          0      32768 i
```

ISP#

**Step 10: Establish a default route.**

The final step is to establish a default route that uses a policy statement that adjusts to changes in the network.

- Configure ISP to inject a default route to both SanJose1 and SanJose2 using BGP using the **default-originate** command. This command does not require the presence of 0.0.0.0 in the ISP router. Configure the 10.0.0.0/8 network which will not be advertised using BGP. This network will be used to test the default route on SanJose1 and SanJose2.

```
ISP(config)# router bgp 200
ISP(config-router)# neighbor 192.168.1.6 default-originate
ISP(config-router)# neighbor 192.168.1.2 default-originate
ISP(config-router)# exit
ISP(config)# interface loopback 10
ISP(config-if)# ip address 10.0.0.1 255.255.255.0
ISP(config-if)#

```

- Verify that both routers have received the default route by examining the routing tables on SanJose1 and SanJose2. Notice that both routers prefer the route between SanJose1 and ISP.

```
SanJose1# show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      a - application route
      + - replicated route, % - next hop override
```

Gateway of last resort is 192.168.1.5 to network 0.0.0.0

```
B* 0.0.0.0/0 [20/0] via 192.168.1.5, 00:00:36
  172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S    172.16.0.0/16 is directly connected, Null0
C    172.16.1.0/24 is directly connected, Serial0/0/1
L    172.16.1.1/32 is directly connected, Serial0/0/1
D    172.16.32.0/24 [90/2297856] via 172.16.1.2, 05:47:24, Serial0/0/1
C    172.16.64.0/24 is directly connected, Loopback0
L    172.16.64.1/32 is directly connected, Loopback0
  192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C    192.168.1.4/30 is directly connected, Serial0/0/0
L    192.168.1.6/32 is directly connected, Serial0/0/0
SanJose1#
```

```
SanJose2# show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
      E1 - OSPF external type 1, E2 - OSPF external type 2
      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
      ia - IS-IS inter area, * - candidate default, U - per-user static route
      o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP
      a - application route
      + - replicated route, % - next hop override
```

Gateway of last resort is 172.16.64.1 to network 0.0.0.0

```
B*   0.0.0.0/0 [200/0] via 172.16.64.1, 00:00:45
    172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S       172.16.0.0/16 is directly connected, Null0
C       172.16.1.0/24 is directly connected, Serial0/0/1
L       172.16.1.2/32 is directly connected, Serial0/0/1
C       172.16.32.0/24 is directly connected, Loopback0
L       172.16.32.1/32 is directly connected, Loopback0
D       172.16.64.0/24 [90/2297856] via 172.16.1.1, 05:47:33, Serial0/0/1
    192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C       192.168.1.0/30 is directly connected, Serial0/0/0
L       192.168.1.2/32 is directly connected, Serial0/0/0
SanJose2#
```

- c. The preferred default route is by way of SanJose1 because of the higher local preference attribute configured on SanJose1 earlier.

```
SanJose2# show ip bgp
BGP table version is 38, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
               x best-external, a additional-path, c RIB-compressed,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found
```

Network	Next Hop	Metric	LocPrf	Weight	Path
*>i 0.0.0.0	172.16.64.1	0	150	0	200 i
*	192.168.1.1		125	0	200 i
* i 172.16.0.0	172.16.64.1	0	100	0	i
*>	0.0.0.0	0		32768	i
*>i 192.168.100.0	172.16.64.1	0	150	0	200 i
*	192.168.1.1	0	125	0	200 i

SanJose2#

- d. Using the traceroute command verify that packets to 10.0.0.1 is using the default route through SanJose1.

```
SanJose2# traceroute 10.0.0.1
Type escape sequence to abort.
Tracing the route to 10.0.0.1
VRF info: (vrf in name/id, vrf out name/id)
  1 172.16.1.1 8 msec 4 msec 8 msec
  2 192.168.1.5 [AS 200] 12 msec * 12 msec
SanJose2#
```

- e. Next, test how BGP adapts to using a different default route when the path between SanJose1 and ISP goes down.

```
ISP(config)# interface serial 0/0/0
ISP(config-if)# shutdown
ISP(config-if)#
```

- f. Verify that both routers are modified their routing tables with the default route using the path between SanJose2 and ISP.

```
SanJose1# show ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
```

N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2  
 E1 - OSPF external type 1, E2 - OSPF external type 2  
 i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2  
 ia - IS-IS inter area, \* - candidate default, U - per-user static route  
 o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP  
 a - application route  
 + - replicated route, % - next hop override

Gateway of last resort is 172.16.32.1 to network 0.0.0.0

```
B* 0.0.0.0/0 [200/0] via 172.16.32.1, 00:00:06
  172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S    172.16.0.0/16 is directly connected, Null0
C    172.16.1.0/24 is directly connected, Serial0/0/1
L    172.16.1.1/32 is directly connected, Serial0/0/1
D    172.16.32.0/24 [90/2297856] via 172.16.1.2, 05:49:25, Serial0/0/1
C    172.16.64.0/24 is directly connected, Loopback0
L    172.16.64.1/32 is directly connected, Loopback0
B  192.168.100.0/24 [200/0] via 172.16.32.1, 00:00:06
SanJose1#
```

SanJose2# **show ip route**

Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP  
 D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area  
 N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2  
 E1 - OSPF external type 1, E2 - OSPF external type 2  
 i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2  
 ia - IS-IS inter area, \* - candidate default, U - per-user static route  
 o - ODR, P - periodic downloaded static route, H - NHRP, l - LISP  
 a - application route  
 + - replicated route, % - next hop override

Gateway of last resort is 192.168.1.1 to network 0.0.0.0

```
B* 0.0.0.0/0 [20/0] via 192.168.1.1, 00:00:30
  172.16.0.0/16 is variably subnetted, 6 subnets, 3 masks
S    172.16.0.0/16 is directly connected, Null0
C    172.16.1.0/24 is directly connected, Serial0/0/1
L    172.16.1.2/32 is directly connected, Serial0/0/1
C    172.16.32.0/24 is directly connected, Loopback0
L    172.16.32.1/32 is directly connected, Loopback0
D    172.16.64.0/24 [90/2297856] via 172.16.1.1, 05:49:49, Serial0/0/1
  192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C    192.168.1.0/30 is directly connected, Serial0/0/0
L    192.168.1.2/32 is directly connected, Serial0/0/0
B  192.168.100.0/24 [20/0] via 192.168.1.1, 00:00:30
SanJose2#
```

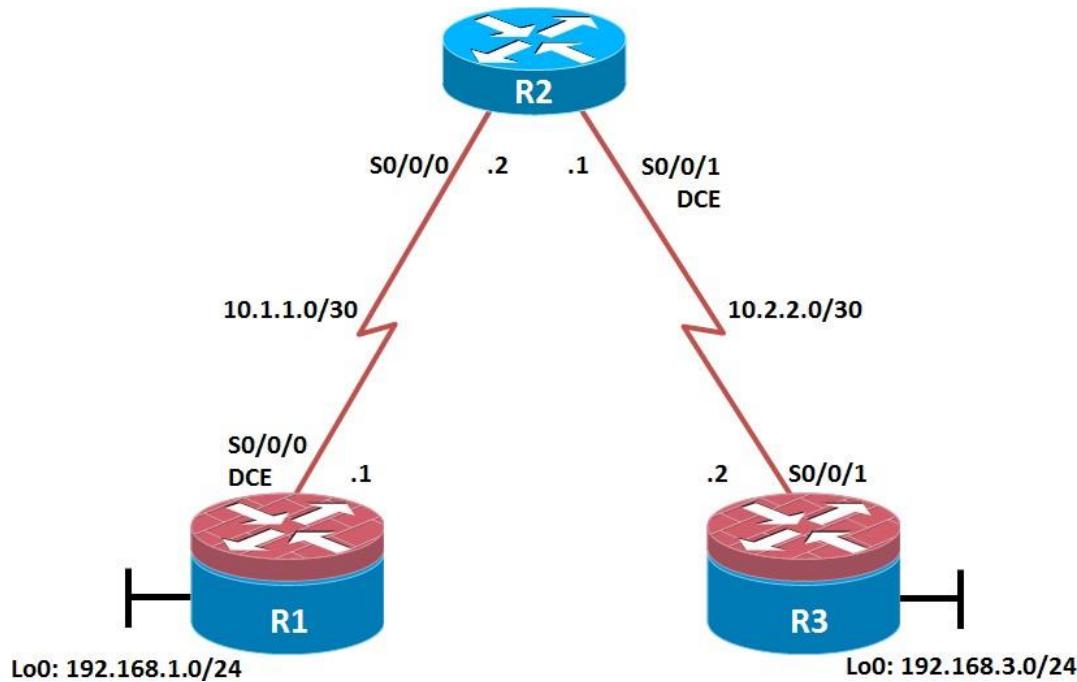
- g. Verify the new path using the traceroute command to 10.0.0.1 from SanJose1. Notice the default route is now through SanJose2.

```
SanJose1# trace 10.0.0.1
Type escape sequence to abort.
Tracing the route to 10.0.0.1
VRF info: (vrf in name/id, vrf out name/id)
  1 172.16.1.2 8 msec 8 msec
  2 192.168.1.1 [AS 200] 12 msec * 12 msec
SanJose1#
```

## Practical No.04

**Aim:** Secure the Management Plane

### Topology



### Objectives

- Secure management access.
- Configure enhanced username password security.
- Enable AAA RADIUS authentication.
- Enable secure remote management.

### Background

The management plane of any infrastructure device should be protected as much as possible. Controlling access to routers and enabling reporting on routers are critical to network security and should be part of a comprehensive security policy.

In this lab, you build a multi-router network and secure the management plane of routers R1 and R3.

**Note:** This lab uses Cisco 1941 routers with Cisco IOS Release 15.2 with IP Base. Depending on the router or switch model and Cisco IOS Software version, the commands available and output produced might vary from what is shown in this lab.

## Required Resources

- 3 routers (Cisco IOS Release 15.2 or comparable)
- Serial and Ethernet cables

### Step 1: Configure loopbacks and assign addresses.

Cable the network as shown in the topology diagram. Erase the startup configuration and reload each router to clear previous configurations. Using the addressing scheme in the diagram, apply the IP addresses to the interfaces on the R1, R2, and R3 routers.

You can copy and paste the following configurations into your routers to begin.

**Note:** Depending on the router model, interfaces might be numbered differently than those listed. You might need to alter the designations accordingly.

#### R1

```
hostname R1

interface Loopback 0
description R1 LAN
ip address 192.168.1.1 255.255.255.0
exit
!
interface Serial0/0/0
description R1 --> R2
ip address 10.1.1.1 255.255.255.252
clock rate 128000
no shutdown
exit
!
end
```

#### R2

```
hostname R2
!
interface Serial0/0/0
description R2 --> R1
ip address 10.1.1.2 255.255.255.252
no shutdown
exit

interface Serial0/0/1
description R2 --> R3
ip address 10.2.2.1 255.255.255.252
clock rate 128000
no shutdown
exit
!
end
```

**R3**

```

hostname R3
!
interface Loopback0
  description R3 LAN
  ip address 192.168.3.1 255.255.255.0
exit

interface Serial0/0/1
  description R3 --> R2
  ip address 10.2.2.2 255.255.255.252
  no shutdown
exit
!
end

```

**Step 2: Configure static routes.**

- a. On R1, configure a default static route to ISP.

i. R1(config)# **ip route 0.0.0.0 0.0.0.0 10.1.1.2**

- b. On R3, configure a default static route to ISP.

i. R3(config)# **ip route 0.0.0.0 0.0.0.0 10.2.2.1**

- c. On R2, configure two static routes.

i. R2(config)# **ip route 192.168.1.0 255.255.255.0 10.1.1.1**

ii. R2(config)# **ip route 192.168.3.0 255.255.255.0 10.2.2.2**

- d. From the R1 router, run the following Tcl script to verify connectivity.

```

foreach address {
192.168.1.1
10.1.1.1
10.1.1.2
10.2.2.1
10.2.2.2
192.168.3.1
} { ping $address }

```

R1# **tclsh**

```

R1(tcl)#foreach address {
+>(tcl)#192.168.1.1
+>(tcl)#10.1.1.1
+>(tcl)#10.1.1.2
+>(tcl)#10.2.2.1
+>(tcl)#10.2.2.2
+>(tcl)#192.168.3.1
+>(tcl)#{ ping $address }

```

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 192.168.1.1, timeout is 2 seconds:  
!!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/1 ms  
Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 10.1.1.1, timeout is 2 seconds:  
!!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 1/2/4 ms  
Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 10.1.1.2, timeout is 2 seconds:  
!!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/4 ms  
Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 10.2.2.1, timeout is 2 seconds:  
!!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/4 ms  
Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 10.2.2.2, timeout is 2 seconds:  
!!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 12/14/16 ms  
Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 192.168.3.1, timeout is 2 seconds:  
!!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 12/15/16 ms  
R1(tcl) #

Are the pings now successful?

---

Yes. If not, troubleshoot.

---

### Step 3: Secure management access.

- On R1, use the **security passwords** command to set a minimum password length of 10 characters.

R1(config)# **security passwords min-length 10**

- Configure the enable secret encrypted password on both routers.

R1(config)# **enable secret class12345**

How does configuring an enable secret password help protect a router from being compromised by an attack?

---

–

---

–

---

–

The goal is to always prevent unauthorized users from accessing a device using Telnet, SSH, or via the console. If attackers are able to penetrate this first layer of defense, using an enable secret password prevents them from being able to alter the configuration of the device. Unless the enable secret password is known, a user cannot go into privileged EXEC mode where they can display the running config and enter various configuration commands to make changes to the router. This provides an additional layer of security.

**Note:** Passwords in this task are set to a minimum of 10 characters but are relatively simple for the benefit of performing the lab. More complex passwords are recommended in a production network.

- c. Configure a console password and enable login for routers. For additional security, the **exec-timeout** command causes the line to log out after 5 minutes of inactivity. The **logging synchronous** command prevents console messages from interrupting command entry.

**Note:** To avoid repetitive logins during this lab, the **exec-timeout** command can be set to 0 0, which prevents it from expiring. However, this is not considered a good security practice.

```
R1(config)# line console 0
R1(config-line)# password ciscoconpass
R1(config-line)# exec-timeout 5 0
R1(config-line)# login
R1(config-line)# logging synchronous
R1(config-line)# exit
R1(config) #
```

- d. Configure the password on the vty lines for router R1.

```
R1(config)# line vty 0 4
R1(config-line)# password ciscovtypass
R1(config-line)# exec-timeout 5 0
R1(config-line)# login
R1(config-line)# exit
R1(config) #
```

- e. The aux port is a legacy port used to manage a router remotely using a modem and is hardly ever used. Therefore, disable the aux port.

```
R1(config)# line aux 0
R1(config-line)# no exec
R1(config-line)# end
R1#
```

- f. Enter privileged EXEC mode and issue the **show run** command. Can you read the enable secret password? Why or why not?

No. The enable secret password is encrypted automatically using the MD5 or SHA hash algorithm. . IOS 15.0(1)S and later default to SHA256 hashing algorithm. SHA256 which is considered to be a very strong hashing algorithm and is extremely difficult to reverse. Earlier IOS versions use the weaker MD5 hashing algorithm.

**Note:** If the **enable secret** password command is lost or forgotten, it must be replaced using the Cisco router password recovery procedure. Refer to [cisco.com](http://cisco.com) for more information.

Can you read the console, aux, and vty passwords? Why or why not?

---



---

Yes. They are all in clear text.

- g. Use the **service password-encryption** command to encrypt the line console and vty passwords.

```
R1(config)# service password-encryption
R1(config) #
```

**Note:** Password encryption is applied to all the passwords, including the **username** passwords, the authentication key passwords, the privileged command password, the console and the virtual terminal line access passwords, and the BGP neighbor passwords.

- h. Issue the **show run** command. Can you read the console, aux, and vty passwords? Why or why not?
- 

No. The passwords are now encrypted.

**Note:** Type 7 passwords are encrypted using a Vigenère cipher which can be easily reversed. Therefore this command primarily protects from shoulder surfing attacks.

- i. Configure a warning to unauthorized users with a message-of-the-day (MOTD) banner using the **banner motd** command. When a user connects to one of the routers, the MOTD banner appears before the login prompt. In this example, the dollar sign (\$) is used to start and end the message.

```
R1(config)# banner motd $Unauthorized access strictly prohibited!$  
R1(config)# exit
```

- j. Issue the **show run** command. What does the \$ convert to in the output?
- 

The \$ is converted to ^C when the running-config is displayed.

- k. Exit privileged EXEC mode using the **disable** or **exit** command and press **Enter** to get started. Does the MOTD banner look like what you created with the **banner motd** command? If the MOTD banner is not as you wanted it, recreate it using the **banner motd** command.

h.

- l. Repeat the configuration portion of steps 3a through 3k on router R3.

#### Step 4: Configure enhanced username password security.

To increase the encryption level of console and VTY lines, it is recommended to enable authentication using the local database. The local database consists of usernames and password combinations that are created locally on each device. The local and VTY lines are configured to refer to the local database when authenticating a user.

- a. To create local database entry encrypted to level 4 (SHA256), use the **username name secret password** global configuration command. In global configuration mode, enter the following command:

```
R1(config)# username JR-ADMIN secret class12345
R1(config)# username ADMIN secret class54321
```

**Note:** An older method for creating local database entries is to use the **username name password**

*password* command.

- b. Set the console line to use the locally defined login accounts.

```
R1(config)# line console 0
R1(config-line)# login local
R1(config-line)# exit
R1(config) #
```

- c. Set the vty lines to use the locally defined login accounts.

```
R1(config)# line vty 0 4
R1(config-line)# login local
R1(config-line)# end
R1(config) #
```

- d. Repeat the steps 4a to 4c on R3.

i.

- e. To verify the configuration, telnet to R3 from R1 and login using the ADMIN local database account.

```
R1# telnet 10.2.2.2
Trying 10.2.2.2 ... Open
Unauthorized access strictly prohibited!
User Access Verification

Username: ADMIN
Password:
R3>
```

## Step 5: Enabling AAA RADIUS Authentication with Local User for Backup.

Authentication, authorization, and accounting (AAA) is a standards-based framework that can be implemented to control who is permitted to access a network (authenticate), what they can do on that network (authorize), and audit what they did while accessing the network (accounting).

Users must authenticate against an authentication database which can be stored:

- **Locally:** Users are authenticated against the local device database which is created using the `username secret` command. Sometimes referred to self-contained AAA.
- **Centrally:** A client-server model where users are authenticated against AAA servers. This provides improved scalability, manageability and control. Communication between the device and AAA servers is secured using either the RADIUS or TACACS+ protocols.

In this step, we will configure AAA authentication to use a RADIUS server and the local database as a backup. Specifically, the authentication will be validated against one of two RADIUS servers. If the servers are not available, then authentication will be validated against the local database.

- a. Always have local database accounts created before enabling AAA. Since we created two local database accounts in the previous step, then we can proceed and enable AAA on R1.

```
R1(config)# aaa new-model
```

**Note:** Although the following configuration refers to two RADIUS servers, the actual RADIUS server implementation is beyond the scope. Therefore, the goal of this step is to provide an example of how to configure a router to access the servers.

- b. Configure the specifics for the first RADIUS server located at 192.168.1.101. Use **RADIUS-1-pa55w0rd** as the server password.

```
R1(config)# radius server RADIUS-1
R1(config-radius-server)# address ipv4 192.168.1.101
R1(config-radius-server)# key RADIUS-1-pa55w0rd
R1(config-radius-server)# exit
R1(config) #
```

- c. Configure the specifics for the second RADIUS server located at 192.168.1.102. Use **RADIUS-2-pa55w0rd** as the server password.

```
R1(config)# radius server RADIUS-2
R1(config-radius-server)# address ipv4 192.168.1.102
R1(config-radius-server)# key RADIUS-2-pa55w0rd
R1(config-radius-server)# exit
R1(config) #
```

- d. Assign both RADIUS servers to a server group.

```
R1(config)# aaa group server radius RADIUS-GROUP
R1(config-sg-radius)# server name RADIUS-1
R1(config-sg-radius)# server name RADIUS-2
R1(config-sg-radius)# exit
R1(config) #
```

- e. Enable the default AAA authentication login to attempt to validate against the server group. If they are not available, then authentication should be validated against the local database..

```
R1(config)# aaa authentication login default group RADIUS-GROUP local
R1(config) #
```

**Note:** Once this command is configured, all line access methods default to the default authentication method. The **local** option enables AAA to refer to the local database. Only the password is case sensitive.

- f. Enable the default AAA authentication Telnet login to attempt to validate against the server group. If they are not available, then authentication should be validated against a case sensitive local database.

```
R1(config)# aaa authentication login TELNET-LOGIN group RADIUS-GROUP local-case
R1(config) #
```

**Note:** Unlike the **local** option that makes the password is case sensitive, local-case makes the username and password case sensitive.

- g. Alter the VTY lines to use the TELNET-LISTEN AAA authentication method.

```
R1(config)# line vty 0 4
R1(config-line)# login authentication TELNET-LISTEN
R1(config-line)# exit
R1(config) #
```

- h. Repeat the steps 5a to 5g on R3.

j.

- i. To verify the configuration, telnet to R3 from R1 and login using the ADMIN local database account.

```
R1# telnet 10.2.2.2
Trying 10.2.2.2 ... Open
Unauthorized access strictly prohibited!
```

User Access Verification

Username: **admin**

Password:

% Authentication failed

Username: **ADMIN**

Password:

R3>

**Note:** The first login attempt did not use the correct username (i.e., ADMIN) which is why it failed.

**Note:** The actual login time is longer since the RADIUS servers are not available.

## Step 6: Enabling secure remote management using SSH.

Traditionally, remote access on routers was configured using Telnet on TCP port 23. However, Telnet was developed in the days when security was not an issue; therefore, all Telnet traffic is forwarded in plaintext.

Secure Shell (SSH) is a network protocol that establishes a secure terminal emulation connection to a router or other networking device. SSH encrypts all information that passes over the network link and provides authentication of the remote computer. SSH is rapidly replacing Telnet as the remote login tool of choice for network professionals.

**Note:** For a router to support SSH, it must be configured with local authentication, (AAA services, or username) or password authentication. In this task, you configure an SSH username and local authentication.

In this step, you will enable R1 and R3 to support SSH instead of Telnet.

- a. SSH requires that a device name and a domain name be configured. Since the router already has a name assigned, configure the domain name.

```
R1(config)# ip domain-name ccnasecurity.com
```

- b. The router uses the RSA key pair for authentication and encryption of transmitted SSH data. Although optional it may be wise to erase any existing key pairs on the router.

```
R1(config)# crypto key zeroize rsa
```

**Note:** If no keys exist, you might receive this message: % No Signature RSA Keys found in configuration.

- c. Generate the RSA encryption key pair for the router. Configure the RSA keys with **1024** for the number of modulus bits. The default is 512, and the range is from 360 to 2048.

```
R1(config)# crypto key generate rsa general-keys modulus 1024
```

The name for the keys will be: R1.ccnasecurity.com

```
% The key modulus size is 1024 bits
```

```
% Generating 1024 bit RSA keys, keys will be non-exportable... [OK]
```

```
R1(config) #
```

```
Jan 10 13:44:44.711: %SSH-5-ENABLED: SSH 1.99 has been enabled
```

```
R1(config) #
```

k.

- d. Cisco routers support two versions of SSH:

- **SSH version 1 (SSHv1):** Original version but has known vulnerabilities.
- **SSH version 2 (SSHv2):** Provides better security using the Diffie-Hellman key exchange and the strong integrity-checking message authentication code (MAC).

- i. The default setting for SSH is SSH version 1.99. This is also known as compatibility mode and is merely an indication that the server supports both SSH version 2 and SSH version 1. However, best practices are to enable version 2 only.

- m. Configure SSH version 2 on R1.

```
R1(config)# ip ssh version 2
```

```
R1(config) #
```

n.

- e. Configure the vty lines to use only SSH connections.

```
R1(config)# line vty 0 4
```

```
R1(config-line)# transport input ssh
```

```
R1(config-line)# end
```

**Note:** SSH requires that the **login local** command be configured. However, in the previous step we enabled AAA authentication using the TELNET-LOGIN authentication method, therefore **login local** is not necessary.

**Note:** If you add the keyword **telnet** to the **transport input** command, users can log in using Telnet as well as SSH. However, the router will be less secure. If only SSH is specified, the connecting host must have an SSH client installed.

- f. Verify the SSH configuration using the **show ip ssh** command.

```
R1# show ip ssh
SSH Enabled - version 2.0
Authentication timeout: 120 secs; Authentication retries: 3
Minimum expected Diffie Hellman key size : 1024 bits
IOS Keys in SECSH format(ssh-rsa, base64 encoded):
ssh-rsa AAAAB3NzaC1yc2EAAAQABAAAAgQC3Lehh7ReYlgyDzls6wq+mFzxqzoaZFr9XGx+Q/yio
dFYw00hQo80tZy1W1FF3Pz6q7Qi0y00urwddHZ0kBZceZK9EZJ6wZ+9a87KKDETCWrGSLi6c81E/y4K+
Z/oVrMMZk7bpTM1MFdP41YgkTf35utYv+TcqbsYo++KJiYk+xw==
R1#
```

- g. Repeat the steps 6a to 6f on R3.

o.

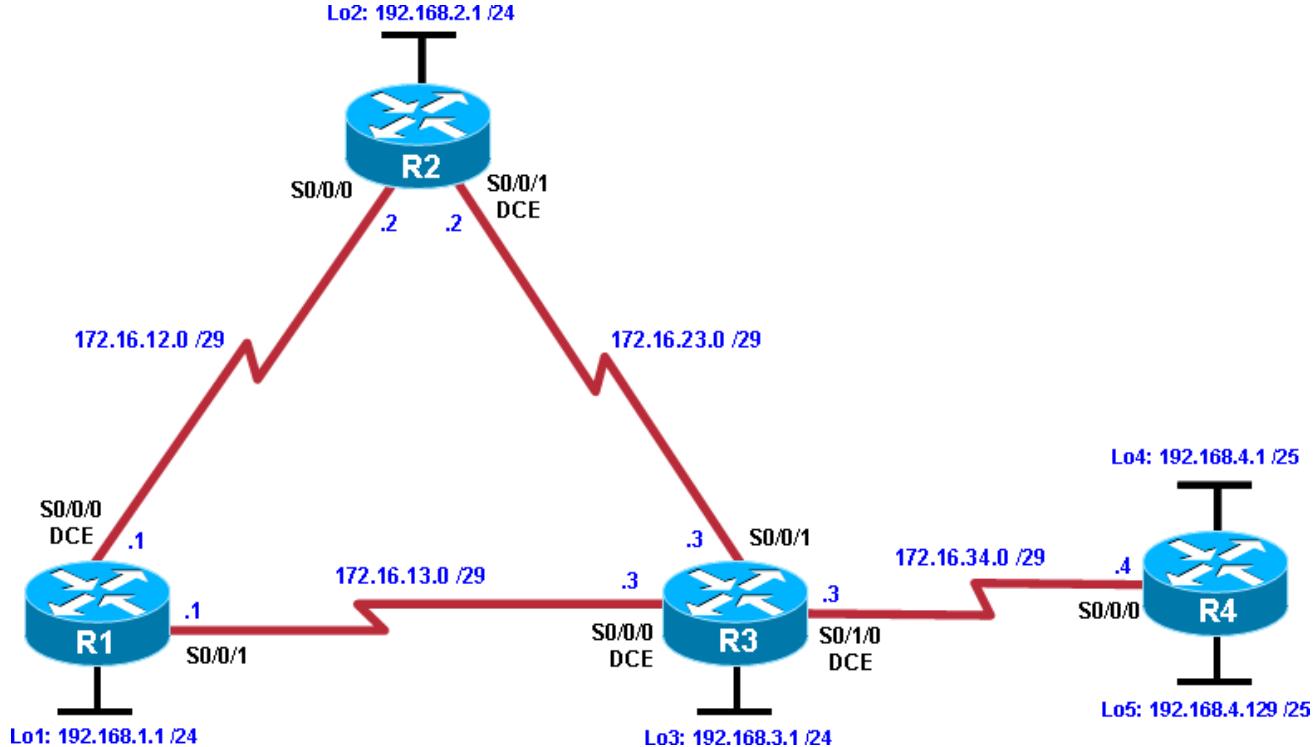
- h. Although a user can SSH from a host using the SSH option of TeraTerm or PuTTY, a router can also SSH to another SSH enabled device. SSH to R3 from R1.

```
R1# ssh -l ADMIN 10.2.2.2
Password:
Unauthorized access strictly prohibited!
R3>
R3> en
Password:
R3#
```

## Practical No.05

**Aim: Configure and Verify Path Control**

### Topology



### Objectives

- Configure and verify policy-based routing.
- Select the required tools and commands to configure policy-based routing operations.
- Verify the configuration and operation by using the proper **show** and **debug** commands.

### Background

You want to experiment with policy-based routing (PBR) to see how it is implemented and to study how it could be of value to your organization. To this end, you have interconnected and configured a test network with four routers. All routers are exchanging routing information using EIGRP.

**Note:** This lab uses Cisco 1841 routers with Cisco IOS Release 12.4(24)T1, and the Advanced IP Services image c1841-adipservicesk9-mz.124-24.T1.bin. You can use other routers (such as 2801 or 2811) and Cisco IOS Software versions if they have comparable capabilities and features. Depending on the router and software version, the commands available and output produced might vary from what is shown in this lab.

### Required Resources

- 4 routers (Cisco 1841 with Cisco IOS Release 12.4(24)T1 Advanced IP Services or comparable)
- Serial and console cables

### **Step 1: Prepare the routers for the lab.**

Cable the network as shown in the topology diagram. Erase the startup configuration, and reload each router to clear previous configurations.

### **Step 2: Configure router hostname and interface addresses.**

- Using the addressing scheme in the diagram, create the loopback interfaces and apply IP addresses to these and the serial interfaces on R1, R2, R3, and R4. On the serial interfaces connecting R1 to R3 and R3 to R4, specify the bandwidth as 64 Kb/s and set a clock rate on the DCE using the **clock rate 64000** command. On the serial interfaces connecting R1 to R2 and R2 to R3, specify the bandwidth as 128 Kb/s and set a clock rate on the DCE using the **clock rate 128000** command.

You can copy and paste the following configurations into your routers to begin.

**Note:** Depending on the router model, interfaces might be numbered differently than those listed. You might need to alter them accordingly.

#### **Router R1**

```
hostname R1
!
interface Lo1 description
    R1 LAN
    ip address 192.168.1.1 255.255.255.0
!
interface Serial0/0/0
    description R1 --> R2
    ip address 172.16.12.1 255.255.255.248
    clock rate 128000
    bandwidth 128no
    shutdown
!
interface Serial0/0/1
    description R1 --> R3
    ip address 172.16.13.1 255.255.255.248
    bandwidth 64no
    shutdown
!
end
```

#### **Router R2**

```
hostname R2
!
interface Lo2 description
    R2 LAN
    ip address 192.168.2.1 255.255.255.0
!
interface Serial0/0/0
    description R2 --> R1
    ip address 172.16.12.2 255.255.255.248
    bandwidth 128no
    shutdown
!
interface Serial0/0/1
    description R2 --> R3
    ip address 172.16.23.2 255.255.255.248
    clock rate 128000
```

```

bandwidth 128no
shutdown
!
end

Router R3

hostname R3
!
interface Lo3 description
  R3 LAN
  ip address 192.168.3.1 255.255.255.0
!
interface Serial0/0/0
  description R3 --> R1
  ip address 172.16.13.3 255.255.255.248
  clock rate 64000
  bandwidth 64no
  shutdown
!
interface Serial0/0/1
  description R3 --> R2
  ip address 172.16.23.3 255.255.255.248
  bandwidth 128no
  shutdown
!
interface Serial0/1/0
  description R3 --> R4
  ip address 172.16.34.3 255.255.255.248
  clock rate 64000
  bandwidth 64no
  shutdown
!
end

```

**Router R4**

```

hostname R4
!
interface Lo4 description R4
  LAN A
  ip address 192.168.4.1 255.255.255.128
!
interface Lo5 description R4
  LAN B
  ip address 192.168.4.129 255.255.255.128
!
interface Serial0/0/0
  description R4 --> R3
  ip address 172.16.34.4 255.255.255.248
  bandwidth 64no
  shutdown
!
end

```

- b. Verify the configuration with the **show ip interface brief**, **show protocols**, and **show interfaces description** commands. The output from router R3 is shown here as an example.

Interface	IP-Address	OK? Method Status	Protocol
FastEthernet0/0	unassigned	YES manual administratively down down	
FastEthernet0/1	unassigned	YES unset administratively down down	
Serial0/0/0	172.16.13.3	YES manual up	up
Serial0/0/1	172.16.23.3	YES manual up	up
Serial0/1/0	172.16.34.3	YES manual up	up
Serial0/1/1	unassigned	YES unset administratively down down	
Loopback3	192.168.3.1	YES manual up	up

R3# **show protocols**

Global values:

Internet Protocol routing is enabled  
 FastEthernet0/0 is administratively down, line protocol is down  
 FastEthernet0/1 is administratively down, line protocol is down  
 Serial0/0/0 is up, line protocol is up  
 Internet address is 172.16.13.3/29  
 Serial0/0/1 is up, line protocol is up  
 Internet address is 172.16.23.3/29  
 Serial0/1/0 is up, line protocol is up  
 Internet address is 172.16.34.3/29  
 Serial0/1/1 is administratively down, line protocol is down  
 Loopback3 is up, line protocol is up  
 Internet address is 192.168.3.1/24

R3# **show interfaces description**

Interface	Status	Protocol Description	
Fa0/0	admin down	down	
Fa0/1	admin down	down	
Se0/0/0	up	up	R3 --> R1
Se0/0/1	up	up	R3 --> R2
Se0/1/0	up	up	R3 --> R4
Se0/1/1	admin down	down	
Lo3	up	up	R3 LAN

### Step 3: Configure basic EIGRP.

- Implement EIGRP AS 1 over the serial and loopback interfaces as you have configured it for the other EIGRP labs.
- Advertise networks 172.16.12.0/29, 172.16.13.0/29, 172.16.23.0/29, 172.16.34.0/29, 192.168.1.0/24, 192.168.2.0/24, 192.168.3.0/24, and 192.168.4.0/24 from their respective routers.

You can copy and paste the following configurations into your routers.

#### Router R1

```
router eigrp 1
network 192.168.1.0
network 172.16.12.0 0.0.0.7
network 172.16.13.0 0.0.0.7
no auto-summary
```

**Router R2**

```
router eigrp 1
  network 192.168.2.0
  network 172.16.12.0 0.0.0.7
  network 172.16.23.0 0.0.0.7
  no auto-summary
```

**Router R3**

```
router eigrp 1
  network 192.168.3.0
  network 172.16.13.0 0.0.0.7
  network 172.16.23.0 0.0.0.7
  network 172.16.34.0 0.0.0.7
  no auto-summary
```

**Router R4**

```
router eigrp 1
  network 192.168.4.0
  network 172.16.34.0 0.0.0.7
  no auto-summary
```

You should see EIGRP neighbor relationship messages being generated.

**Step 4: Verify EIGRP connectivity.**

- Verify the configuration by using the **show ip eigrp neighbors** command to check which routers have EIGRP adjacencies.

```
R1# show ip eigrp neighbors
IP-EIGRP neighbors for process 1
H   Address           Interface      Hold Uptime    SRTT     RTO   Q   Seq
   Address           Interface      (sec)   (ms)       Cnt Num
1   172.16.13.3      Se0/0/1        12 00:00:58  127   2280  0   16
0   172.16.12.2      Se0/0/0        13 00:01:20  8     1140  0   17
```

```
R2# show ip eigrp neighbors
IP-EIGRP neighbors for process 1
H   Address           Interface      Hold Uptime    SRTT     RTO   Q   Seq
   Address           Interface      (sec)   (ms)       Cnt Num
1   172.16.23.3      Se0/0/1        10 00:01:30  15     1140  0   17
0   172.16.12.1      Se0/0/0        11 00:01:43  14     1140  0   180
```

```
R3# show ip eigrp neighbors
IP-EIGRP neighbors for process 1
H   Address           Interface      Hold Uptime    SRTT     RTO   Q   Seq
   Address           Interface      (sec)   (ms)       Cnt Num
2   172.16.34.4      Se0/1/0        10 00:02:51  27     2280  0   3
0   172.16.13.1      Se0/0/0        12 00:03:08  45     2280  0   19
1   172.16.23.2      Se0/0/1        12 00:03:13  12     1140  0   16
```

```
R4# show ip eigrp neighbors
```

LPI-EIGRP neighbors for process 1		Interface	Hold (sec)	Uptime	SRTT (ms)	RTO	Q Cnt	Seq Num
H	Address							
0	172.16.34.3	Se0/0/0	13	00:03:33	40	2280	0	15

Did you receive the output you expected?

- b. Run the following Tcl script on all routers to verify full connectivity.

R1# **tclsh**

```
foreach address {
172.16.12.1
172.16.12.2
172.16.13.1
172.16.13.3
172.16.23.2
172.16.23.3
172.16.34.3
172.16.34.4
192.168.1.1
192.168.2.1
192.168.3.1
192.168.4.1
192.168.4.129
} { ping $address }
```

You should get ICMP echo replies for every address pinged. Make sure to run the Tcl script on each router.

## Step 5: Verify the current path.

Before you configure PBR, verify the routing table on R1.

- a. On R1, use the **show ip route** command. Notice the next-hop IP address for all networks discovered by EIGRP.

R1# **show ip route**

Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP  
D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2 E1 - OSPF external type 1, E2 - OSPF external type 2  
i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2 ia - IS-IS inter area, \* - candidate default, U - per-user static  
route o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

```
172.16.0.0/29 is subnetted, 4 subnets
D      172.16.34.0 [90/41024000] via 172.16.13.3, 00:05:18, Serial0/0/1D    172.16.23.0
[90/21024000] via 172.16.12.2, 00:05:18, Serial0/0/0
C      172.16.12.0 is directly connected, Serial0/0/0C
          172.16.13.0 is directly connected, Serial0/0/1
192.168.4.0/25 is subnetted, 2 subnets
D      192.168.4.0 [90/41152000] via 172.16.13.3, 00:05:06, Serial0/0/1 D 192.168.4.128
[90/41152000] via 172.16.13.3, 00:05:06, Serial0/0/1
C      192.168.1.0/24 is directly connected, Loopback1
D      192.168.2.0/24 [90/20640000] via 172.16.12.2, 00:05:18, Serial0/0/0D    192.168.3.0/24
[90/21152000] via 172.16.12.2, 00:05:18, Serial0/0/0
```

- b. On R4, use the **traceroute** command to the R1 LAN address and source the ICMP packet from R4 LAN A and LAN B.

**Note:** You can specify the source as the interface address (for example 192.168.4.1) or the interface designator (for example, Fa0/0).

```
R4# traceroute 192.168.1.1 source 192.168.4.1
```

Type escape sequence to abort. Tracing the route to 192.168.1.1

```
1 172.16.34.3 12 msec 12 msec 16 msec
2 172.16.23.2 20 msec 20 msec 20 msec
3 172.16.12.1 28 msec 24 msec *
```

```
R4# traceroute 192.168.1.1 source 192.168.4.129
```

Type escape sequence to abort. Tracing the route to 192.168.1.1

```
1 172.16.34.3 12 msec 12 msec 16 msec
2 172.16.23.2 20 msec 20 msec 20 msec
3 172.16.12.1 28 msec 24 msec *
```

Notice that the path taken for the packets sourced from the R4 LANs are going through R3 --> R2 --> R1.

Why are the R4 interfaces not using the R3 --> R1 path?

- c. On R3, use the **show ip route** command and note that the preferred route from R3 to R1 LAN 192.168.1.0/24 is via R2 using the R3 exit interface S0/0/1.

R3# **show ip route**

Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP  
D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2 E1 - OSPF external type 1, E2 - OSPF external type 2  
i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2 ia - IS-IS inter area, \* - candidate default, U - per-user static  
route  
o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

```
172.16.0.0/29 is subnetted, 4 subnets
C      172.16.34.0 is directly connected, Serial0/1/0C
          172.16.23.0 is directly connected, Serial0/0/1
D      172.16.12.0 [90/21024000] via 172.16.23.2, 00:15:07, Serial0/0/1
C      172.16.13.0 is directly connected, Serial0/0/0192.168.4.0/25 is
          subnetted, 2 subnets
D      192.168.4.0 [90/40640000] via 172.16.34.4, 00:14:55, Serial0/1/0
D      192.168.4.128 [90/40640000] via 172.16.34.4, 00:14:55, Serial0/1/0 D192.168.1.0/24
[90/21152000] via 172.16.23.2, 00:15:07, Serial0/0/1 D      192.168.2.0/24 [90/20640000] via
172.16.23.2, 00:15:07, Serial0/0/1
C      192.168.3.0/24 is directly connected, Loopback3
```

- d. On R3, use the **show interfaces serial 0/0/0** and **show interfaces s0/0/1** commands.

```
R3# show interfaces s0/0/0
Serial0/0/0 is up, line protocol is upHardware is
  GT96K Serial Description: R3 --> R1
  Internet address is 172.16.13.3/29
  MTU 1500 bytes, BW 64 Kbit/sec, DLY 20000 usec, reliability
    255/255, txload 1/255, rxload 1/255
  Encapsulation HDLC, loopback not setKeepalive set
  (10 sec)
  CRC checking enabled
  Last input 00:00:00, output 00:00:00, output hang neverLast clearing of
  "show interface" counters never
  Input queue: 0/75/0/0 (size/max/drops/flushes); Total output drops: 0Queueing strategy: weighted
  fair
  Output queue: 0/1000/64/0 (size/max total/threshold/drops)Conversations
    0/1/256 (active/max active/max total) Reserved Conversations 0/0
    (allocated/max allocated) Available Bandwidth 48 kilobits/sec
  5 minute input rate 0 bits/sec, 0 packets/sec
  5 minute output rate 0 bits/sec, 0 packets/sec771 packets input,
  53728 bytes, 0 no buffer
  Received 489 broadcasts, 0 runts, 0 giants, 0 throttles
  0 input errors, 0 CRC, 0 frame, 0 overrun, 0 ignored, 0 abort768 packets output,
  54404 bytes, 0 underruns
  0 output errors, 0 collisions, 6 interface resets
  0 unknown protocol drops
  0 output buffer failures, 0 output buffers swapped out
  1 carrier transitions
  DCD=up DSR=up DTR=up RTS=up CTS=up

R3# show interfaces s0/0/1
Serial0/0/1 is up, line protocol is upHardware is
  GT96K Serial Description: R3 --> R2
  Internet address is 172.16.23.3/29
  MTU 1500 bytes, BW 128 Kbit/sec, DLY 20000 usec, reliability 255/255, txload
    1/255, rxload 1/255
  Encapsulation HDLC, loopback not setKeepalive set
  (10 sec)
  CRC checking enabled
  Last input 00:00:00, output 00:00:01, output hang neverLast clearing of
  "show interface" counters never
  Input queue: 0/75/0/0 (size/max/drops/flushes); Total output drops: 0Queueing strategy: weighted
  fair
  Output queue: 0/1000/64/0 (size/max total/threshold/drops)Conversations
    0/1/256 (active/max active/max total) Reserved Conversations 0/0
    (allocated/max allocated) Available Bandwidth 1158 kilobits/sec
  5 minute input rate 0 bits/sec, 0 packets/sec
  5 minute output rate 0 bits/sec, 0 packets/sec894 packets input,
  65653 bytes, 0 no buffer
  Received 488 broadcasts, 0 runts, 0 giants, 0 throttles
  0 input errors, 0 CRC, 0 frame, 0 overrun, 0 ignored, 0 abort895 packets output,
  66785 bytes, 0 underruns
  0 output errors, 0 collisions, 6 interface resets
  0 unknown protocol drops
```

```
0 output buffer failures, 0 output buffers swapped out
1 carrier transitions
DCD=up DSR=up DTR=up RTS=up CTS=up
```

Notice that the bandwidth of the serial link between R3 and R1 (S0/0/0) is set to 64 Kb/s, while the bandwidth of the serial link between R3 and R2 (S0/0/1) is set to 128 Kb/s.

- Confirm that R3 has a valid route to reach R1 from its serial 0/0/0 interface using the **show ip eigrp topology 192.168.1.0** command.

```
R3# show ip eigrp topology 192.168.1.0
IP-EIGRP (AS 1): Topology entry for 192.168.1.0/24
  State is Passive, Query origin flag is 1, 1 Successor(s), FD is 21152000Routing Descriptor
  Blocks:
    172.16.23.2 (Serial0/0/1), from 172.16.23.2, Send flag is 0x0Composite metric is
      (21152000/20640000), Route is InternalVector metric:
        Minimum bandwidth is 128 Kbit Total delay is
        45000 microsecondsReliability is 255/255
        Load is 1/255 Minimum
        MTU is 1500Hop count is
          2
    172.16.13.1 (Serial0/0/0), from 172.16.13.1, Send flag is 0x0Composite metric is
      (40640000/128256), Route is InternalVector metric:
        Minimum bandwidth is 64 Kbit Total delay is
        25000 microsecondsReliability is 255/255
        Load is 1/255 Minimum
        MTU is 1500Hop count is
          1
```

As indicated, R4 has two routes to reach 192.168.1.0. However, the metric for the route to R1 (172.16.13.1) is much higher (40640000) than the metric of the route to R2 (21152000), making the route through R2 the successor route.

## Step 6: Configure PBR to provide path control.

Now you will deploy source-based IP routing by using PBR. You will change a default IP routing decision based on the EIGRP-acquired routing information for selected IP source-to-destination flows and apply a different next-hop router.

Recall that routers normally forward packets to destination addresses based on information in their routing table. By using PBR, you can implement policies that selectively cause packets to take different paths based on source address, protocol type, or application type. Therefore, PBR overrides the router's normal routing behavior.

Configuring PBR involves configuring a route map with **match** and **set** commands and then applying the route map to the interface.

The steps required to implement path control include the following:

- Choose the path control tool to use. Path control tools manipulate or bypass the IP routing table. For PBR, **route-map** commands are used.
- Implement the traffic-matching configuration, specifying which traffic will be manipulated. The **match** commands are used within route maps.
- Define the action for the matched traffic using **set** commands within route maps.

- Apply the route map to incoming traffic.

As a test, you will configure the following policy on router R3:

- All traffic sourced from R4 LAN A must take the R3 --> R2 --> R1 path.
- All traffic sourced from R4 LAN B must take the R3 --> R1 path.

- On router R3, create a standard access list called **PBR-ACL** to identify the R4 LAN B network.

```
R3(config)# ip access-list standard PBR-ACL
R3(config-std-nacl)# remark ACL matches R4 LAN B traffic
R3(config-std-nacl)# permit 192.168.4.128 0.0.0.127
R3(config-std-nacl)# exit
```

- Create a route map called **R3-to-R1** that matches PBR-ACL and sets the next-hop interface to the R1 serial 0/0/1 interface.

```
R3(config)# route-map R3-to-R1 permit
R3(config-route-map)# match ip address PBR-ACL
R3(config-route-map)# set ip next-hop 172.16.13.1
R3(config-route-map)# exit
```

- Apply the R3-to-R1 route map to the serial interface on R3 that receives the traffic from R4. Use the **ip policy route-map** command on interface S0/1/0.

```
R3(config)# interface s0/1/0
R3(config-if)# ip policy route-map R3-to-R1
R3(config-if)# end
```

- On R3, display the policy and matches using the **show route-map** command.

```
R3# show route-map
route-map R3-to-R1, permit, sequence 10
Match clauses:
    ip address (access-lists): PBR-ACL
    Set clauses:
        ip next-hop 172.16.13.1
Policy routing matches: 0 packets, 0 bytes
```

**Note:** There are currently no matches because no packets matching the ACL have passed through R3 S0/1/0.

## Step 7: Test the policy.

Now you are ready to test the policy configured on R3. Enable the **debug ip policy** command on R3 so that you can observe the policy decision-making in action. To help filter the traffic, first create a standard ACL that identifies all traffic from the R4 LANs.

- On R3, create a standard ACL which identifies all of the R4 LANs.

```
R3# conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)# access-list 1 permit 192.168.4.0 0.0.0.255
R3(config)# exit
```

- Enable PBR debugging only for traffic that matches the R4 LANs.

```
R3# debug ip policy ?
<1-199> Access list
dynamic dynamic PBR
<cr>
```

```
R3# debug ip policy 1
Policy routing debugging is on for access list 1
```

- c. Test the policy from R4 with the **traceroute** command, using R4 LAN A as the source network.

```
R4# traceroute 192.168.1.1 source 192.168.4.1
```

Type escape sequence to abort. Tracing the route  
to 192.168.1.1

```
1 172.16.34.3 0 msec 0 msec 4 msec
2 172.16.23.2 0 msec 0 msec 4 msec
3 172.16.12.1 4 msec 0 msec *
```

Notice the path taken for the packet sourced from R4 LAN A is still going through R3 --> R2 --> R1.

As the traceroute was being executed, router R3 should be generating the following debug output.

R3#

```
*Feb 23 06:59:20.931: IP: s=192.168.4.1 (Serial0/1/0), d=192.168.1.1, len
28, policy rejected -- normal forwarding
*Feb 23 06:59:29.935: IP: s=192.168.4.1 (Serial0/1/0), d=192.168.1.1, len
28, policy rejected -- normal forwarding
*Feb 23 06:59:29.939: IP: s=192.168.4.1 (Serial0/1/0), d=192.168.1.1, len
28, policy rejected -- normal forwarding
*Feb 23 06:59:29.939: IP: s=192.168.4.1 (Serial0/1/0), d=192.168.1.1, len28,
*Feb 23 06:59:30.939: IP: s=192.168.4.1 (Serial0/1/0)68nrm(Serial0/1/0)gd=192.168.1.1, len28,
FIB policy rejected(no match) - normal forwarding
```

Why is the traceroute traffic not using the R3 --> R1 path as specified in the R3-to-R1 policy?

- d. Test the policy from R4 with the **traceroute** command, using R4 LAN B as the source network.

```
R4# traceroute 192.168.1.1 source 192.168.4.129
```

Type escape sequence to abort. Tracing the route  
to 192.168.1.1

```
1 172.16.34.3 12 msec 12 msec 16 msec
2 172.16.13.1 28 msec 28 msec *
```

Now the path taken for the packet sourced from R4 LAN B is R3 --> R1, as expected.

The debug output on R3 also confirms that the traffic meets the criteria of the R3-to-R1 policy.

R3#

```
*Feb 23 07:07:46.467: IP: s=192.168.4.129 (Serial0/1/0), d=192.168.1.1, le
n 28, policy match
*Feb 23 07:07:46.467: IP: route map R3-to-R1, item 10, permit
```

\*Feb 23 07:07:46.467: IP: s=192.168.4.129 (Serial0/1/0), d=192.168.1.1 (Serial0/0/0), len 28, policy routed  
 \*Feb 23 07:07:46.467: IP: Serial0/1/0 to Serial0/0/0 172.16.13.1  
 \*Feb 23 07:07:55.471: IP: s=192.168.4.129 (Serial0/1/0), d=192.168.1.1, len 28, policy match  
 \*Feb 23 07:07:55.471: IP: route map R3-to-R1, item 10, permit  
 \*Feb 23 07:07:55.471: IP: s=192.168.4.129 (Serial0/1/0), d=192.168.1.1 (Serial0/0/0), len 28, policy routed  
 \*Feb 23 07:07:55.471: IP: Serial0/1/0 to Serial0/0/0 172.16.13.1  
 \*Feb 23 07:07:55.471: IP: s=192.168.4.129 (Serial0/1/0), d=192.168.1.1, len 28, policy match  
 \*Feb 23 07:07:55.471: IP: route map R3-to-R1, item 10, permit  
 \*Feb 23 07:07:55.475: IP: s=192.168.4.129 (Serial0/1/0), d=192.168.1.1 (Serial0/0/0), len 28, policy routed  
 \*Feb 23 07:07:55.475: IP: Serial0/1/0 to Serial0/0/0 172.16.13.1  
 \*Feb 23 07:07:55.475: IP: s=192.168.4.129 (Serial0/1/0), d=192.168.1.1, len 28, FIB policy match  
 \*Feb 23 07:07:55.475: IP: s=192.168.4.129 (Serial0/1/0), d=192.168.1.1, g=172.16.13.1, len 28, FIB policy routed  
 \*Feb 23 07:08:04.483: IP: s=192.168.4.129 (Serial0/1/0), d=192.168.1.1, len 28, FIB policy match  
 \*Feb 23 07:08:04.483: IP: s=192.168.4.129 (Serial0/1/0), d=192.168.1.1, g=172.16.13.1, len 28, FIB policy routed  
 \*Feb 23 07:08:04.491: IP: s=192.168.4.129 (Serial0/1/0), d=192.168.1.1, len 28, FIB policy match  
 \*Feb 23 07:08:04.491: IP: s=192.168.4.129 (Serial0/1/0), d=192.168.1.1, g=172.16.13.1, len 28, FIB policy routed

- e. On R3, display the policy and matches using the **show route-map** command.

R3# **show route-map**

route-map R3-to-R1, permit, sequence 10

Match clauses:

ip address (access-lists): PBR-

ACLSet clauses:

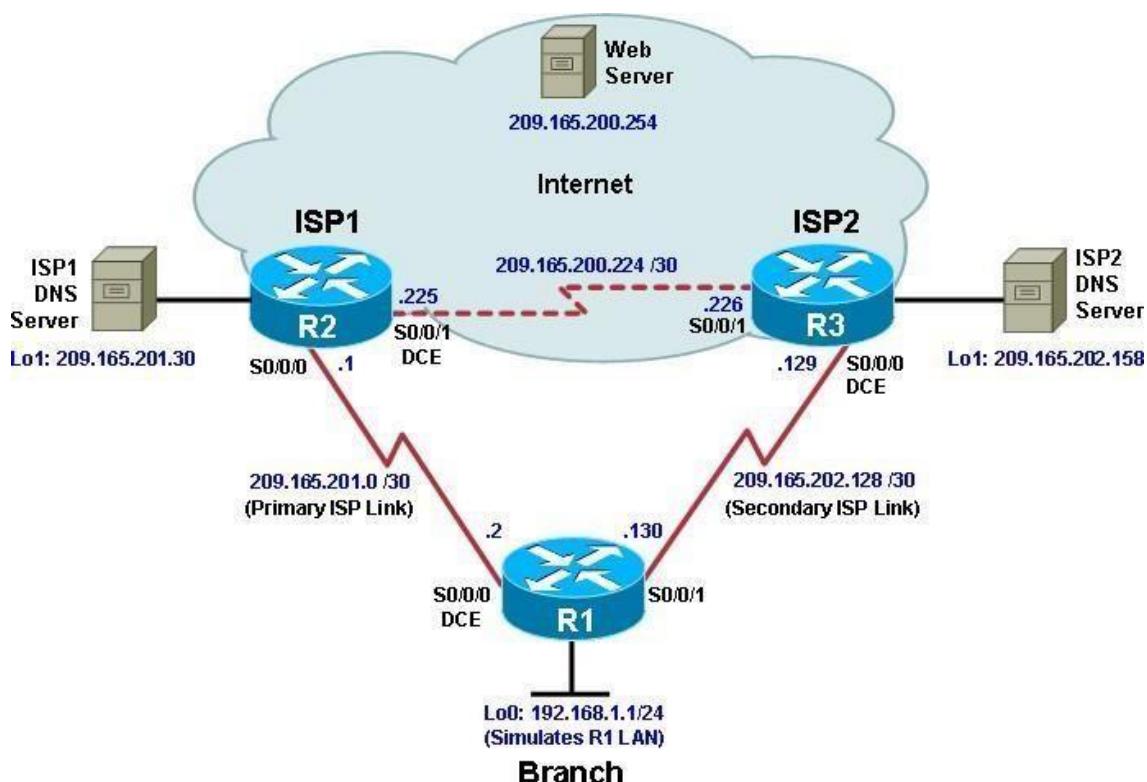
ip next-hop 172.16.13.1

Policy routing matches: 12 packets, 384 bytes

**Note:** There are now matches to the policy because packets matching the ACL have passed through R3S0/1/0.

**Aim: Configure IP SLA Tracking and Path Control with gateway**

## Topology



## Objectives

- Configure and verify the IP SLA feature.
- Test the IP SLA tracking feature.
- Verify the configuration and operation using **show** and **debug** commands.

## Background

You want to experiment with the Cisco IP Service Level Agreement (SLA) feature to study how it could be of value to your organization.

At times, a link to an ISP could be operational, yet users cannot connect to any other outside Internet resources. The problem might be with the ISP or downstream from them. Although policy-based routing (PBR) can be implemented to alter path control, you will implement the Cisco IOS SLA feature to monitor this behavior and intervene by injecting another default route to a backup ISP.

To test this, you have set up a three-router topology in a lab environment. Router R1 represents a branch office connected to two different ISPs. ISP1 is the preferred connection to the Internet, while ISP2 provides a backup link. ISP1 and ISP2 can also interconnect, and both can reach the web server. To monitor ISP1 for failure, you will configure IP SLA probes to track the reachability to the ISP1 DNS server. If connectivity to the ISP1 server fails, the SLA probes detect the failure and alter the default static route to point to the ISP2 server.

**Note:** This lab uses Cisco 1941 routers with Cisco IOS Release 15.2 with IP Base. Depending on the

## Required Resources

- 3 routers (Cisco IOS Release 15.2 or comparable)
- Serial and Ethernet cables

## Step 1: Configure loopbacks and assign addresses.

- p. Cable the network as shown in the topology diagram. Erase the startup configuration and reload each router to clear the previous configurations. Using the addressing scheme in the diagram, create the loopback interfaces and apply IP addresses to them as well as the serial interfaces on R1, ISP1, and ISP2.

You can copy and paste the following configurations into your routers to begin.

**Note:** Depending on the router model, interfaces might be numbered differently than those listed. You might need to alter them accordingly.

### Router R1

```
hostname R1

interface Loopback 0
description R1 LAN
ip address 192.168.1.1 255.255.255.0

interface Serial0/0/0
description R1 --> ISP1
ip address 209.165.201.2 255.255.255.252
clock rate 128000
bandwidth 128
no shutdown

interface Serial0/0/1
description R1 --> ISP2
ip address 209.165.202.130 255.255.255.252
bandwidth 128
no shutdown
```

### Router ISP1 (R2)

```
hostname ISP1

interface Loopback0
description Simulated Internet Web Server
ip address 209.165.200.254 255.255.255.255

interface Loopback1
description ISP1 DNS Server
ip address 209.165.201.30 255.255.255.255

interface Serial0/0/0
description ISP1 --> R1
ip address 209.165.201.1 255.255.255.252
bandwidth 128
no shutdown

interface Serial0/0/1
description ISP1 --> ISP2
ip address 209.165.200.225 255.255.255.252
clock rate 128000
```

### Router ISP2 (R3)

```
hostname ISP2
interface Loopback0
description Simulated Internet Web Server
ip address 209.165.200.254 255.255.255.255

interface Loopback1
description ISP2 DNS Server
ip address 209.165.202.158 255.255.255.255

interface Serial0/0/0
description ISP2 --> R1
ip address 209.165.202.129 255.255.255.252
clock rate 128000
bandwidth 128
no shutdown

interface Serial0/0/1
description ISP2 --> ISP1
ip address 209.165.200.226 255.255.255.252
bandwidth 128
no shutdown
```

- q. Verify the configuration by using the **show interfaces description** command. The output from router R1 is shown here as an example.

```
R1# show interfaces description | include up
Se0/0/0          up           up      R1 --> ISP1
Se0/0/1          up           up      R1 --> ISP2
Lo0              up           up      R1 LAN
R1#
```

All three interfaces should be active. Troubleshoot if necessary.

## Step 2: Configure static routing.

The current routing policy in the topology is as follows:

- Router R1 establishes connectivity to the Internet through ISP1 using a default static route.
- ISP1 and ISP2 have dynamic routing enabled between them, advertising their respective public address pools.
- ISP1 and ISP2 both have static routes back to the ISP LAN.

**Note:** For the purpose of this lab, the ISPs have a static route to an RFC 1918 private network address on the branch router R1. In an actual branch implementation, Network Address Translation (NAT) would be configured for all traffic exiting the branch LAN. Therefore, the static routes on the ISP routers would be pointing to the provided public pool of the branch office.

- a. Implement the routing policies on the respective routers. You can copy and paste the following configurations.

### Router R1

```
R1(config)# ip route 0.0.0.0 0.0.0.0 209.165.201.1
R1(config) #
```

### Router ISP1 (R2)

```
ISP1(config)# router eigrp 1
```

```

ISP1(config-router) # network 209.165.200.224 0.0.0.3
ISP1(config-router) # network 209.165.201.0 0.0.0.31
ISP1(config-router) # no auto-summary
ISP1(config-router) # exit
ISP1(config)#
ISP1(config-router) # ip route 192.168.1.0 255.255.255.0 209.165.201.2
ISP1(config)#
Router ISP2 (R3)

```

```

ISP2(config) # router eigrp 1
ISP2(config-router) # network 209.165.200.224 0.0.0.3
ISP2(config-router) # network 209.165.202.128 0.0.0.31
ISP2(config-router) # no auto-summary
ISP2(config-router) # exit
ISP2(config)#
ISP2(config) # ip route 192.168.1.0 255.255.255.0 209.165.202.130
ISP2(config)#

```

EIGRP neighbor relationship messages on ISP1 and ISP2 should be generated. Troubleshoot if necessary.

- b. The Cisco IOS IP SLA feature enables an administrator to monitor network performance between Cisco devices (switches or routers) or from a Cisco device to a remote IP device. IP SLA probes continuously check the reachability of a specific destination, such as a provider edge router interface, the DNS server of the ISP, or any other specific destination, and can conditionally announce a default route only if the connectivity is verified.

Before implementing the Cisco IOS SLA feature, you must verify reachability to the Internet servers. From router R1, ping the web server, ISP1 DNS server, and ISP2 DNS server to verify connectivity. You can copy the following Tcl script and paste it into R1.

```

foreach address {
209.165.200.254
209.165.201.30
209.165.202.158
} {
ping $address source 192.168.1.1
}

```

All pings should be successful. Troubleshoot if necessary.

- c. Trace the path taken to the web server, ISP1 DNS server, and ISP2 DNS server. You can copy the following Tcl script and paste it into R1.

```

foreach address {
209.165.200.254
209.165.201.30
209.165.202.158
} {
trace $address source 192.168.1.1
}

```

Through which ISP is traffic flowing?

---



---

### **Step 3: Configure IP SLA probes.**

When the reachability tests are successful, you can configure the Cisco IOS IP SLAs probes.

Different types of probes can be created, including FTP, HTTP, and jitter probes.

In this scenario, you will configure ICMP echo probes.

- d. Create an ICMP echo probe on R1 to the primary DNS server on ISP1 using the **ip sla** command.

```
R1(config)# ip sla 11
R1(config-ip-sla)# icmp-echo 209.165.201.30
R1(config-ip-sla-echo)# frequency 10
R1(config-ip-sla-echo)# exit
R1(config)#
R1(config)# ip sla schedule 11 life forever start-time now
R1(config)#

```

The operation number of 11 is only locally significant to the router. The **frequency 10** command schedules the connectivity test to repeat every 10 seconds. The probe is scheduled to start now and to run forever.

- e. Verify the IP SLAs configuration of operation 11 using the **show ip sla configuration 11** command.

```
R1# show ip sla configuration 11
IP SLAs Infrastructure Engine-III
Entry number: 11
Owner:
Tag:
Operation timeout (milliseconds): 5000
Type of operation to perform: icmp-echo
Target address/Source address: 209.165.201.30/0.0.0.0
Type Of Service parameter: 0x0
Request size (ARR data portion): 28
Verify data: No
Vrf Name:
Schedule:
  Operation frequency (seconds): 10 (not considered if randomly
  scheduled)
  Next Scheduled Start Time: Start Time already passed
  Group Scheduled : FALSE
  Randomly Scheduled : FALSE
  Life (seconds): Forever
  Entry Ageout (seconds): never
  Recurring (Starting Everyday): FALSE
  Status of entry (SNMP RowStatus): Active
Threshold (milliseconds): 5000
Distribution Statistics:
  Number of statistic hours kept: 2
  Number of statistic distribution buckets kept: 1
  Statistic distribution interval (milliseconds): 20
Enhanced History:
History Statistics:
  Number of history Lives kept: 0
  Number of history Buckets kept: 15
  History Filter Type: None

```

R1#

The output lists the details of the configuration of operation 11. The operation is an ICMP echo to 209.165.201.30, with a frequency of 10 seconds, and it has already started (the start time has already passed).

- f. Issue the **show ip sla statistics** command to display the number of successes, failures, and results of the latest operations.

```
R1# show ip sla statistics
IPSLAs Latest Operation Statistics

IPSLA operation id: 11
Latest RTT: 8 milliseconds
Latest operation start time: 10:33:18 UTC Sat Jan 10 2015
Latest operation return code: OK
Number of successes: 51
Number of failures: 0
Operation time to live: Forever
R1#
```

You can see that operation 11 has already succeeded five times, has had no failures, and the last operation returned an OK result.

- g. Although not actually required because IP SLA session 11 alone could provide the desired fault tolerance, create a second probe, 22, to test connectivity to the second DNS server located on router ISP2.

```
R1(config)# ip sla 22
R1(config-ip-sla)# icmp-echo 209.165.202.158
R1(config-ip-sla-echo)# frequency 10
R1(config-ip-sla-echo)# exit
R1(config)#
R1(config)# ip sla schedule 22 life forever start-time now
R1(config)# end
R1#
```

- h. Verify the new probe using the **show ip sla configuration** and **show ip sla statistics** commands.

```
R1# show ip sla configuration 22
IP SLAs Infrastructure Engine-III
Entry number: 22
Owner:
Tag:
Operation timeout (milliseconds): 5000
Type of operation to perform: icmp-echo
Target address/Source address: 209.165.202.158/0.0.0.0
Type Of Service parameter: 0x0
Request size (ARR data portion): 28
Verify data: No
Vrf Name:
Schedule:
  Operation frequency (seconds): 10 (not considered if randomly
  scheduled)
  Next Scheduled Start Time: Start Time already passed
  Group Scheduled : FALSE
  Randomly Scheduled : FALSE
  Life (seconds): Forever
  Entry Ageout (seconds): never
  Recurring (Starting Everyday): FALSE
  Status of entry (SNMP RowStatus): Active
Threshold (milliseconds): 5000
Distribution Statistics:
  Number of statistic hours kept: 2
  Number of statistic distribution buckets kept: 1
  Statistic distribution interval (milliseconds): 20
Enhanced History:
History Statistics:
  Number of history Lives kept: 0
```

Number of history Buckets kept: 15

History Filter Type: None

R1#

```
R1# show ip sla configuration 22
IP SLAs, Infrastructure Engine-II.
Entry number: 22
Owner:
Tag:
Type of operation to perform: icmp-echo
Target address/Source address: 209.165.201.158/0.0.0.0
Type Of Service parameter: 0x0
Request size (ARR data portion): 28
Operation timeout (milliseconds): 5000
Verify data: No
Vrf Name:
Schedule:
  Operation frequency (seconds): 10 (not considered if randomly
  scheduled)
    Next Scheduled Start Time: Start Time already passed
    Group Scheduled : FALSE
    Randomly Scheduled : FALSE
    Life (seconds): Forever
    Entry Ageout (seconds): never
    Recurring (Starting Everyday): FALSE
    Status of entry (SNMP RowStatus): Active
Threshold (milliseconds): 5000 (not considered if react RTT is
configured)
Distribution Statistics:
  Number of statistic hours kept: 2
  Number of statistic distribution buckets kept: 1
  Statistic distribution interval (milliseconds): 20
History Statistics:
  Number of history Lives kept: 0
  Number of history Buckets kept: 15
  History Filter Type: None
Enhanced History:
```

```
R1#
R1# show ip sla statistics 22
IPSLAs Latest Operation Statistics

IPSLA operation id: 22
  Latest RTT: 16 milliseconds
Latest operation start time: 10:38:29 UTC Sat Jan 10 2015
Latest operation return code: OK
Number of successes: 82
Number of failures: 0
Operation time to live: Forever
```

R1#

The output lists the details of the configuration of operation 22. The operation is an ICMP echo to 209.165.202.158, with a frequency of 10 seconds, and it has already started (the start time has

### Step 4: Configure tracking options.

Although PBR could be used, you will configure a floating static route that appears or disappears depending on the success or failure of the IP SLA.

- On R1, remove the current default route and replace it with a floating static route having an administrative distance of 5.

```
R1(config)# no ip route 0.0.0.0 0.0.0.0 209.165.201.1
R1(config)# ip route 0.0.0.0 0.0.0.0 209.165.201.1 5
R1(config)# exit
```

- Verify the routing table.

```
R1# show ip route | begin Gateway
Gateway of last resort is 209.165.201.1 to network 0.0.0.0
S*   0.0.0.0/0 [5/0] via 209.165.201.1
      192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C        192.168.1.0/24 is directly connected, Loopback0
L        192.168.1.1/32 is directly connected, Loopback0
      209.165.201.0/24 is variably subnetted, 2 subnets, 2 masks
C        209.165.201.0/30 is directly connected, Serial0/0/0
L        209.165.201.2/32 is directly connected, Serial0/0/0
      209.165.202.0/24 is variably subnetted, 2 subnets, 2 masks
C        209.165.202.128/30 is directly connected, Serial0/0/1
L        209.165.202.130/32 is directly connected, Serial0/0/1
R1#
```

Notice that the default static route is now using the route with the administrative distance of 5. The first tracking object is tied to IP SLA object 11.

- From global configuration mode on R1, use the **track 1 ip sla 11 reachability** command to enter the config-track subconfiguration mode.

```
R1(config)# track 1 ip sla 11 reachability
R1(config-track)#
```

- Specify the level of sensitivity to changes of tracked objects to 10 seconds of down delay and 1 second of up delay using the **delay down 10 up 1** command. The delay helps to alleviate the effect of flapping objects—objects that are going down and up rapidly. In this situation, if the DNS server fails momentarily and comes back up within 10 seconds, there is no impact.

```
R1(config-track)# delay down 10 up 1
R1(config-track)# exit
R1(config)#
```

- To view routing table changes as they happen, first enable the **debug ip routing** command.

```
R1# debug ip routing
IP routing debugging is on
R1#
```

- Configure the floating static route that will be implemented when tracking object 1 is active. Use the **ip route 0.0.0.0 0.0.0.0 209.165.201.1 2 track 1** command to create a floating static default route via 209.165.201.1 (ISP1). Notice that this command references the tracking object number 1, which in turn references IP SLA operation number 11.

```
R1(config)# ip route 0.0.0.0 0.0.0.0 209.165.201.1 2 track 1
R1(config)#
Jan 10 10:45:39.119: RT: updating static 0.0.0.0/0 (0x0) :
      via 209.165.201.1    0 1048578
```

```

Jan 10 10:45:39.119: RT: closer admin distance for 0.0.0.0, flushing 1
routes
Jan 10 10:45:39.119: RT: add 0.0.0.0/0 via 209.165.201.1, static metric
[2/0]
Jan 10 10:45:39.119: RT: updating static 0.0.0.0/0 (0x0) :
    via 209.165.201.1    0 1048578

Jan 10 10:45:39.119: RT: rib update return code: 17
Jan 10 10:45:39.119: RT: updating static 0.0.0.0/0 (0x0) :
    via 209.165.201.1    0 1048578

Jan 10 10:45:39.119: RT: rib update return code: 17
R1(config)#

```

Notice that the default route with an administrative distance of 5 has been immediately flushed because of a route with a better admin distance. It then adds the new default route with the admin distance of 2.

- o. Repeat the steps for operation 22, track number 2, and assign the static route an admin distance higher than track 1 and lower than 5. On R1, copy the following configuration, which sets an admin distance of 3.

```

R1(config)# track 2 ip sla 22 reachability
R1(config-track)# delay down 10 up 1
R1(config-track)# exit
R1(config)#
R1(config)# ip route 0.0.0.0 0.0.0.0 209.165.202.129 3 track 2
R1(config)#

```

- p. Verify the routing table again.

```

R1#show ip route | begin Gateway
Gateway of last resort is 209.165.201.1 to network 0.0.0.0

```

```

S*   0.0.0.0/0 [2/0] via 209.165.201.1
    192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C      192.168.1.0/24 is directly connected, Loopback0
L      192.168.1.1/32 is directly connected, Loopback0
    209.165.201.0/24 is variably subnetted, 2 subnets, 2 masks
C      209.165.201.0/30 is directly connected, Serial0/0/0
L      209.165.201.2/32 is directly connected, Serial0/0/0
    209.165.202.0/24 is variably subnetted, 2 subnets, 2 masks
C      209.165.202.128/30 is directly connected, Serial0/0/1
L      209.165.202.130/32 is directly connected, Serial0/0/1
R1#

```

Although a new default route was entered, its administrative distance is not better than 2. Therefore, it does not replace the previously entered default route.

### **Step 5: Verify IP SLA operation.**

In this step you observe and verify the dynamic operations and routing changes when tracked objects fail. The following summarizes the process:

- Disable the DNS loopback interface on ISP1 (R2).
- Observe the output of the **debug** command on R1.
- Verify the static route entries in the routing table and the IP SLA statistics of R1.
- Re-enable the loopback interface on ISP1 (R2) and again observe the operation of the IP SLA tracking feature.

- q. On ISP1, disable the loopback interface 1.

```
ISP1(config-if)# int lo1
ISP1(config-if)# shutdown
ISP1(config-if)#
Jan 10 10:53:25.091: %LINK-5-CHANGED: Interface Loopback1, changed
state to administratively down
Jan 10 10:53:26.091: %LINEPROTO-5-UPDOWN: Line protocol on Interface
Loopback1, changed state to down
ISP1(config-if)#

```

- r. On R1, observe the **debug** output being generated. Recall that R1 will wait up to 10 seconds before initiating action therefore several seconds will elapse before the output is generated.

```
R1#
Jan 10 10:53:59.551: %TRACK-6-STATE: 1 ip sla 11 reachability Up ->
Down
Jan 10 10:53:59.551: RT: del 0.0.0.0 via 209.165.201.1, static metric
[2/0]
Jan 10 10:53:59.551: RT: delete network route to 0.0.0.0/0
Jan 10 10:53:59.551: RT: default path has been cleared
Jan 10 10:53:59.551: RT: updating static 0.0.0.0/0 (0x0) :
via 209.165.202.129    0 1048578

Jan 10 10:53:59.551: RT: add 0.0.0.0/0 via 209.165.202.129, static
metric [3/0]
Jan 10 10:53:59.551: RT: default path is now 0.0.0.0 via
209.165.202.129
Jan 10 10:53:59.551: RT: updating static 0.0.0.0/0 (0x0) :
via 209.165.201.1    0 1048578

Jan 10 10:53:59.551: RT: rib update return code: 17
Jan 10 10:53:59.551: RT: updating static 0.0.0.0/0 (0x0) :
via 209.165.202.129    0 1048578

Jan 10 10:53:59.551: RT: updating static 0.0.0.0/0 (0x0) :
via 209.165.201.1    0 1048578

Jan 10 10:53:59.551: RT: rib update return code: 17
R1#
```

The tracking state of track 1 changes from up to down. This is the object that tracked reachability for IP SLA object 11, with an ICMP echo to the ISP1 DNS server at 209.165.201.30.

R1 then proceeds to delete the default route with the administrative distance of 2 and installs the next highest default route to ISP2 with the administrative distance of 3.

- s. On R1, verify the routing table.

```
R1# show ip route | begin Gateway
Gateway of last resort is 209.165.202.129 to network 0.0.0.0

S*   0.0.0.0/0 [3/0] via 209.165.202.129
      192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
C       192.168.1.0/24 is directly connected, Loopback0
L       192.168.1.1/32 is directly connected, Loopback0
      209.165.201.0/24 is variably subnetted, 2 subnets, 2 masks
C       209.165.201.0/30 is directly connected, Serial0/0/0
L       209.165.201.2/32 is directly connected, Serial0/0/0
```

```

209.165.202.0/24 is variably subnetted, 2 subnets, 2 masks
C          209.165.202.128/30 is directly connected, Serial0/0/1
L          209.165.202.130/32 is directly connected, Serial0/0/1
R1#

```

The new static route has an administrative distance of 3 and is being forwarded to ISP2 as it should.

- t. Verify the IP SLA statistics.

```

R1# show ip sla statistics
IPSLAs Latest Operation Statistics

IPSLA operation id: 11
    Latest RTT: NoConnection/Busy/Timeout
Latest operation start time: 11:01:08 UTC Sat Jan 10 2015
Latest operation return code: Timeout
Number of successes: 173
Number of failures: 45
Operation time to live: Forever
IPSLA operation id: 22
    Latest RTT: 8 milliseconds
Latest operation start time: 11:01:09 UTC Sat Jan 10 2015
Latest operation return code: OK
Number of successes: 218
Number of failures: 0
Operation time to live: Forever

```

R1#

Notice that the latest return code is **Timeout** and there have been 45 failures on IP SLA object 11.

- u. On R1, initiate a trace to the web server from the internal LAN IP address.

```

R1# trace 209.165.200.254 source 192.168.1.1
Type escape sequence to abort.
Tracing the route to 209.165.200.254
VRF info: (vrf in name/id, vrf out name/id)
  1 209.165.202.129 4 msec * *
R1#

```

This confirms that traffic is leaving router R1 and being forwarded to the ISP2 router.

- v. On ISP1, re-enable the DNS address by issuing the **no shutdown** command on the loopback 1 interface to examine the routing behavior when connectivity to the ISP1 DNS is restored.

```

ISP1(config-if)# no shutdown
Jan 10 11:05:45.847: %LINK-3-UPDOWN: Interface Loopback1, changed state
to up
Jan 10 11:05:46.847: %LINEPROTO-5-UPDOWN: Line protocol on Interface
Loopback1, changed state to up
ISP1(config-if)#

```

Notice the output of the **debug ip routing** command on R1.

```

R1#
Jan 10 11:06:20.551: %TRACK-6-STATE: 1 ip sla 11 reachability Down ->
Up
Jan 10 11:06:20.551: RT: updating static 0.0.0.0/0 (0x0) :

```

```
via 209.165.201.1      0 1048578
```

```
Jan 10 11:06:20.551: RT: closer admin distance for 0.0.0.0, flushing 1
routes
Jan 10 11:06:20.551: RT: add 0.0.0.0/0 via 209.165.201.1, static metric
[2/0]
Jan 10 11:06:20.551: RT: updating static 0.0.0.0/0 (0x0) :
via 209.165.202.129      0 1048578

Jan 10 11:06:20.551: RT: rib update return code: 17
Jan 10 11:06:20.551: RT: u
R1#pdating static 0.0.0.0/0 (0x0) :
via 209.165.202.129      0 1048578

Jan 10 11:06:20.551: RT: rib update return code: 17
Jan 10 11:06:20.551: RT: updating static 0.0.0.0/0 (0x0) :
via 209.165.201.1      0 1048578
```

```
Jan 10 11:06:20.551: RT: rib update return code: 17
R1#
```

Now the IP SLA 11 operation transitions back to an up state and reestablishes the default static route to ISP1 with an administrative distance of 2.

- w. Again examine the IP SLA statistics.

```
R1# show ip sla statistics
IPSLAs Latest Operation Statistics

IPSLA operation id: 11
Latest RTT: 8 milliseconds
Latest operation start time: 11:07:38 UTC Sat Jan 10 2015
Latest operation return code: OK
Number of successes: 182
Number of failures: 75
Operation time to live: Forever
```

```
IPSLA operation id: 22
Latest RTT: 16 milliseconds
Latest operation start time: 11:07:39 UTC Sat Jan 10 2015
Latest operation return code: OK
Number of successes: 257
Number of failures: 0
Operation time to live: Forever
```

```
R1#
```

The IP SLA 11 operation is active again, as indicated by the OK return code, and the number of successes is incrementing.

- x. Verify the routing table.

```
R1# show ip route | begin Gateway
Gateway of last resort is 209.165.201.1 to network 0.0.0.0
```

```
S*      0.0.0.0/0 [2/0] via 209.165.201.1
192.168.1.0/24 is variably subnetted, 2 subnets, 2 masks
```

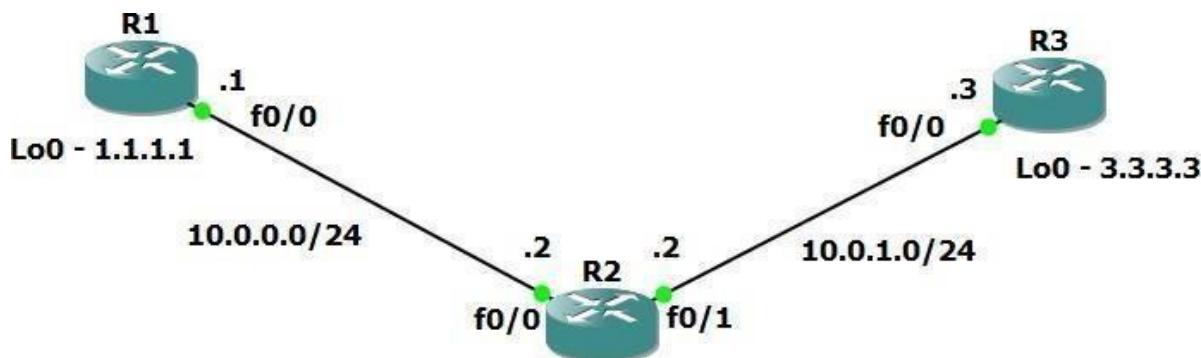
```
C      192.168.1.0/24 is directly connected, Loopback0
L      192.168.1.1/32 is directly connected, Loopback0
      209.165.201.0/24 is variably subnetted, 2 subnets, 2 masks
C      209.165.201.0/30 is directly connected, Serial0/0/0
L      209.165.201.2/32 is directly connected, Serial0/0/0
      209.165.202.0/24 is variably subnetted, 2 subnets, 2 masks
C      209.165.202.128/30 is directly connected, Serial0/0/1
L      209.165.202.130/32 is directly connected, Serial0/0/1
R1#
```

## Practical No. 07

### Aim: Configuring Basic MPLS Using OSPF

#### Step 1 – IP addressing of MPLS Core and OSPF

First bring 3 routers into your topology R1, R2, R3 position them as below. We are going to address the routers and configure ospf to ensure loopback to loopback connectivity between R1 and R3



```

R1
hostname
R1int lo0
ip add 1.1.1.1 255.255.255.255
ip ospf 1 area 0

int f0/0
ip add 10.0.0.1 255.255.255.0
no shut
ip ospf 1 area 0

R2
hostname
r2
int lo0

```

```
ip add 2.2.2.2 255.255.255.255
```

```
ip ospf 1 area 0
```

```
int f0/0
```

```
ip add 10.0.0.2 255.255.255.0
```

```
no shut
```

```
ip ospf 1 area 0
```

```
int f0/1
```

```
ip add 10.0.1.2 255.255.255.0
```

```
no shut
```

```
ip ospf 1 area 0
```

R3

```
hostname R3int
```

```
lo0
```

```
ip add 3.3.3.3 255.255.255.255
```

```
ip ospf 1 area 0
```

```
int f0/0
```

```
ip add 10.0.1.3 255.255.255.0
```

```
no shut
```

```
ip ospf 1 area 0
```

You should now have full ip connectivity between R1, R2, R3 to verify this we need to see if we can ping between the loopbacks of R1 and R3

```
R1#ping 3.3.3.3 source lo0

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 3.3.3.3, timeout is 2seconds:
Packet sent with a source address of 1.1.1.1
!!!!!

Success rate is 100 percent (5/5), round-tripmin/avg/max =
40/52/64 ms

R1#
```

You could show the routing table here, but the fact that you can ping between the loopbacks is verification enough and it is safe to move on.

## Step 2 – Configure LDP on all the interfaces in the MPLS Core

In order to run MPLS you need to enable it, there are two ways to do this.

- At each interface enter the **mpls ip** command
- Under the ospf process use the **mpls ldp autoconfig** command

For this tutorial we will be using the second option, so go into the ospf process and enter mpls ldp autoconfig – this will enable mpls label distribution protocol on every interface running ospf under that specific process.

```
R1

router ospf 1

mpls ldp autoconfig
```

```
R2

router ospf 1
```

```
mpls ldp autoconfig
```

```
R3
```

```
router ospf 1
```

```
mpls ldp autoconfig
```

You should see log messages coming up showing the LDP neighbors are up.

```
R2#
```

```
*Mar 1 00:31:53.643: %SYS-5-CONFIG_I: Configured fromconsole
```

```
*Mar 1 00:31:54.423: %LDP-5-NBRCHG: LDP Neighbor  
1.1.1.1:0 (1) is UPR2#
```

```
*Mar 1 00:36:09.951: %LDP-5-NBRCHG: LDP Neighbor  
3.3.3.3:0 (2) is UP
```

To verify the mpls interfaces the command is very simple – **sh mpls interface**

This is done on R2 and you can see that both interfaces are running mpls and using LDP

```
R2#sh mpls interface
```

Interface	IP	Tunnel
Operational		
l		
FastEthernet0/0	Yes (ldp)	No Yes
FastEthernet0/1	Yes (ldp)	No Yes

You can also verify the LDP neighbors with the **sh mpls ldp neighbors** command.

---

```
R2#sh mpls ldp neigh

    Peer LDP Ident: 1.1.1.1:0; Local LDP Ident 2.2.2.2:0TCP connection:
        1.1.1.1.646 - 2.2.2.2.37909
    State: Oper; Msgs sent/rcvd: 16/17;
    DownstreamUp time: 00:07:46
    LDP discovery sources:
        FastEthernet0/0, Src IP addr:
            10.0.0.1Addresses bound to peer LDP
            10.0.0.          1.1.1.

    Peer LDP Ident: 3.3.3.3:0; Local LDP Ident
        2.2.2.2:0TCP connection: 3.3.3.3.22155 -
        2.2.2.2.646
    State: Oper; Msgs sent/rcvd: 12/11;
    DownstreamUp time: 00:03:30
    LDP discovery sources:
        FastEthernet0/1, Src IP addr:
            10.0.1.          3.3.3.
```

One more verification to confirm LDP is running ok is to do a trace between R1 and R3 and verify if you get MPLS Labels show up in the trace.

```
R1#trace 3.3.3.3

Type escape sequence to abort. Tracing the
route to 3.3.3.3

1 10.0.0.2 [MPLS: Label 17 Exp 0] 84 msec 72 msec 44msec
2 10.0.1.3 68 msec 60 msec *
```

As you can see the trace to R2 used an MPLS Label in the path, as this is a very small MPLS core only one label was used as R3 was the final hop.

So to review we have now configured IP addresses on the MPLS core, enabled OSPF and full IP connectivity between all routers and finally enabled mpls on all the interfaces in the core and have established Idpneighbors between all routers.

The next step is to configure MP-BGP between R1 and R3

This is when you start to see the layer 3 vpn configuration come to life

### Step 3 – MPLS BGP Configuration between R1 and R3

We need to establish a Multi Protocol BGP session between R1 and R3 this is done by configuring the vpng4 address family as below

```
R1#  
  
router bgp 1  
  
neighbor 3.3.3.3 remote-as 1  
  
neighbor 3.3.3.3 update-source Loopback0 no auto-summary  
!  
  
address-family vpnv4 neighbor  
    3.3.3.3 activate
```

```
R3#  
  
router bgp 1  
  
neighbor 1.1.1.1 remote-as 1  
  
neighbor 1.1.1.1 update-source Loopback0 no auto-summary  
!  
  
address-family vpnv4 neighbor  
    1.1.1.1 activate
```

```
*Mar 1 00:45:01.047: %BGP-5-ADJCHANGE: neighbor 1.1.1.1
Up
```

You should see log messages showing the BGP sessions coming up.

To verify the BGP session between R1 and R3 issue the command **sh bgp vpnv4 unicast all summary**

```
R1#sh bgp vpnv4 unicast all summary
BGP router identifier 1.1.1.1, local AS number 1
BGP table version is 1, main routing table

BGP peer statistics
Peer          V     AS MsgRcvd      TblVer
OutQ          State/PfxRc
Up/Down        4       1      21      21      1      0
3.3.3.3          0
```

You can see here that we do have a bgp vpnv4 peering to R3 – looking at the PfxRcd you can see it says 0 this is because we have not got any routes in BGP. We are now going to add two more routers to the topology. These will be the customer sites connected to R1 and R3. We will then create a VRF on each router and put the interfaces connected to each site router into that VRF.

#### Step 4 – Add two more routers, create VRFs

We will add two more routers into the topology so it now looks like the final topology

Router 4 will peer OSPF using process number 2 to a VRF configured on R1. It will use the local site addressing of 192.168.1.0/24.

```
R4
int lo0
ip add 4.4.4.4 255.255.255.255
ip ospf 2 area 2
int f0/0
```

```
ip add 192.168.1.4 255.255.255.0

ip ospf 2 area 2no
shut

R1

int f0/1
no shut
ip add 192.168.1.1 255.255.255.0
```

Now at this point we have R4 peering to R1 but in the global routing table of R1 which is not what we want.

## We are now going to start using VRF's

What is a VRF in networking?

Virtual routing and forwarding (**VRF**) is a technology included in IP (InternetProtocol) that allows multiple instances of a routing table to co-exist in a router and work together but not interfere with each other.. This increases functionality by allowing network paths to be segmented without using multiple devices.

As an example if R1 was a PE Provider Edge router of an ISP and it had two customers that were both addressed locally with the 192.168.1.0/24 addressspace it could accommodate both their routing tables in different VRFs – it distinguishes between the two of them using a Route Distinguisher

So back to the topology – we now need to create a VRF on R1For this mpls tutorial I will be using VRF RED

```
R1

ip vrf RED
rd 4:4
route-target both 4:4
```

The RD and route-target do not need to be the same – Route Distinguisher

**Route Distinguisher vs Route Target** before proceeding.

So now we have configured the VRF on R1 we need to move the interface F0/1 into that VRF

```
R1
int f0/1
ip vrf forwarding RED
```

Now notice what happens when you do that – the IP address is removed

```
R1(config-if)#ip vrf fo
R1(config-if)#ip vrf forwarding RED
% Interface FastEthernet0/1 IP address 192.168.1.1 removed due to
enabling VRF RED
```

You just need to re-apply it

```
R1
int f0/1
ip address 192.168.1.1 255.255.255.0
```

Now if we view the config on R1 int f0/1 you can see the VRF configured.

```
R1
R1#sh run int f0/1 Building
configuration...
Current configuration : 119 bytes
P!
```

```
interface FastEthernet0/1ip vrf
    forwarding RED
    ip address 192.168.1.1 255.255.255.0
    duplex auto
    speed auto
end

R1#
```

Now we can start to look int VRF's and how they operate – you need to understand now that there are 2 routing tables within R1

- The Global Routing Table
- The Routing Table for VRF RED

If you issue the command **sh ip route** this shows the routes in the globaltable and you will notice that you do not see 192.168.1.0/24

```
R1#sh ip route

Codes: C - connected, S - static, R - RIP, M - mobile, B
      - BGP

      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF
      inter area

      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2

      E1 - OSPF external type 1, E2 - OSPF external type 2

      i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 -
      IS-IS level-2

      ia - IS-IS inter area, * - candidate default, U - per-user static route

      o - ODR, P - periodic downloaded static route

Gateway of last resort is not set
```

```
1.0.0.0/32 is subnetted, 1 subnets
C 1.1.1.1 is directly connected, Loopback02.0.0.0/32 is
subnetted, 1 subnets
O 2.2.2.2 [110/11] via 10.0.0.2, 01:03:48,
FastEthernet0/0
3.0.0.0/32 is subnetted, 1 subnets
O 3.3.3.3 [110/21] via 10.0.0.2, 01:02:29,
FastEthernet0/0
10.0.0.0/24 is subnetted, 2 subnets
C 10.0.0.0 is directly connected, FastEthernet0/0 10.0.1.0 [110/20]
via 10.0.0.2, 01:02:39,
FastEthernet0/0R1#
```

If you now issue the command `sh ip route vrf red` – this will show the routes in the routing table for VRF RED

```
R1#sh ip route vrf red
% IP routing table red does not exist
```

**NOTE:** The VRF name is case sensitive!

```
R1#sh ip route vrf RED
Routing Table: RED
Codes: C - connected, S - static, R - RIP, M - mobile, B
      - BGP
      D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF
          inter area
      N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
```

```
E1 - OSPF external type 1, E2 - OSPF external type 2
```

```
i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 -  
IS-IS level-2
```

```
ia - IS-IS inter area, * - candidate default, U - per-user static route
```

```
o - ODR, P - periodic downloaded static route
```

```
Gateway of last resort is not set
```

```
C 192.168.1.0/24 is directly connected, FastEthernet0/1R1#
```

We just need to enable OSPF on this interface and get the loopback address for R4 in the VRF RED routing table before proceeding.

```
R1
```

```
int f0/1  
ip ospf 2 area 2
```

You should see a log message showing the OSPF neighbor come up

```
R1(config-if)#  
  
*Mar 1 01:12:54.323: %OSPF-5-ADJCHG: Process 2, Nbr  
4.4.4.4  
  
on FastEthernet0/1 from LOADING to FULL, Loading Done
```

If we now check the routes in the VRF RED routing table you should see **4.4.4.4** in there as well.

```
R1#sh ip route vrf RED
```

Routing Table: RED

Codes: C - connected, S - static, R - RIP, M - mobile,B - BGP

D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area

N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2

E1 - OSPF external type 1, E2 - OSPF external type 2

i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2

ia - IS-IS inter area, \* - candidate default, U - per-user static route

o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

4.0.0.0/32 is subnetted, 1 subnets

**O 4.4.4.4 [110/11] via 192.168.1.4, 00:00:22,  
FastEthernet0/1**

C 192.168.1.0/24 is directly connected, FastEthernet0/1R1#

We now need to repeat this process for R3 & R6

Router 6 will peer OSPF using process number 2 to a VRF configured on R3. It will use the local site addressing of 192.168.2.0/24.

```
R6
int lo0
ip add 6.6.6.6 255.255.255.255
ip ospf 2 area 2
int f0/0
ip add 192.168.2.6 255.255.255.0
ip ospf 2 area 2
no shut

R3
int f0/1
no shut
ip add 192.168.2.3 255.255.255.0
```

We also need to configure a VRF onto R3 as well.

```
R3
ip vrf RED
rd 4:4
route-target both 4:4
```

So now we have configured the VRF on R3 we need to move the interface F0/1 into that VRF

```
R3
int f0/1
ip vrf forwarding RED
```

Now notice what happens when you do that – the IP address is removed

```
R3(config-if)#ip vrf forwarding RED
% Interface FastEthernet0/1 IP address 192.168.2.1 removed due to
enabling VRF RED
```

You just need to re-apply it

```
R3
int f0/1
ip address 192.168.2.1 255.255.255.0
```

Now if we view the config on R3 int f0/1 you can see the VRF configured.

```
R3
```

```
R3#sh run int f0/1 Building
configuration...

Current configuration : 119 bytes

!
interface FastEthernet0/1ip vrf
forwarding RED
ip address 192.168.2.1 255.255.255.0
duplex auto
speed auto
end
```

Finally we just need to enable OSPF on that interface and verify the routes are in the RED routing table.

```
R3
int f0/1
ip ospf 2 area 2
```

Check the routes in vrf RED

```
R3
R3#sh ip route vrf RED

Routing Table: RED

Codes: C - connected, S - static, R - RIP, M - mobile, B
- BGP

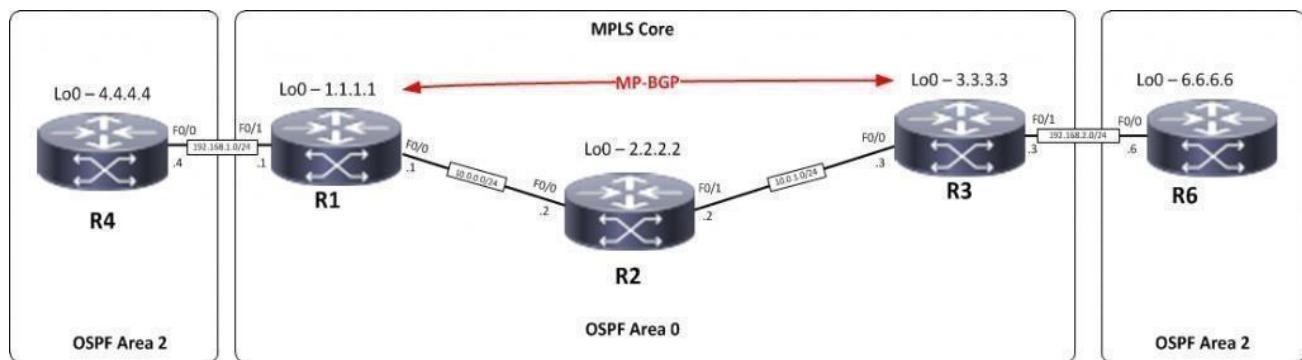
Gateway of last resort is not set
```

```
6.0.0.0/32 is subnetted, 1 subnets
```

```
O      6.6.6.6 [110/11] via 192.168.2.6,
FastEthernet0/
```

```
C      192.168.2.0/24 is directly
FastEthernet0/
1
```

Ok so we have come a long way now let's review the current situation. We now have this setup



R1,R2,R3 form the MPLS Core and are running OSPF with all loopbacks running a /32 address and all have full connectivity. R1 and R3 are peering with MP-BGP. LDP is enabled on all the internal interfaces. The external interfaces of the MPLS core have been placed into a VRF called RED and then a site router has been joined to that VRF on each side of the MPLS core – (These represent a small office)

The final step to get full connectivity across the MPLS core is to redistribute the routes in OSPF on R1 and R3 into MP-BGP and MP-BGP into OSPF, this is what we are going to do now.

We need to redistribute the OSPF routes from R4 into BGP in the VRF on R1, the OSPF routes from R6 into MP-BGP in the VRF on R3 and then the routes in MP-BGP in R1 and R3 back out to OSPF

Before we start lets do some verifications

### Check the routes on R4

```
R4#sh ip route
4.0.0.0/32 is subnetted, 1 subnets
C 4.4.4.4 is directly connected, Loopback0
```

```
C 192.168.1.0/24 is directly connected, FastEthernet0/0
```

As expected we have the local interface and the loopback address.

When we are done we want to see 6.6.6.6 in there so we can ping across the MPLS

### Check the routes on R1

```
R1#sh ip route

1.0.0.0/32 is subnetted, 1 subnets
C 1.1.1.1 is directly connected, Loopback02.0.0.0/32 is
subnetted, 1 subnets

O 2.2.2.2 [110/11] via 10.0.0.2, 00:01:04,
FastEthernet0/0

3.0.0.0/32 is subnetted, 1 subnets
O 3.3.3.3 [110/21] via 10.0.0.2, 00:00:54,
FastEthernet0/0

10.0.0.0/24 is subnetted, 2 subnets

C 10.0.0.0 is directly connected, FastEthernet0/0 10.0.1.0 [110/20]
via 10.0.0.2, 00:00:54,
FastEthernet0/0
```

Remember we have a VRF configured on this router so this command will show routes in the global routing table (the MPLS Core) and it will not show the 192.168.1.0/24 route as that is in VRF RED – to see that we run the following command

```
R1#sh ip route vrf RED

Routing Table: RED

4.0.0.0/32 is subnetted, 1 subnets
```

```
O 4.4.4.4 [110/11] via 192.168.1.4, 00:02:32,  
FastEthernet0/1  
  
C 192.168.1.0/24 is directly connected, FastEthernet0/1
```

Here you can see Routing Table: RED is shown and the routes to R4 are now visible with 4.4.4.4 being in OSPF. So we need to do

the following;

- Redistribute OSPF into MP-BGP on R1
- Redistribute MP-BGP into OSPF on R1
- Redistribute OSPF into MP-BGP on R3
- Redistribute MP-BGP into OSPF on R3

### Redistribute OSPF into MP-BGP on R1

```
R1  
  
router bgp 1  
  
address-family ipv4 vrf RED  
redistribute ospf 2
```

### Redistribute OSPF into MP-BGP on R3

```
R3  
  
router bgp 1  
  
address-family ipv4 vrf RED  
redistribute ospf 2
```

This has enabled redistribution of the OSPF routes into BGP. We can check the routes from R4 and R6 are now showing in the BGP table for their VRF with this command

**sh ip bgp vpnv4 vrf RED**

```
R1#sh ip bgp vpnv4 vrf RED  
  
BGP table version is 9, local router ID is 1.1.1.1  
  
Status codes: s suppressed, d damped, h history, *valid, > best,
```

```
r RIB-failure, S Stale
```

```
Origin codes: i IGP, e - EGP, ? - incomplete
-
```

Network	Next Hop	Metric	LocPrf	Weight	Path
---------	----------	--------	--------	--------	------

Route Distinguisher: 4:4 (default for vrf RED)

**\*> 4.4.4.4/32 192.168.1.4 11 32768 ?**

**\*>i6.6.6.6/32 3.3.3.3 100 0 ?**

**11**

**\*> 192.168.1.0 0.0.0.0 32768 ?**

**0**

**\*>i192.168.2.0 3.3.3.3 100 0 ?**

**0**

Here we can see that 4.4.4.4 is now in the BGP table in VRF RED on R1 with a next hop of 192.168.1.4 (R4) and also 6.6.6.6 is in there as wellwith a next hop of 3.3.3.3 (which is the loopback of R3 – showing that it isgoing over the MPLS and R1 is not in the picture)

The same should be true on R3

```
R3#sh ip bgp vpng4 vrf RED
```

BGP table version is 9, local router ID is 3.3.3.3

Status codes: s suppressed, d damped, h history, \*valid, > best, i - internal,

```
r RIB-failure, S Stale
```

Origin codes: i - IGP, e - EGP, ? - incomplete

Network Next Hop Metric LocPrf Weight Path Route Distinguisher: 4:4

(default for vrf RED)

**\*>i4.4.4.4/32 1.1.1.1 11 100 0 ?**

**\*> 6.6.6.6/32 192.168.2.6 11 32768 ?**

**\*>i192.168.1.0 1.1.1.1 0 100 0 ?**

```
*> 192.168.2.0 0.0.0.0 0 32768 ?
```

Which it is! 6.6.6.6 is now in the BGP table in VRF RED on R3 with a next hop of 192.168.2.6 (R6) and also 4.4.4 is in there as well with a next hop of 1.1.1.1 (which is the loopback of R1 – showing that it is going over the MPLS and R2 is not in the picture). The final step is to get the routes that have come across the MPLS back into OSPF and then we can get end to end connectivity.

```
R1
router ospf 2 redistribute bgp 1
subnets

R3
router ospf 2 redistribute bgp 1
subnets
```

If all has worked we should be now able to ping 6.6.6.6 from R4

### **Before we do let's see what the routing table looks like on R4**

```
R4#sh ip route
4.0.0.0/32 is subnetted, 1 subnets
C 4.4.4.4 is directly connected, Loopback0 6.0.0.0/32 is
subnetted, 1 subnets
O IA 6.6.6.6 [110/21] via 192.168.1.1, 00:01:31,
FastEthernet0/0
C 192.168.1.0/24 is directly connected, FastEthernet0/0
```

```
O E2 192.168.2.0/24 [110/1] via 192.168.1.1, 00:01:31,  
FastEthernet0/0
```

Great we have 6.6.6.6 in there

### Also check the routing table on R6

```
R6#sh ip route  
  
4.0.0.0/32 is subnetted, 1 subnets  
  
O IA 4.4.4.4 [110/21] via 192.168.2.1, 00:01:22,  
FastEthernet0/0  
  
6.0.0.0/32 is subnetted, 1 subnets  
  
C 6.6.6.6 is directly connected, Loopback00 IA  
  
192.168.1.0/24 [110/11] via  
192.168.2.1, 00:01:22, FastEthernet0/0  
  
C 192.168.2.0/24 is directly connected, FastEthernet0/0
```

Brilliant we have 4.4.4.4 in there so we should be able to ping across the MPLS

```
R4#ping 6.6.6.6  
  
Type escape sequence to abort.  
  
Sending 5, 100-byte ICMP Echos to 6.6.6.6, timeout is 2seconds:  
!!!!!  
  
Success rate is 100 percent (5/5), round-tripmin/avg/max=  
40/48/52ms
```

Which we can – to prove this is going over the MPLS and be label switched and not routed, lets do a trace

```
R4#trace 6.6.6.6
```

```
Type escape sequence to abort. Tracing the  
route to 6.6.6.6
```

```
1 192.168.1.1 20 msec 8 msec 8 msec  
2 10.0.0.2 [MPLS: Labels 17/20 Exp 0] 36 msec 40 msec  
36 msec  
3 192.168.2.1 [MPLS: Label 20 Exp 0] 16 msec 40 msec 16 msec  
4 192.168.2.6 44 msec 40 msec 56 msecR4#
```

**University of Mumbai**

**Practical Journal of  
Computer Vision**

**M.Sc. (Information Technology) Part-I**

**Submitted by**

**SINGH MANASI RAKESH**

**Seat No: 1312502**



**DEPARTMENT OF INFORMATION TECHNOLOGY  
PILLAI HOC COLLEGE OF ARTS, SCIENCE & COMMERCE, RASAYANI  
(Affiliated to Mumbai University)  
RASAYANI, 410207  
MAHARASHTRA  
2023-2024**

**Mahatma Education Society's  
Pillai Hoc College of Arts, Science & Commerce, Rasayani  
(Affiliated to Mumbai University)  
RASAYANI – MAHARASHTRA - 410207**

**DEPARTMENT OF  
INFORMATION TECHNOLOGY**



**CERTIFICATE**

This is to certify that the experiment work entered in this journal is as per the syllabus in **M.Sc. (Information Technology) Part-I, Semester-II**; class prescribed by University of Mumbai for the subject **Computer Vision** was done in computer lab of Mahatma Education Society's Pillai HOC College of Arts, Science & Commerce, Rasayani by **MANASI SINGH** during Academic year 2023-2024.

**Exam Seat No: 1312502**

**Subject In-Charge**

**Coordinator**

**External Examiner**

**Principal**

**Date:**

**College Seal**

# **COMPUTER VISION**

Sr. No	Title	Date	Signature
1	Perform Geometric transformation. A. Image Scaling. B. Image Shrinking. C. Image Rotation. D. Affine Transformation. E. Perspective Transformation.	26/03/2024	
	F. Shearing X-axis. G. Shearing Y-axis. H. Reflected Image. I. Cropped Image.	30/03/2024	
2	Perform Image Stitching.	02/04/2024	
3	Perform Camera Calibration.	15/04/2024	
4	A. Perform the following Face detection.  B. Perform the following: i. Object detection ii. Line Detection iii. Hough transform  C. Perform the following Pedestrian detection.	20/04/2024 25/04/2024 29/04/2024	
	D. Perform the following Face Recognition.	01/05/2024	
5	A. Implement object detection and tracking from video.  B. Implement object detection and tracking from video. (Count number of Faces using Python)	07/05/2024 15/05/2024	
	C. Implement object detection and tracking from video. (Object Tracking using Homography)	21/05/2024	
6	Perform Colorization.	25/05/2024	
7	Perform Text Detection and Recognition.	30/05/2024	
8	Construct 3D model from Images.	01/06/2024	
9	Perform Feature extraction using RANSAC.	12/06/2024	

MSC IT Part 1 Sem 2

Roll No.: 00

# Practical - 1

**Aim:** Perform Geometric transformation.

## Theory:

### Geometric Transformations:

Geometric transformation is a fundamental technique used in image processing that involves manipulating the spatial arrangement of pixels in an image. It is used to modify the geometric properties of an image, such as its size, shape, position, and orientation.

OpenCV provides two transformation functions, `cv2.warpAffine` and `cv2.warpPerspective`, with which you can have all kinds of transformations. `cv2.warpAffine` takes a  $2 \times 3$  transformation matrix while `cv2.warpPerspective` takes a  $3 \times 3$  transformation matrix as input.

### Translation

Translation is the shifting of object's location. If you know the shift in  $(x,y)$  direction, let it be  $(t_x, t_y)$ , you can create the transformation matrix  $M$  as follows:

$$M = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix}$$

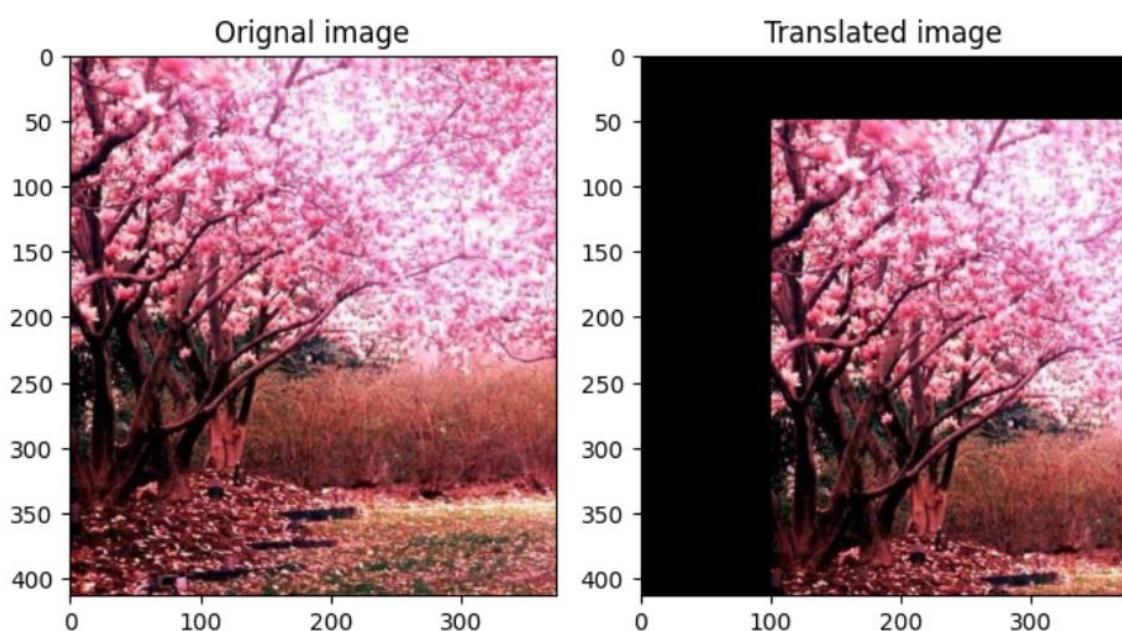
You can take make it into a NumPy array of type `np.float32` and pass it into `cv2.warpAffine()` function. See below example for a shift of (100,50):

### Warning:

Third argument of the `cv2.warpAffine()` function is the size of the output image, which should be in the form of **(width, height)**. Remember width = number of columns, and height = number of rows.

**Code -**

```
import cv2  
  
import matplotlib.pyplot as plt  
  
import numpy as np  
  
img=cv2.imread("/content/cherry blossom 1.jpg")  
img_rqb=cv2.cvtColor(img,cv2.COLOR_BGR2RGB)  
rows,cols,channels=img_rqb.shape  
M=np.float32([[1,0,100],[0,1,50]])  
dst=cv2.warpAffine(img_rqb,M,(cols,rows))  
  
fig,axs=plt.subplots(1,2,figsize=(7,4))  
axs[0].imshow(img_rqb)  
axs[0].set_title('Original image')  
axs[1].imshow(dst)  
axs[1].set_title('Translated image')  
plt.tight_layout()  
plt.show()
```

**Output -**

## Practical – 1(A)

**Aim:** Image Scaling

**Theory:**

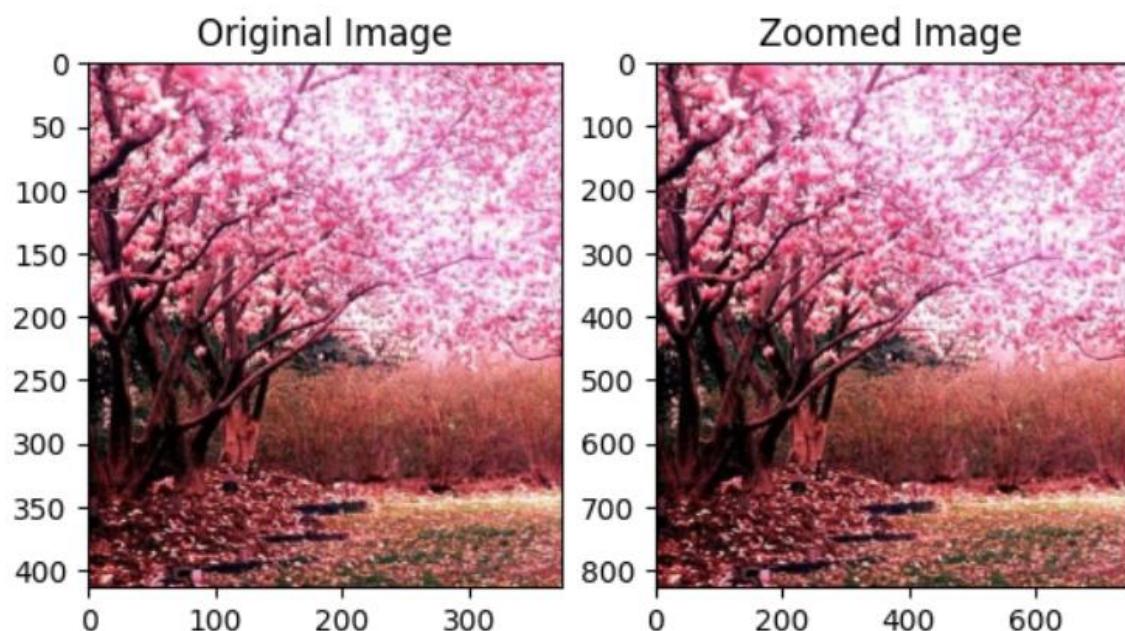
**Scaling:**

Scaling is just resizing of the image. OpenCV comes with a function **cv2.resize()** for this purpose. The size of the image can be specified manually, or you can specify the scaling factor. Different interpolation methods are used. Preferable interpolation methods are **cv2.INTER\_AREA** for shrinking and **cv2.INTER\_CUBIC** (slow) & **cv2.INTER\_LINEAR** for zooming. By default, interpolation method used is **cv2.INTER\_LINEAR** for all resizing purposes. You can resize an input image either of following methods:

**Code –**

```
import cv2  
  
import matplotlib.pyplot as plt  
  
import numpy as np  
  
img=cv2.imread("/content/cherry blossom 1.jpg")  
  
img_rgb=cv2.cvtColor(img,cv2.COLOR_BGR2RGB)  
  
rows, cols, channels=img_rgb.shape  
  
resize_img = cv2.resize(img_rgb,(0,0), fx=2, fy=2, interpolation=cv2.INTER_CUBIC)  
  
plt.subplot(121), plt.imshow(img_rgb), plt.title('Original Image')  
  
plt.subplot(122), plt.imshow(resize_img), plt.title('Zoomed Image')  
  
plt.show()
```

**Output –**



## Practical – 1(B)

**Aim:** Image Shrinking

**Theory:**

**Shrinking:**

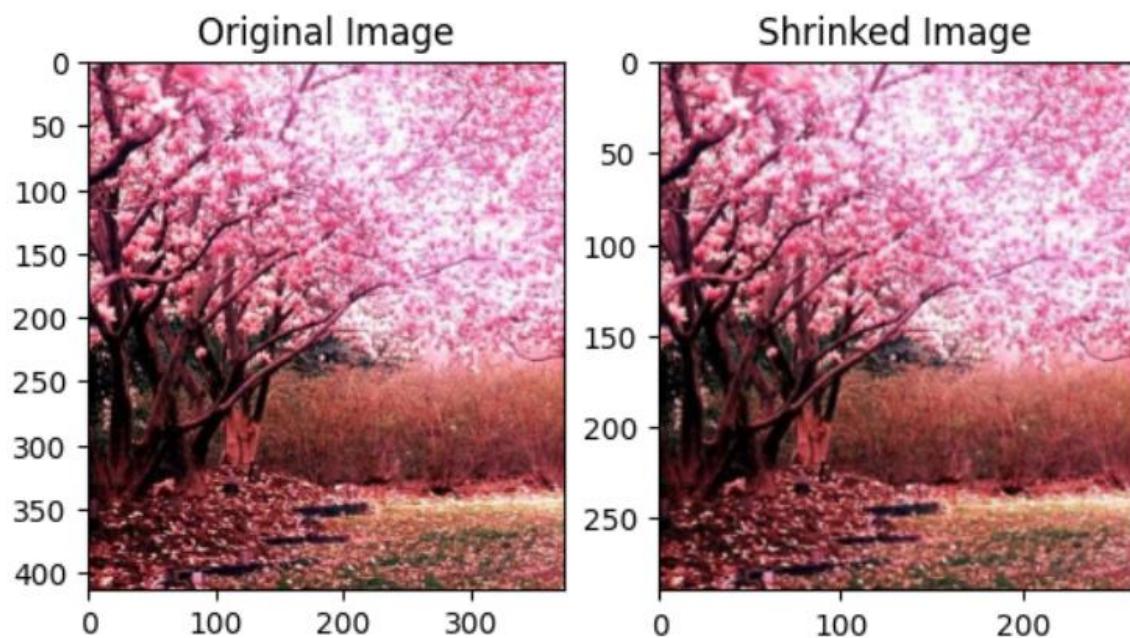
Image shrinking in Python involves reducing the dimensions of an image, effectively making it smaller while maintaining its aspect ratio. This process is commonly done using libraries like PIL (Python Imaging Library) or its fork, Pillow. By resizing the image to smaller dimensions, either by specifying new dimensions or a scaling factor, you can achieve the desired reduction in size.

This is useful for tasks like image compression, thumbnail generation, or preparing images for web display, where smaller file sizes are advantageous. The process typically involves opening the image, calculating the new dimensions, resizing the image accordingly, and then saving the resized image to a new file.

**Code –**

```
import cv2
import matplotlib.pyplot as plt
import numpy as np
img = cv2.imread("/content/cherry blossom 1.jpg")
img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
rows, cols, channels = img_rgb.shape
resize_img = cv2.resize(img_rgb,(0,0), fx=0.7, fy=0.7, interpolation=cv2.INTER_AREA)
plt.subplot(121), plt.imshow(img_rgb), plt.title('Original Image')
plt.subplot(122), plt.imshow(resize_img), plt.title('Shrunked Image')
plt.show()
```

**Output –**



# Practical – 1(C)

**Aim:** Image Rotation

**Theory:**

**Rotation:**

Rotation of an image for an angle  $\theta$  is achieved by the transformation matrix of the form

$$M = \begin{bmatrix} \cos\theta & -\sin\theta \\ \sin\theta & \cos\theta \end{bmatrix}$$

But OpenCV provides scaled rotation with adjustable center of rotation so that you can rotate at any location you prefer. Modified transformation matrix is given by

$$\begin{bmatrix} \alpha & \beta & (1 - \alpha) \cdot center.x - \beta \cdot center.y \\ -\beta & \alpha & \beta \cdot center.x + (1 - \alpha) \cdot center.y \end{bmatrix}$$

*where:*

$$\begin{aligned} \alpha &= scale \cdot \cos\theta, \\ \beta &= scale \cdot \sin\theta \end{aligned}$$

To find this transformation matrix, OpenCV provides a function, **cv2.getRotationMatrix2D**. Check below example which rotates the image by 90 degree with respect to center without any scaling.

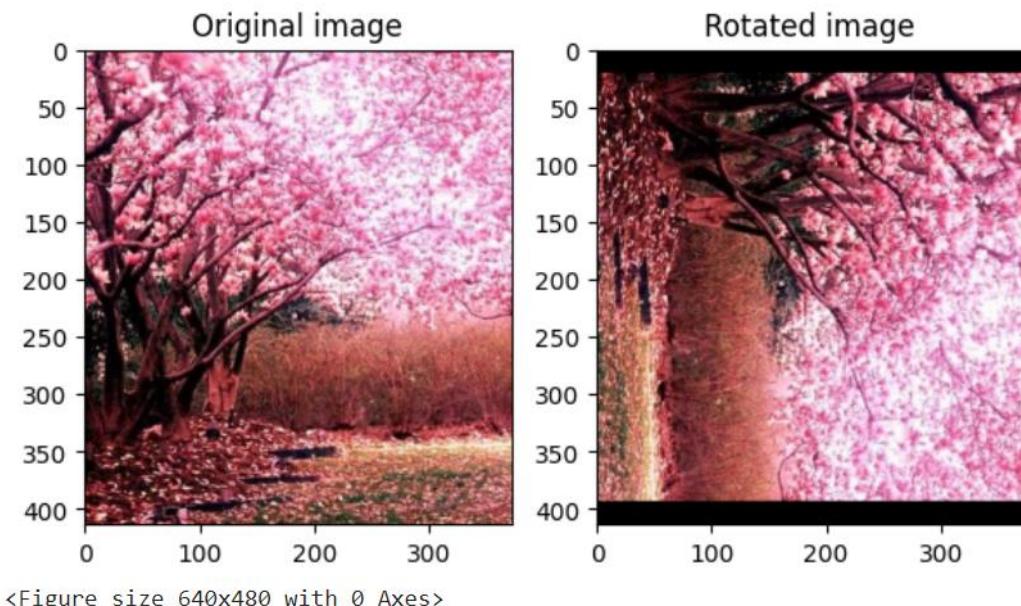
**Code –**

```
import cv2
import matplotlib.pyplot as plt
import numpy as np

img=cv2.imread("/content/cherry blossom 1.jpg")
img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
rows, cols, channels = img_rgb.shape
center = (cols // 2, rows // 2)
angle = -90
scale= 1
```

```
rotation_matrix = cv2.getRotationMatrix2D (center, angle, scale)
rotated_image = cv2.warpAffine(img_rgb, rotation_matrix, (cols,rows))
fig, axs = plt.subplots(1,2, figsize=(7, 4))
axs[0].imshow(img_rgb)
axs[0].set_title("Original image")
axs[1].imshow(rotated_image)
axs[1].set_title("Rotated image")
plt.show()
plt.tight_layout()
```

## Output –



# Practical – 1(D)

**Aim:** Affine Transformation

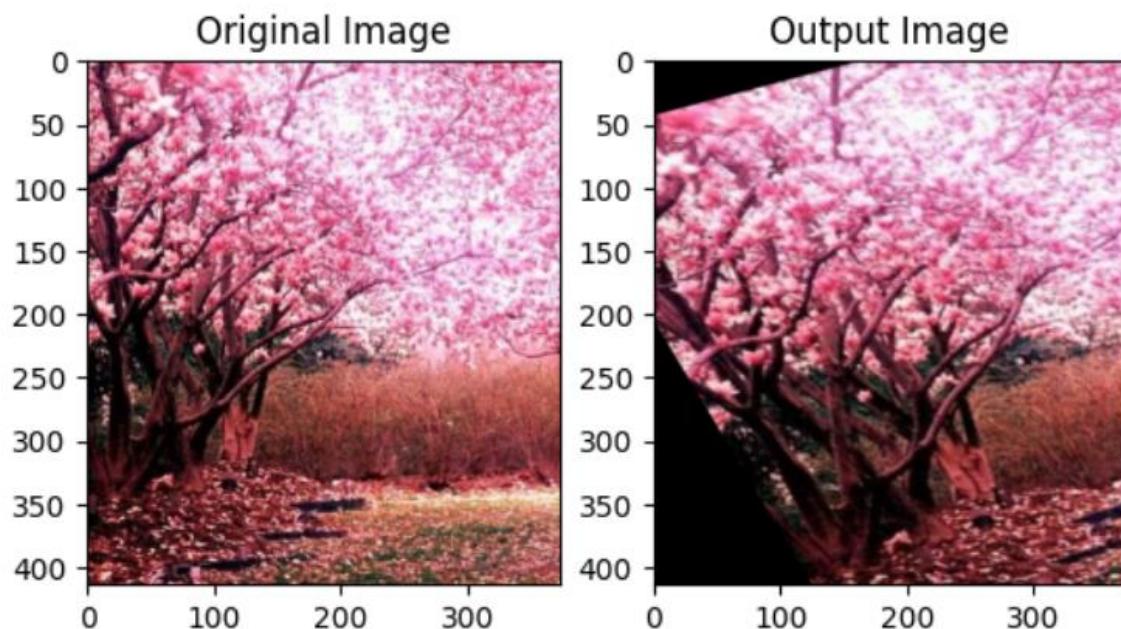
## Theory:

In affine transformation, all parallel lines in the original image will still be parallel in the output image. To find the transformation matrix, we need three points from input image and their corresponding locations in output image. Then **cv2.getAffineTransform** will create a 2x3 matrix which is to be passed to **cv2.warpAffine**.

## Code –

```
import cv2  
import matplotlib.pyplot as plt  
import numpy as np  
  
img= cv2.imread("/content/cherry blossom 1.jpg")  
img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)  
rows, cols, channels = img_rgb.shape  
pts1= np.float32([[50,50], [200,50],[50,200]])  
pts2 =np.float32([[10,100],[200,50],[100,250]])  
M= cv2.getAffineTransform(pts1,pts2)  
dst = cv2.warpAffine(img_rgb,M,(cols,rows))  
plt.subplot(121), plt.imshow(img_rgb), plt.title('Original Image')  
plt.subplot(122), plt.imshow(dst), plt.title('Output Image')  
plt.show()
```

**Output –**



## Practical – 1(E)

**Aim:** Perspective Transformation.

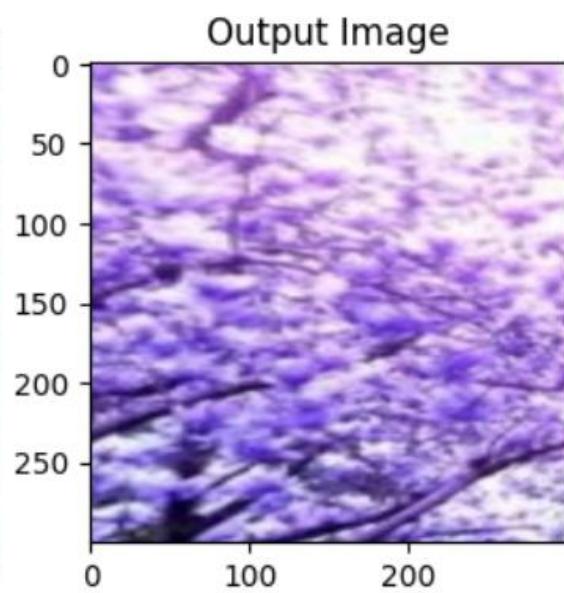
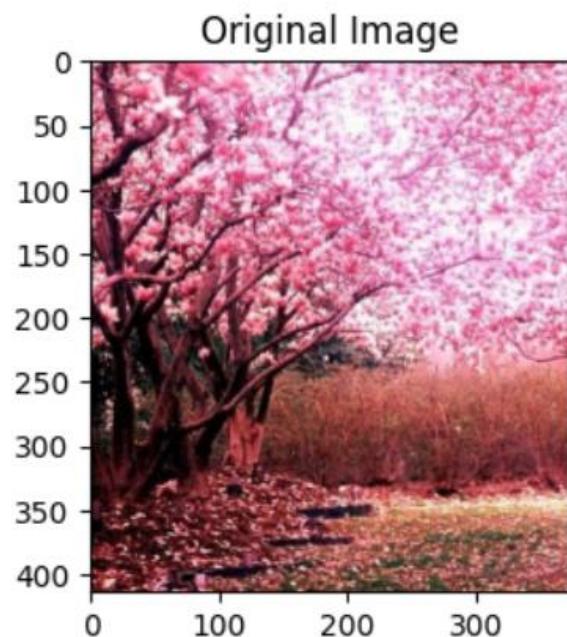
### Theory:

For perspective transformation, you need a 3x3 transformation matrix. Straight lines will remain straight even after the transformation. To find this transformation matrix, you need 4 points on the input image and corresponding points on the output image. Among these 4 points, 3 of them should not be collinear. Then transformation matrix can be found by the function **cv2.getPerspectiveTransform**. Then apply **cv2.warpPerspective** with this 3x3 transformation matrix.

### Code –

```
import cv2  
import matplotlib.pyplot as plt  
import numpy as np  
  
img = cv2.imread("/content/cherry blossom 1.jpg")  
img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)  
rows, cols, channels = img_rgb.shape  
pts1= np.float32([[133,34], [226,16], [133,206], [226,219]])  
pts2= np.float32([[0,0],[300,0], [0,300],[300,300]])  
M = cv2.getPerspectiveTransform(pts1,pts2)  
dst= cv2.warpPerspective (img, M, (300,300))  
plt.subplot(121), plt.imshow(img_rgb), plt.title('Original Image')  
plt.subplot(122), plt.imshow(dst), plt.title('Output Image')  
plt.show()
```

**Output –**



## Practical – 1(F)

**Aim:** Shearing X-axis.

### Theory:

**Shearing** deals with changing the shape and size of the 2D object along x-axis and y- axis. It is similar to sliding the layers in one direction to change the shape of the 2D object. It is an ideal technique to change the shape of an existing object in a two dimensional plane. In a two-dimensional plane, the object size can be changed along X direction as well as Y direction.

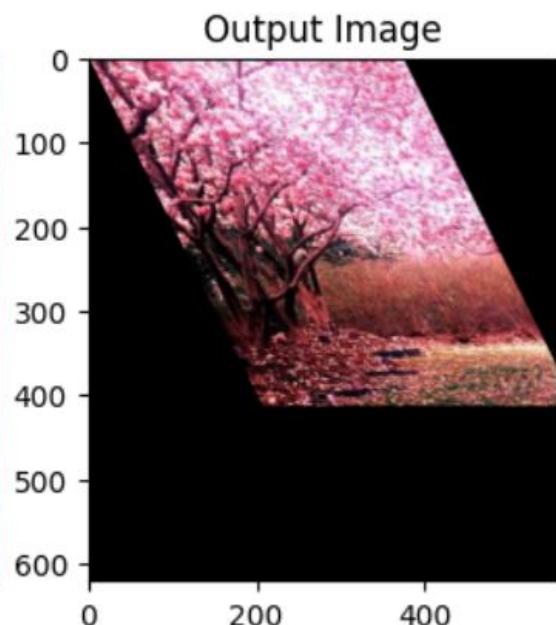
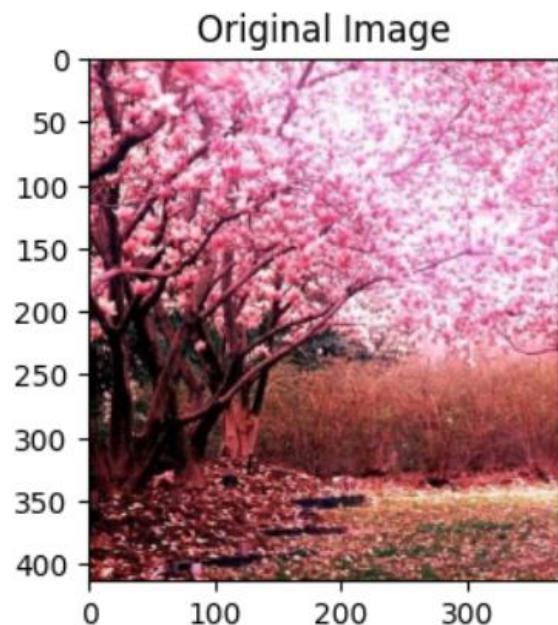
### X-Shear:

In x shear, the y co-ordinates remain the same but the x co-ordinates changes.

### Code-

```
import cv2
import matplotlib.pyplot as plt
import numpy as np
img=cv2.imread("/content/cherry blossom 1.jpg")
img_rgb=cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
rows, cols, channels=img_rgb.shape
M=np.float32([[1, 0.5, 0], [0, 1, 0], [0, 0, 1]])
dst= cv2.warpPerspective(img_rgb, M, (int(cols*1.5), int(rows*1.5)))
plt.subplot(121), plt.imshow(img_rgb), plt.title('Original Image')
plt.subplot(122), plt.imshow(dst),
plt.title('Output Image')
plt.show()
```

**Output –**



## Practical – 1(G)

**Aim:** Shearing Y-axis.

### Theory:

Shearing deals with changing the shape and size of the 2D object along x-axis and y- axis. It is similar to sliding the layers in one direction to change the shape of the 2D object. It is an ideal technique to change the shape of an existing object in a two dimensional plane. In a two-dimensional plane, the object size can be changed along X direction as well as Y direction.

### Y-Shear:

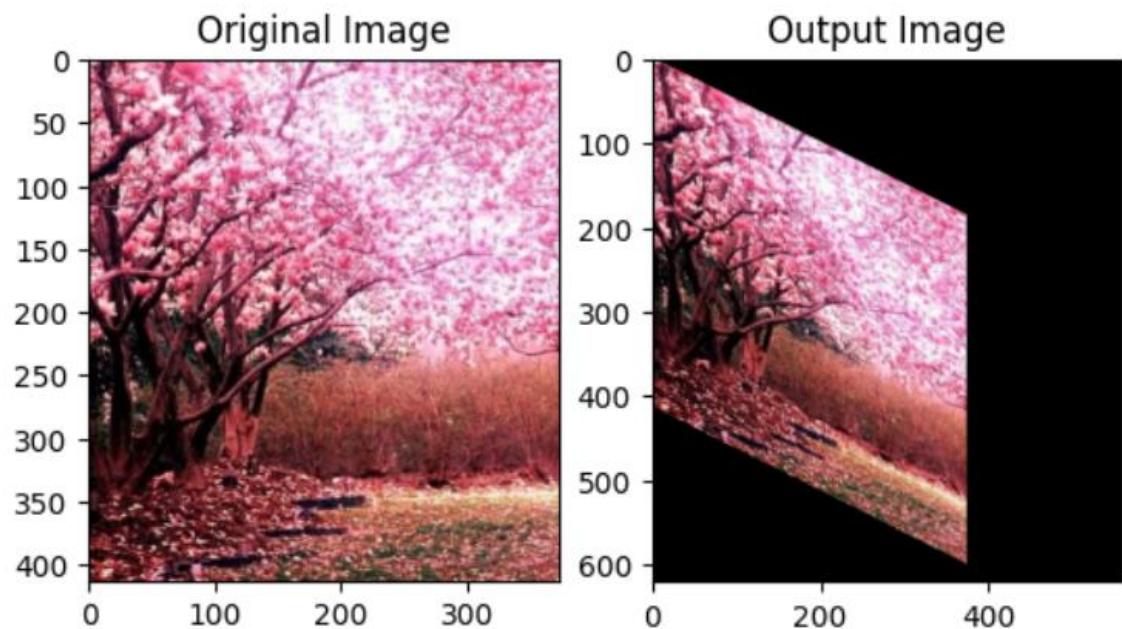
In y shear, the x co-ordinates remain the same but the y co-ordinates changes.

### Code –

```
import cv2
import matplotlib.pyplot as plt
import numpy as np

img=cv2.imread("/content/cherry blossom 1.jpg")
img_rgb=cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
rows, cols, channels=img_rgb.shape
M= np.float32([[1,0,0],[0.5, 1,0],[0, 0,1]])
dst = cv2.warpPerspective(img_rgb, M, (int(cols*1.5), int(rows*1.5)))
plt.subplot(121),
plt.imshow(img_rgb),
plt.title('Original Image')
plt.subplot(122),
plt.imshow(dst),
plt.title('Output Image')
plt.show()
```

**Output –**



## Practical – 1(H)

**Aim:** Reflected Image.

### Theory

#### Image Reflection

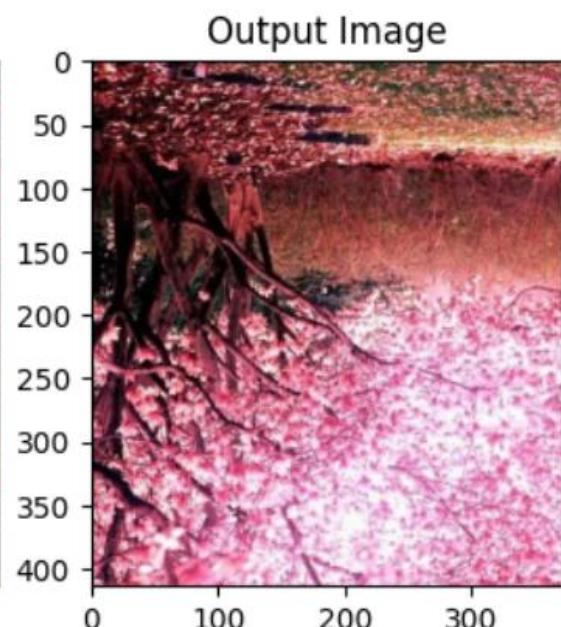
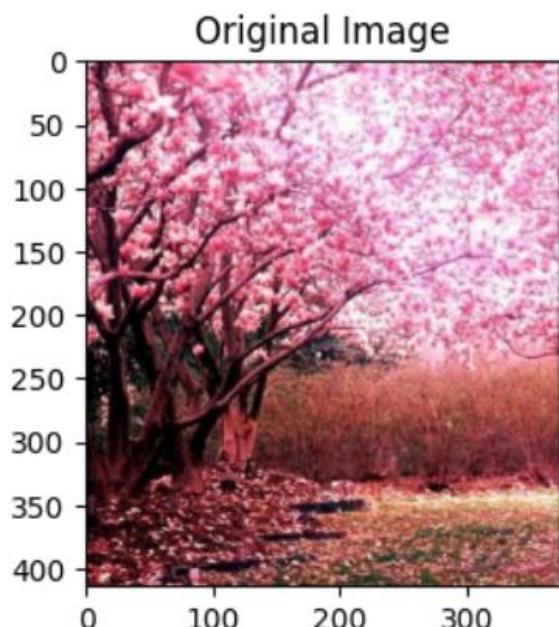
Image reflection is used to flip the image vertically or horizontally. For reflection along the x-axis, we set the value of Sy to -1, Sx to 1, and vice-versa for the y-axis reflection.

#### Code –

```
import cv2
import matplotlib.pyplot as plt
import numpy as np

img=cv2.imread("/content/cherry blossom 1.jpg")
img_rgb=cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
rows,cols, channels=img_rgb.shape
M = np.float32([[1,0,0],[0,-1,rows], [0,0,1]])
dst = cv2.warpPerspective (img_rgb, M,((cols),(rows)))
plt.subplot(121),
plt.imshow(img_rgb),
plt.title('Original Image')
plt.subplot(122),
plt.imshow(dst),
plt.title('Output Image')
plt.show()
```

**Output –**



## Practical – 1(I)

**Aim:** Cropped Image.

**Theory:**

### Image Cropping

Cropping is the removal of unwanted outer areas from an image.

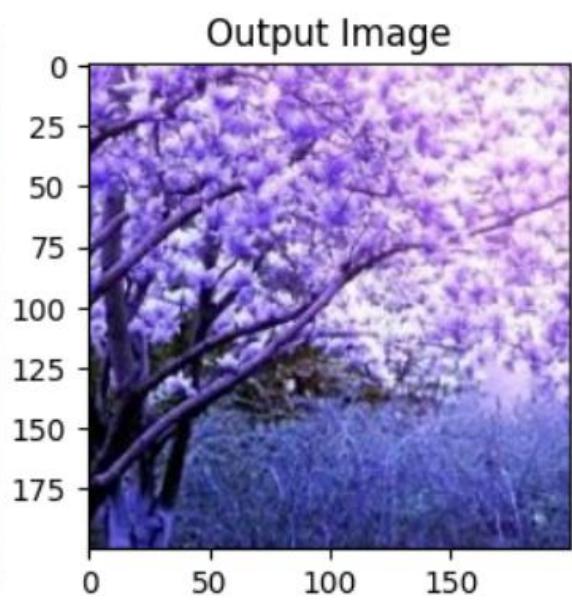
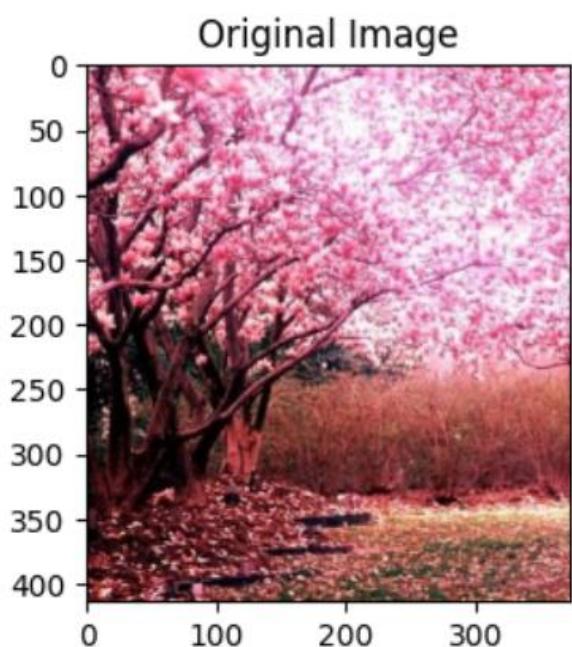
**Code-**

```
import cv2
import matplotlib.pyplot as plt
import numpy as np

img=cv2.imread("/content/cherry blossom 1.jpg")
img_rgb=cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
rows, cols, channels=img_rgb.shape
dst =img[100:300, 100:300]

plt.subplot(121),
plt.imshow(img_rgb),
plt.title('Original Image')
plt.subplot(122),
plt.imshow(dst),
plt.title('Output Image')
plt.show()
```

## Output



## Practical – 2

**Aim:** Perform Image Stitching.

### Theory:

**Image stitching** is the process of combining multiple overlapping images to create a seamless, high-resolution output image. This technique is commonly used to create panoramic images, virtual tours, and even some medical imaging applications.

Image stitching involves several steps:

**Feature detection:** Identifying and extracting unique features (e.g., corners, edges) from each input image. Compute the SIFT-key points and descriptors for both the images.

1. **Feature matching:** Finding correspondences between features in the overlapping regions of the input images. Compute distances between every descriptor in one image and every descriptor in the other image. Select the top  $m$  matches for each descriptor of an image.
2. **Homography estimation:** Estimating the transformation (e.g., rotation, scaling, translation) that aligns the input images. Run RANSAC to estimate homography
3. **Warping:** Applying the estimated transformation to the input images. Warp to align for stitching
4. **Blending:** Combining the warped images into a single seamless output image. Now stitch them together

### Explanation of Code:

Firstly, we have to *find out the features matching in both the images*. These best matched features act as the basis for stitching. *We extract the key points and sift descriptors for both the images as follows:*

```
sift = cv2.SIFT_create()
# find the keypoints and descriptors with SIFT
kp1, des1 = sift.detectAndCompute(img1,None)
kp2, des2 = sift.detectAndCompute(img2,None)
```

kp1 and kp2 are keypoints, des1 and des2 are the descriptors of the respective images. Now, the obtained descriptors in one image are to be recognized in the image too. We do that as follows:

```
bf = cv2.BFMatcher()
matches = bf.knnMatch(des1,des2, k=2)
```

The [BFMatcher\(\)](#) matches the features which are more similar. When we set parameter  $k=2$ , we are asking the knnMatcher to give out 2 best matches for each descriptor. `_matches` is a list of list, where each sub-list consists of `_k` objects.

Often in images, there are tremendous chances where the features may be existing in many places of the image. This may mislead us to use trivial features for our experiment. So we filter out through all the matches to obtain the best ones. So we apply ratio test using the top 2 matches obtained above. We consider a match if the ratio defined below is predominantly greater than the specified ratio.

It's time to align the images now. As you know that a homography matrix is needed to perform  
`if len(matches[:,0]) >= 4:`

```
# src = np.float32([ kp1[m.queryIdx].pt for m in matches[:,0] ]).reshape(-1,1,2)
g dst = np.float32([ kp2[m.trainIdx].pt for m in matches[:,0] ]).reshape(-1,1,2)H,
f masked = cv2.findHomography(src, dst, cv2.RANSAC, 5.0) #print H
i]
0 else:
matches = np.asarray(good)
raise AssertionError("Can't find enough keypoints.")
```

the transformation, and the homography matrix requires at least 4 matches, we do the following.

And finally comes the last part, stitching of the images. Now that we found the homography for transformation, we can now proceed to warp and stitch them together:

```
dst = cv2.warpPerspective(img_,H,(img_.shape[1] + img_.shape[1],
img_.shape[0]))
plt.subplot(122),plt.imshow(dst),plt.title("Warped Image")
plt.show()
plt.figure()
dst[0:img_.shape[0], 0:img_.shape[1]] = img_
cv2.imwrite("output.jpg",dst)
plt.imshow(dst)
plt.show()
```

### **Code-**

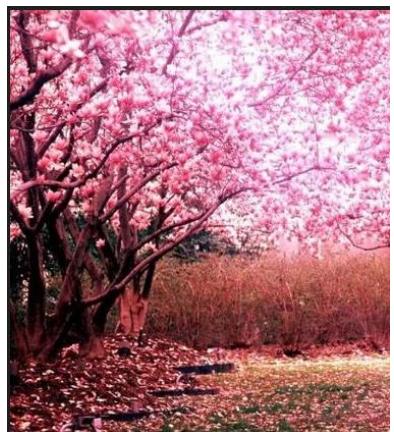
```
import cv2
import numpy as nm
import matplotlib.pyplot as plt
```

```
from random import randrange
img = cv2.imread("/content/cherry blossom 1.jpg")
img1 = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
img = cv2.imread("/content/cherry blossom 2.jpg")
img2 = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
sift = cv2.SIFT_create()
kpl, des1 = sift.detectAndCompute(img1, None)
kp2, des2 = sift.detectAndCompute(img2, None)
bf = cv2.BFM Matcher()
matches = bf.knnMatch(des1, des2, k=2)
good = []
for m in matches:
    if m[0].distance < 0.5 * m[1].distance:
        good.append(m)
matches = np.asarray(good)
if len(matches[:, 0]) >= 4:
    src = np.float32([kpl[m.queryIdx].pt for m in matches[:, 0]]).reshape(-1, 1, 2)
    dst = np.float32([kp2[m.trainIdx].pt for m in matches[:, 0]]).reshape(-1, 1, 2)
    H, masked = cv2.findHomography(src, dst, cv2.RANSAC, 5.0)
    print(H)
else:
    raise AssertionError("Can't find enough keypoints.")
dst = cv2.warpPerspective(img, H, (img.shape[1] + img.shape[1], img.shape[0]))
plt.subplot(122), plt.imshow(dst), plt.title("Warped Image")
plt.show()

plt.figure()
dst[0:img.shape[0], 0:img.shape[1]] = img
cv2.imwrite("resultant_stitched_panorama.jpg", dst)
plt.imshow(dst)
plt.show()
```

**Input Images –**

1

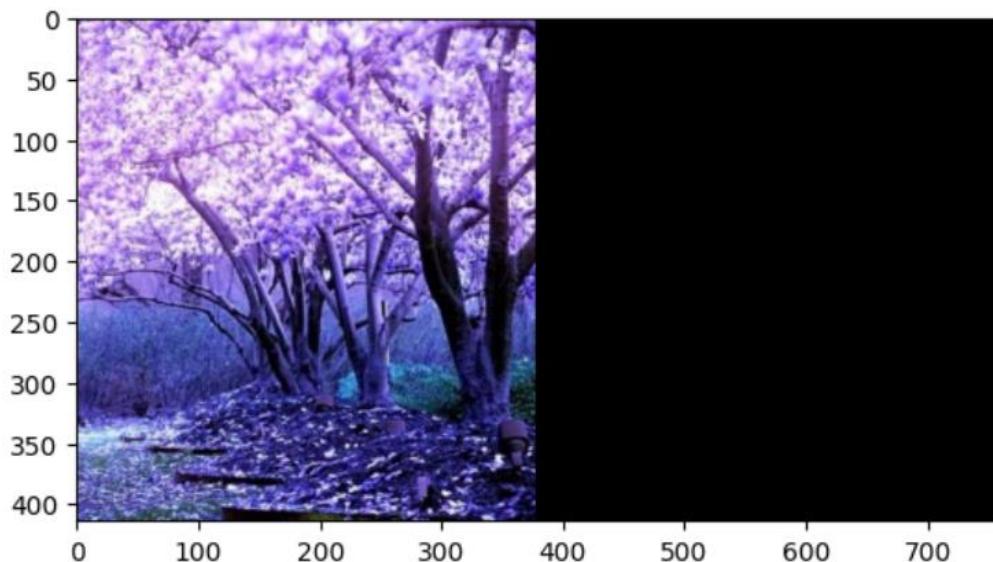
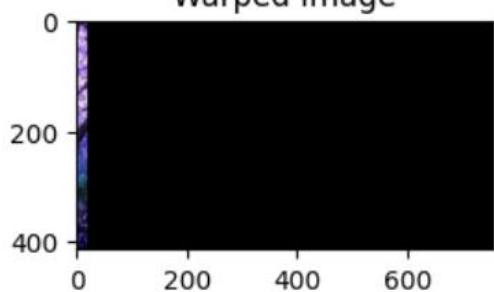


2



**Output –**

Warped Image



## Practical- 3

**Aim:** Perform Camera Calibration.

### Theory:

A camera is an integral part of several domains like robotics, space exploration, etc camera is playing a major role. It helps to capture each and every moment and helpful for many analyses. In order to use the camera as a visual sensor, we should know the parameters of the camera. **Camera Calibration** is nothing but estimating the parameters of a camera, parameters about the camera are required to determine an accurate relationship between a 3D point in the real world and its corresponding 2D projection (pixel) in the image captured by that calibrated camera.

### Code:

```
import numpy as np
import cv2 as cv
import glob
import matplotlib.pyplot as plt
criteria = (cv.TERM_CRITERIA_EPS + cv.TERM_CRITERIA_MAX_ITER, 30, 0.001)
objp = np.zeros((6*7, 3), np.float32)
objp[:, :2] = np.mgrid[0:7, 0:6].T.reshape(-1, 2)
objpoints = []
imgpoints = []

file_images = glob.glob('/content/chess.jpg')
for fname in images:
    img =
        cv.imread(fname) if
    img is None:
```

```
print(f"Failed to load image {fname}")

continue

gray = cv.cvtColor(img, cv.COLOR_BGR2GRAY)

ret, corners = cv.findChessboardCorners(gray, (7, 6),

None) if ret:

    objpoints.append(objp)

    corners2 = cv.cornerSubPix(gray, corners, (11, 11), (-1, -1), criteria)

    imgpoints.append(corners2)

    cv.drawChessboardCorners(img, (7, 6), corners2, ret)

# Display the image with chessboard corners using

matplotlib plt.imshow(cv.cvtColor(img,

cv.COLOR_BGR2RGB)) plt.title(fname)

plt.show()

else:

    print(f"Chessboard corners not found in image {fname}")

empty if len(objpoints) > 0 and len(imgpoints)

> 0:

ret, mtx, dist, rvecs, tvecs =

cv.calibrateCamera(objpoints, imgpoints,

gray.shape[:-1], None, None)

print("Camera matrix:

") print(mtx)

print("Distortion

coefficients:") print(dist)

print("Rotation vectors:")

print(rvecs)

print("Translation

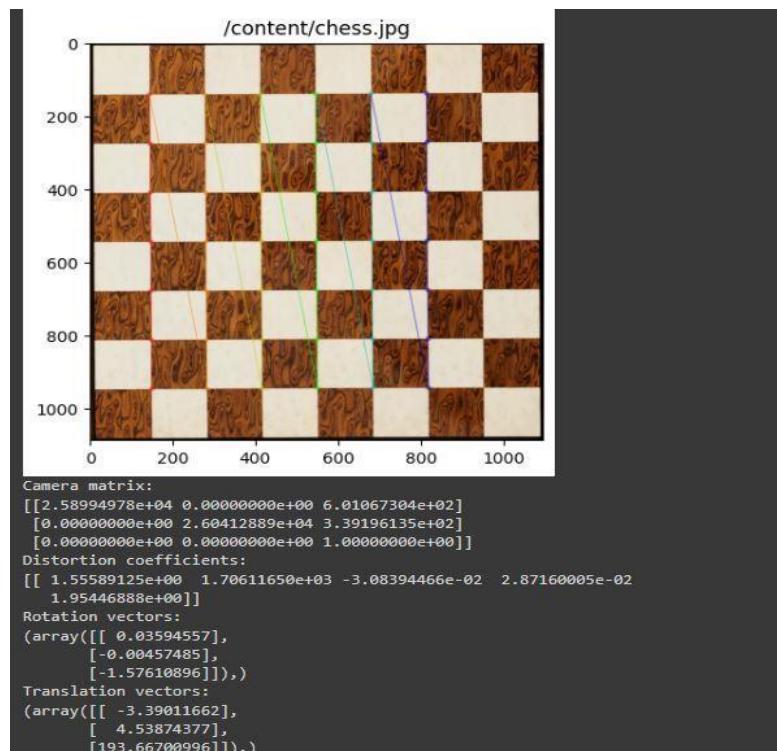
vectors:") print(tvecs)
```

```

img =
cv.imread('/content/left08.jpg') if
img is None:
    print("Failed to load image for undistortion. Please check the
path.") else:
    h, w = img.shape[:2]
    newcameramtx, roi = cv.getOptimalNewCameraMatrix(mtx, dist, (w, h), 1, (w, h))
dst = cv.undistort(img, mtx, dist, None, newcameramtx)
x, y, w, h = roi
dst = dst[y:y+h, x:x+w]
cv.imwrite('/content/calibresult.png', dst)
plt.imshow(cv.cvtColor(dst,
cv.COLOR_BGR2RGB)) plt.title('Undistorted
Image')
plt.show()
else:
print("No chessboard corners were found in any of the images. Calibration cannot be
performed.")

```

## Output:



## Practical 4(A)

**Aim:** Perform the following Face detection.

**Theory:** **Face detection** involves identifying a person's face in an image or video. This is done by analyzing the visual input to determine whether a person's facial features are present. Since human faces are so diverse, face detection models typically need to be trained on large amounts of input data for them to be accurate. The training dataset must contain a sufficient representation of people who come from different backgrounds, genders, and cultures. These algorithms also need to be fed many training samples comprising different lighting, angles, and orientations to make correct predictions in real-world scenarios. These nuances make face detection a non-trivial, time-consuming task that requires hours of model training and millions of data samples.

### face detection

#### Code -

```
import cv2
import matplotlib.pyplot as plt
# Specify the path to the image and cascade classifier
imagePath = '/content/cricket.webp'
classifierPath = '/content/haarcascade_frontalface_default.xml'
img = cv2.imread(imagePath)
correctly if img is None:
    raise ValueError(f"Failed to load image from path:
{imagePath}") else:
    print(f'Image shape: {img.shape}')
gray_image = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
face_classifier = cv2.CascadeClassifier(classifierPath)
```

```
faces = face_classifier.detectMultiScale(gray_image, scaleFactor=1.1, minNeighbors=5,  
minSize=(40, 40))
```

for (x, y, w, h) in faces:

```
    cv2.rectangle(img, (x, y), (x + w, y + h), (0, 255, 0), 4)  
  
matplotlib img_rgb = cv2.cvtColor(img,  
cv2.COLOR_BGR2RGB) plt.figure(figsize=(20, 10))  
plt.imshow(img_rgb)  
plt.axis('off') plt.show()
```

## Output:



## Practical – 4(B-i)

**Aim:** Perform the following Object detection.

### Theory

#### Object Detection

Object detection is a computer technology related to and image processing that deals with detecting instances of semantic objects of a certain class (such as humans, buildings, or cars) in digital images and videos.

**Haar cascade:** Basically, the Haar cascade technique is an approach based on machine learning where we use a lot of positive and negative images to train the classifier to classify between the images. Haar cascade classifiers are considered as the effective way to do object detection with the OpenCV library.

**Positive images:** These are the images that contain the objects which we want to be identified from the classifier.

**Negative Images:** These are the images that do not contain any object that we want to be detected by the classifier, and these can be images of everything else.

**Code:**

```
import cv2

from matplotlib import pyplot as plt

# Load the image
image_path = "/content/images.jfif"
image = cv2.imread(image_path)
if image is None:
    print(f"Failed to load image {image_path}")
else:
    image_gray = cv2.cvtColor(image,
        cv2.COLOR_BGR2GRAY)
    image_rgb =
        cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

# Load the cascade classifier
cascade_path =
    '/content/haarcascade_frontalcatface.xml'
cascade =
    cv2.CascadeClassifier(cascade_path)
if cascade.empty():
    print(f"Failed to load cascade classifier
{cascade_path}") else:
    # Detect objects in the image
    detections = cascade.detectMultiScale(image_gray, minSize=(30,
        30))
    detection_count = len(detections)

    # Draw rectangles around detected
    objects if detection_count != 0:
```

```
for (x, y, width, height) in detections:
```

```
    cv2.rectangle(image_rgb, (x, y), (x + width, y + height), (0, 255, 0), 9)
```

```
# Display the image with detections
```

```
plt.imshow(image_rgb)
```

```
plt.title(f"Detections:
```

```
{detection_count}") plt.axis('off') #
```

```
Hide axis
```

```
plt.show()
```

### Output:



## Practical – 4(B-ii)

**Aim:** Perform the following Line detection.

### Theory:

So, how does the Hough Transform algorithm work? Let's dive into the details.

Each point in the image is transformed into a line in the parameter space. This conversion allows us to represent lines in a different coordinate system, making line detection more feasible. Think of it as a translation from the image domain to a space where lines become easier to analyze.

The **Hough Transform algorithm** is based on the following assumptions:

- Lines in an image can be represented by their equation in polar coordinates, which is given by:  $r = x\cos(\theta) + y\sin(\theta)$ , where  $(r, \theta)$  are the parameters of the line, and  $(x, y)$  are the coordinates of a point on the line.
- A point in an image corresponds to a sinusoidal curve in the parameter space.
- The intersection of curves in the parameter space corresponds to a line in the image.

By considering these assumptions, we can establish that a point in an image corresponds to a sinusoidal curve in the parameter space. As we detect more points, these curves intersect and accumulate at certain locations. These intersections indicate the presence of lines in the image.

## Line Detection Algorithm using OpenCV Python

Now that we have a fundamental grasp of the Hough Transform and Probabilistic Hough Transform algorithm, let's explore their implementation using OpenCV and Python.

To identify lines in an image using OpenCV in Python, we can proceed with the following steps:

1. Import Required Libraries.
2. Load the image.
3. Convert the image to grayscale.
4. Apply Canny Edge Detection to get the edges.
5. Apply Hough Line Transform to detect the lines.
6. Draw the detected lines on the image.

### Line Detection

## Convert the Image to Grayscale

The Hough Transform algorithm works on grayscale images, so we need to convert the image to grayscale using the `cv2.cvtColor()` function.

```
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
```

## Apply Edge Detection

Next, we need to apply an [edge detection](#) algorithm to the grayscale image. In this example, we will use the Canny edge detection algorithm, which is a popular edge detection algorithm that is based on the gradient of the image.

```
edges = cv2.Canny(gray, 50, 150, apertureSize=3)
```

## Apply the Hough Transform Algorithm

OpenCV provides several algorithms for line detection, including **HoughLines** and **HoughLinesP**.

**HoughLines** is a standard Hough transform algorithm that returns an array of lines in the image. **HoughLinesP** is a probabilistic Hough transform algorithm that is faster and more accurate than HoughLines. It returns an array of line segments instead of full lines.

Now, we can apply the Hough Transform algorithm using the `cv2.HoughLines()` or `cv2.HoughLinesP` function.

The function takes the edge image as input and returns a list of lines detected in the image. The parameters of the function are explained below:

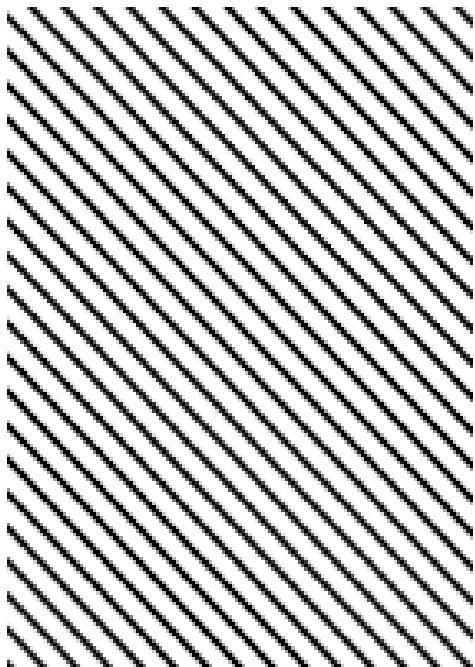
- **rho**: The distance resolution in pixels of the Hough grid.
- **theta**: The angular resolution in radians of the Hough grid.
- **threshold**: The minimum number of votes (intersections in the Hough grid) required to detect a line.
- **minLineLength**: The minimum length of a line. Lines shorter than this will be discarded.
- **maxLineGap**: The maximum allowed gap between two points on the same line.

**Code:**

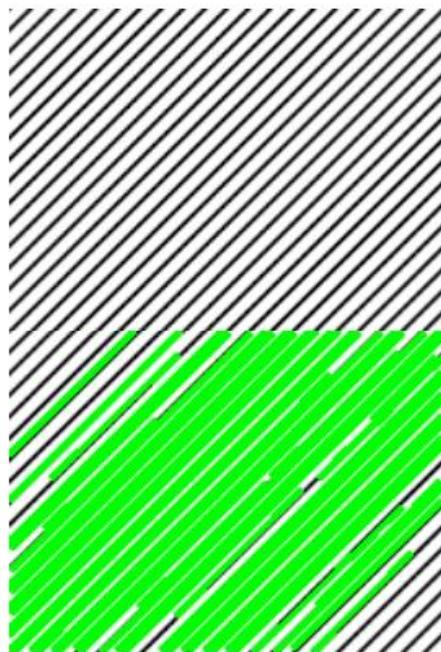
```
import cv2
import numpy as np
from google.colab.patches import cv2_imshow
image = cv2.imread('/content/lines.png')
cv2_imshow(image)
line_list = []
gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
edges = cv2.Canny(gray, 50, 150, apertureSize=3)
lines = cv2.HoughLinesP(
    edges,
    rho=2,
    theta=np.pi / 180,
    threshold=100,
    minLineLength=20,
    maxLineGap=5
)
if lines is not None:
    for points in lines:
        x1, y1, x2, y2 = points[0]
        cv2.line(image, (x1, y1), (x2, y2), (0, 255, 0), 2)
        line_list.append([(x1, y1), (x2, y2)])
cv2_imshow(image)
```

**Output :**

**Input image**



**Output image**



## Practical – 4(B-iii)

**Aim:** Perform the following Hough transform.

**Theory:**

**Hough transform**

- Finds circles in a grayscale image using the Hough transform.
- The function finds circles in a grayscale image using a modification of

```
void cv::HoughCircles ( InputArray image,
                      OutputArray circles,
                      int method,
                      double dp,
                      double minDist,
                      double param1 = 100 ,
                      double param2 = 100 ,
                      int minRadius = 0 ,
                      int maxRadius = 0
                    )
```

the Hough transform.

- Usually the function detects the centers of circles well. However, it may fail to find correct radii.
- You can assist to the function by specifying the radius range ( minRadius and maxRadius ) if you know it.
- Or, in the case of HOUGH\_GRADIENT method you may set maxRadius to a negative number to return centers only without radius search, and find the correct radius using an additional procedure.
- It also helps to smooth image a bit unless it's already soft. For example, GaussianBlur() with 7x7 kernel and 1.5x1.5 sigma or similar blurring may help.

What is Hough accumulator?

- The Hough transform algorithm uses an array, called an accumulator, to detect the existence of a line  $y = mx + b$ . The dimension of the accumulator is equal to the number of unknown parameters of the Hough transform problem.

---

## Overview of np.around()

---

Basically, `np.around()` is NumPy's function to round floats (decimal numbers) or integers. It takes three arguments:

1. The first mandatory argument is the number you want to round
2. `decimals` - optional argument, the number of decimal places you want to round to
3. `out` - optional argument, where to output

By default `np.around()` will round the given number to the nearest whole number and return the rounded value:

```
>>> np.around(5)
5

>>> np.around(5.11212)
5.0

>>> np.around(9.11212)
9.0

>>> np.around(9.5)
10.0
```

Website uses cookies to ensure you get the best experience on our website. [Learn more](#)

---

**OpenCV-Python** is a library of Python bindings designed to solve computer vision problems. `cv2.circle()` method is used to draw a circle on any image. The syntax of `cv2.circle()` method is:

**Syntax:**

```
cv2.circle(image, center_coordinates, radius, color, thickness)
```

**Parameters:**

- **image:** It is the image on which the circle is to be drawn.
- **center\_coordinates:** It is the center coordinates of the circle. The coordinates are represented as tuples of two values i.e. (X coordinate value, Y coordinate value).
- **radius:** It is the radius of the circle.
- **color:** It is the color of the borderline of a circle to be drawn. For **BGR**, we pass a tuple. eg: (255, 0, 0) for blue color.
- **thickness:** It is the thickness of the circle border line in **px**. Thickness of **-1 px** will fill the circle shape by the specified color.

**Return Value:** It returns an image.

**Code:**

```
import cv2
import numpy as np
from google.colab.patches import cv2_imshow

# Read the image
img = cv2.imread('/content/circle.jpeg')

# Display the original image
cv2_imshow(img)

# Convert to grayscale
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

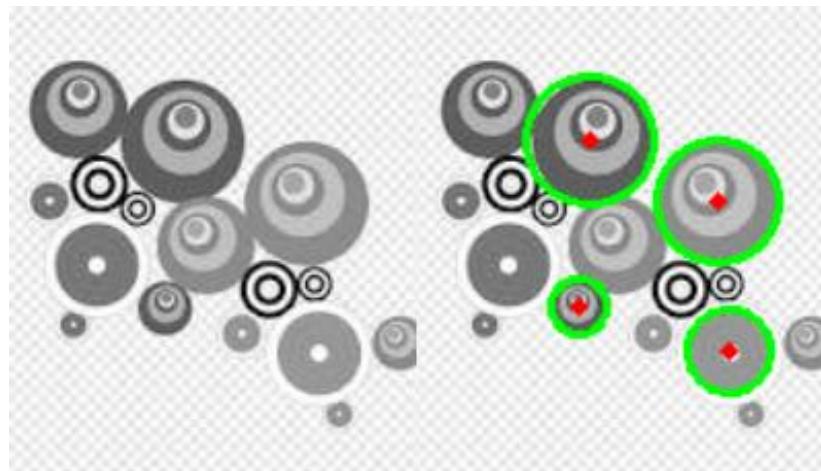
# Apply a blur to the grayscale image
gray_blurred = cv2.blur(gray, (3, 3))

# Detect circles using HoughCircles
detected_circles = cv2.HoughCircles(
    gray_blurred,
    cv2.HOUGH_GRADIENT,
    1,
    20,
    param1=100,
    param2=50,
    minRadius=2,
    maxRadius=80
)

# Draw detected circles on the image
if detected_circles is not None:
    detected_circles = np.uint16(np.around(detected_circles))
    for pt in detected_circles[0, :]:
        a, b, r = pt[0], pt[1], pt[2]
        cv2.circle(img, (a, b), r, (0, 255, 0), 2)
        cv2.circle(img, (a, b), 1, (0, 0, 255), 3)

# Display the image with detected circles
cv2_imshow(img)
```

## Output



## Practical – 4(C)

**Aim:** Perform the following Pedestrian detection.

### Theory:

#### Pedestrian detection

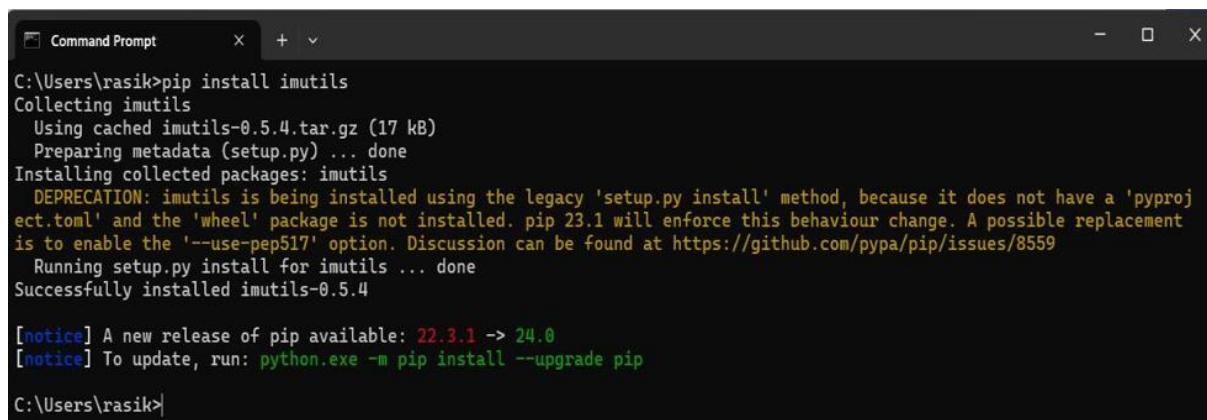
Pedestrian detection is a very important area of research because it can enhance the functionality of a pedestrian protection system in Self Driving Cars. We can extract features like head, two arms, two legs, etc, from an image of a human body and pass them to train a machine learning model. After training, the model can be used to detect and track humans in images and video streams. However, OpenCV has a built-in method to detect pedestrians. It has a pre-trained HOG(Histogram of Oriented Gradients) + Linear SVM model to detect pedestrians in images and video streams.

#### Histogram of Oriented Gradients

This algorithm checks directly surrounding pixels of every single pixel. The goal is to check how darker is the current pixel compared to the surrounding pixels. The algorithm draws and arrows showing the direction of the image getting darker. It repeats the process for each and every pixel in the image. At last, every pixel would be replaced by an arrow, these arrows are called Gradients. These gradients show the flow of light from light to dark. By using these gradients algorithms perform further analysis.

### Requirements

1. opencv-python
2. imutils



```
C:\Users\rasik>pip install imutils
Collecting imutils
  Using cached imutils-0.5.4.tar.gz (17 kB)
    Preparing metadata (setup.py) ... done
Installing collected packages: imutils
  DEPRECATION: imutils is being installed using the legacy 'setup.py install' method, because it does not have a 'pyproject.toml' and the 'wheel' package is not installed. pip 23.1 will enforce this behaviour change. A possible replacement is to enable the '--use-pep517' option. Discussion can be found at https://github.com/pypa/pip/issues/8559
    Running setup.py install for imutils ... done
Successfully installed imutils-0.5.4

[notice] A new release of pip available: 22.3.1 -> 24.0
[notice] To update, run: python.exe -m pip install --upgrade pip
C:\Users\rasik>
```

**Code –**

```
import cv2  
import imutils  
from google.colab.patches import cv2_imshow  
hog = cv2.HOGDescriptor()  
hog.setSVMDetector(cv2.HOGDescriptor_getDefaultPeopleDetector())  
image = cv2.imread('/content/people_walking.jpg')  
image = imutils.resize(image, width=min(400, image.shape[1]))  
(regions, _) = hog.detectMultiScale(image, winStride=(4, 4), padding=(4, 4), scale=1.05)  
for (x, y, w, h) in regions:  
    cv2.rectangle(image, (x, y), (x + w, y + h), (0, 0, 255), 2)  
cv2_imshow(image)
```

**Output :**

## Practical – 4(D)

**Aim:** Perform the following Face Recognition.

### Theory:

Face recognition is different from face detection. In face detection, we had only detected the location of human faces, and we recognized the identity of faces in the face recognition task. In this article, we are going to build a face recognition system using python with the help of face recognition library.

There are many algorithms available in the market for face recognition. This broad challenge is detecting faces from videos and pictures. Many applications can be built on top of recognition systems. Many big companies are adopting recognition systems for their security and authentication purposes.

### Use Cases of Recognition Systems

Face recognition systems are widely used in the modern era, and many new innovative systems are built on top of recognition systems.

#### There are a few used cases :

- Finding Missing Person
- Identifying accounts on social media
- Recognizing Drivers in Cars
- School Attendance System

Several methods and algorithms implement facial recognition systems depending on the performance and accuracy.

### Traditional Face Recognition Algorithm

Traditional face recognition algorithms don't meet modern-day's facial recognition standards. They were designed to recognize faces using old conventional algorithms.

OpenCV provides some traditional facial Recognition Algorithms.

#### Eigenfaces

#### Scale Invariant Feature Transform (SIFT) Fisher faces

#### Local Binary Patterns Histograms (LBPH)

These methods differ in the way they extract image information and match input and output images.

LBPH algorithm is a simple yet very efficient method still in use but it's slow compared to modern days algorithms.

## Deep Learning For Face Recognition

There are various deep learning-based facial recognition algorithms available.

- DeepFace
- DeepID series of systems,
- FaceNet
- VGGFace

Generally, face recognizers that are based on landmarks take face images and try to find essential feature points such as eyebrows, corners of the mouth, eyes, nose, lips, etc. There are more than 60 points.

## Steps Involved in Face Recognition

1. **Face Detection:** Locate the face, note the coordinates of each face located and draw a bounding box around every faces.
2. **Face Alignments.** Normalize the faces in order to attain fast training.
3. **Feature Extraction.** Local feature extraction from facial pictures for training, this step is performed differently by different algorithms.
4. **Face Recognition.** Match the input face with one or more known faces in our dataset.

## Code :

```
import cv2
from google.colab.patches import cv2_imshow
img = cv2.imread('/content/image.jpg')
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
face_cascade = cv2.CascadeClassifier(cv2.data.haarcascades +
'haarcascade_frontalface_default.xml')
faces = face_cascade.detectMultiScale(gray, scaleFactor=1.1, minNeighbors=5)
for (x, y, w, h) in faces:
    cv2.rectangle(img, (x, y), (x+w, y+h), (255, 0, 255), 2)
    font = cv2.FONT_HERSHEY_SIMPLEX
    text = 'I am Manasi Singh'
    font_scale = 1
```

```
font_thickness = 1  
text_size, _ = cv2.getTextSize(text, font, font_scale, font_thickness)  
text_x = 20  
text_y = 30 + text_size[1]  
cv2.putText(img, text, (text_x, text_y), font, font_scale, (0, 0, 0), font_thickness,  
cv2.LINE_AA)  
scale_percent = 20  
width = int(img.shape[1] * scale_percent / 100)  
height = int(img.shape[0] * scale_percent / 100)  
dim = (width, height)  
resized_img = cv2.resize(img, dim, interpolation=cv2.INTER_AREA)  
cv2_imshow(resized_img)
```

## Output



## Practical – 5(A)

**Aim:** Implement object detection and tracking from video.

**Theory:**

**Object detection** is the detection on every single frame and frame after frame.

**Object tracking** does frame-by-frame tracking but keeps the history of where the object is at a time after time

### 1. Importing Libraries and Modules:

```
import cv2:
```

Imports the OpenCV library used for tasks. from tracker import \*:

Imports all functions and classes from a tracker module, which likely contains the implementation of the EuclideanDistTracker.

### 2. Creating Objects and Initializing Video Capture:

```
tracker = EuclideanDistTracker():
```

Instantiates the Euclidean Distance Tracker. cap = cv2.VideoCapture("highway.mp4"):

Initializes video capture with the video file "highway.mp4".

### 3. Object Detection Initialization:

```
object_detector = cv2.createBackgroundSubtractorMOG2(history=100, varThreshold=40):
```

Initializes a background subtractor with MOG2 method to differentiate between foreground (moving objects) and the background.

Background subtraction (BS) is a common and widely used technique for generating a foreground mask (namely, a binary image containing the pixels belonging to moving objects in the scene) by using static cameras.

As the name suggests, BS calculates the foreground mask performing a subtraction between the current frame and a background model, containing the static part of the scene or, more in general, everything that can be considered as background given the characteristics of the observed scene.

#### 1. Processing Video Frames:

The while True loop starts an infinite loop to process video frames until manually stopped.

ret, frame = cap.read(): Reads the next frame from the video.

height, width, \_ = frame.shape: Retrieves the dimensions of the frame.

## 2. Defining Region of Interest (ROI):

roi = frame[340: 720, 500: 800]: Defines a specific area in the video frame to focus the object detection on. This reduces computation and ignores irrelevant areas.

## 3. Object Detection:

mask = object\_detector.apply(roi): Applies the background subtractor to the ROI to get the foreground mask.

\_ , mask = cv2.threshold(mask, 254, 255, cv2.THRESH\_BINARY): Applies a threshold to the mask to make it binary, which helps in identifying distinct

objects.contours, \_ = cv2.findContours(mask, cv2.RETR\_TREE,  
cv2.CHAIN\_APPROX\_SIMPLE):

Finds the contours of the detected objects in the binary mask.

## 4. Filtering and Storing Detections:

It iterates through each contour, calculates its area, and if the area is larger than a threshold (100), it calculates a bounding box for the object. These bounding boxes ([x, y, w, h]) are added to the detections list representing detected objects.

## 5. Object Tracking:

boxes\_ids = tracker.update(detections): The tracker updates with the current frame's detections and returns the tracked objects with their IDs. The loop then iterates through these tracked objects, drawing their ID and bounding box on the ROI.

### Code :

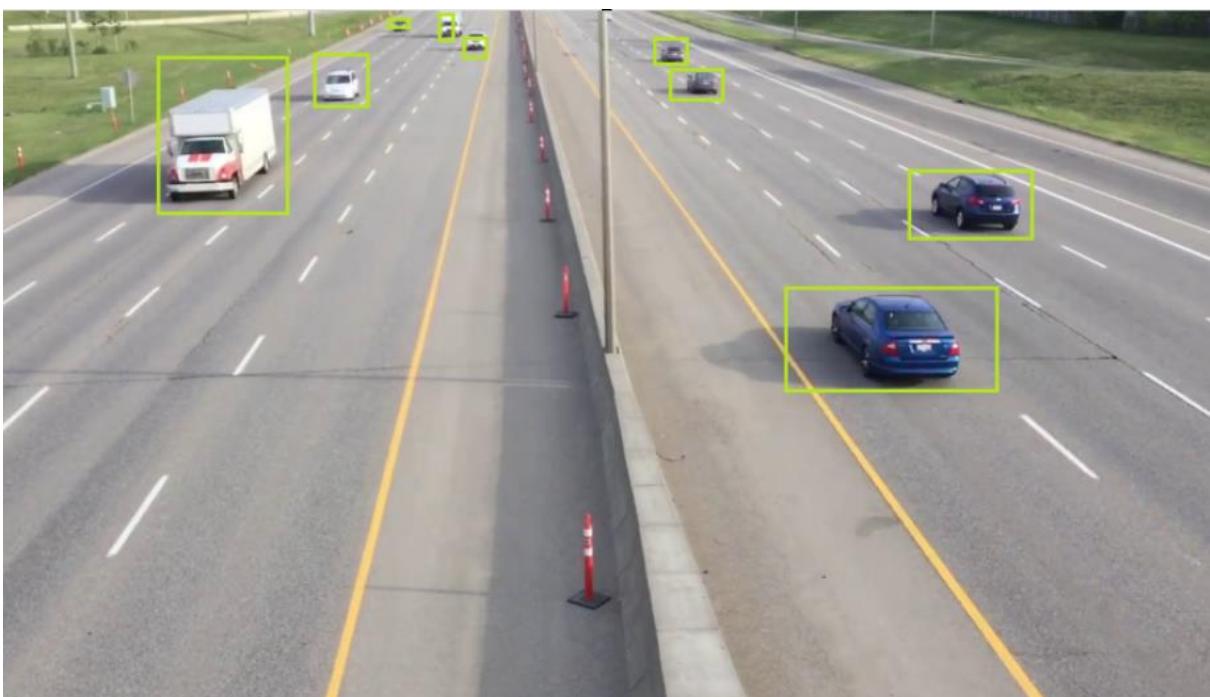
```
import cv2
from google.colab.patches import cv2_imshow
from tracker import *
tracker = EuclideanDistTracker()
cap = cv2.VideoCapture("/content/highway_video.mp4")
object_detector = cv2.createBackgroundSubtractorMOG2(history=100, varThreshold=40)
while True:
    ret, frame = cap.read()
```

```
if not ret:  
    print("End of video or cannot read the frame.")  
  
    break # Exit the loop if the video ends or if there's an issue  
height, width, channels = frame.shape # Unpack all three values  
print(height, width)  
roi = frame[340:720, 500:800]  
mask = object_detector.apply(roi)  
_, mask = cv2.threshold(mask, 254, 255, cv2.THRESH_BINARY)  
contours, _ = cv2.findContours(mask, cv2.RETR_TREE,  
cv2.CHAIN_APPROX_SIMPLE)  
detections = []  
for cnt in contours:  
    area = cv2.contourArea(cnt)  
    if area > 100:  
        x, y, w, h = cv2.boundingRect(cnt)  
        cv2.rectangle(roi, (x, y), (x + w, y + h), (0, 255, 0), 2)  
        detections.append([x, y, w, h])  
boxes_ids = tracker.update(detections)  
for box_id in boxes_ids:  
    x, y, w, h, id = box_id  
    cv2.putText(roi, str(id), (x, y - 15), cv2.FONT_HERSHEY_PLAIN, 1, (255, 0, 0), 2)  
    cv2.rectangle(roi, (x, y), (x + w, y + h), (0, 255, 0), 3)  
# Display images using cv2_imshow  
cv2_imshow(roi)  
cv2_imshow(frame)  
cv2_imshow(mask)  
key = cv2.waitKey(30)  
if key == 27:  
    break
```

```
cap.release()
```

```
cv2.destroyAllWindows()
```

## Output



## Practical – 5(B)

**Aim:** Implement object detection and tracking from video. (Count number of Faces using Python)

### Theory:

#### Count number of Faces using Python

We will use image processing to detect and count the number of faces. We are not supposed to get all the features of the face. Instead, the objective is to obtain the bounding box through some methods i.e. coordinates of the face in the image, depending on different areas covered by the number of the coordinates, number faces that will be computed.

#### Required libraries:

[OpenCV](#) library in python is a library, mostly used for image processing, video processing, and analysis, facial recognition and detection, etc.

- **Dlib** library in python contains the pre-trained facial landmark detector, that is used to detect the (x, y) coordinates that map to facial structures on the face.

[Numpy](#) is a general-purpose array-processing package. It provides a high-performance multidimensional array object and tools for working with these arrays

#### Code :

```
import cv2
import numpy as np
import dlib
from google.colab.patches import cv2_imshow
cap = cv2.VideoCapture("/content/people.mp4")
detector = dlib.get_frontal_face_detector()
while True:
    ret, frame = cap.read()
    if not ret:
        break
    frame = cv2.flip(frame, 1)
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
```

```
faces = detector(gray)

for i, face in enumerate(faces):
    x, y = face.left(), face.top()
    x1, y1 = face.right(), face.bottom()
    cv2.rectangle(frame, (x, y), (x1, y1), (0, 255, 0), 2)
    cv2.putText(frame, 'face num' + str(i + 1), (x - 10, y - 10),
    cv2.FONT_HERSHEY_SIMPLEX, 0.7, (0, 0, 255), 3)
    print(face, i + 1)

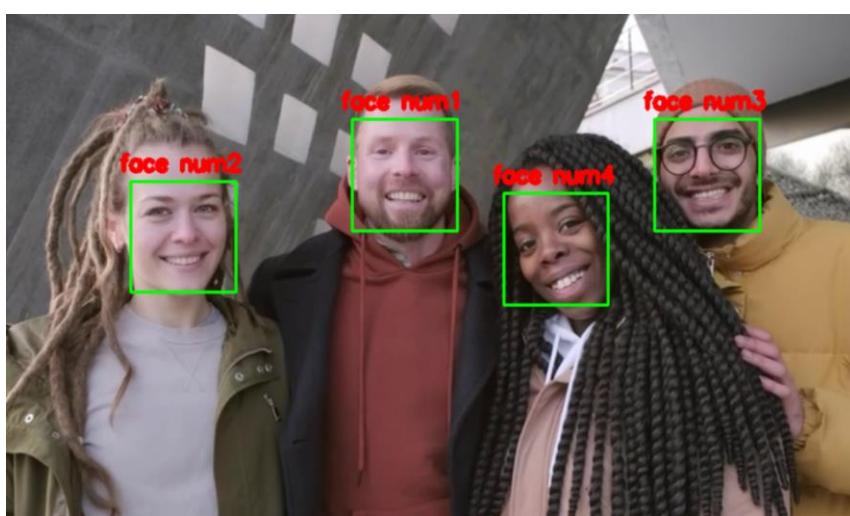
frame = cv2.resize(frame, (1000, 600))
cv2_imshow(frame)

if cv2.waitKey(1) & 0xFF == ord('q'):
    break

cap.release()
cv2.destroyAllWindows()
```

## Output

```
[(146, 146) (250, 250)] 1
[(362, 102) (448, 189)] 2
[(492, 169) (595, 273)] 3
[(630, 89) (734, 192)] 4
```



## Practical – 5(C)

**Aim:** Implement object detection and tracking from video. (Object Tracking using Homography)

**Theory:** Homography is a transformation that maps the points in one point to the corresponding point in another image.

The homography is a  $3 \times 3$  matrix :

$$\mathbf{H} = \begin{vmatrix} h_{00} & h_{01} & h_{02} \\ h_{10} & h_{11} & h_{12} \\ h_{20} & h_{21} & h_{22} \end{vmatrix}$$

If 2 points are not in the same plane then we have to use 2 homographs. Similarly, for n planes, we have to use n homographs. If we have more homographs then we need to handle all of them properly. So that is why we use feature matching.

### Code:

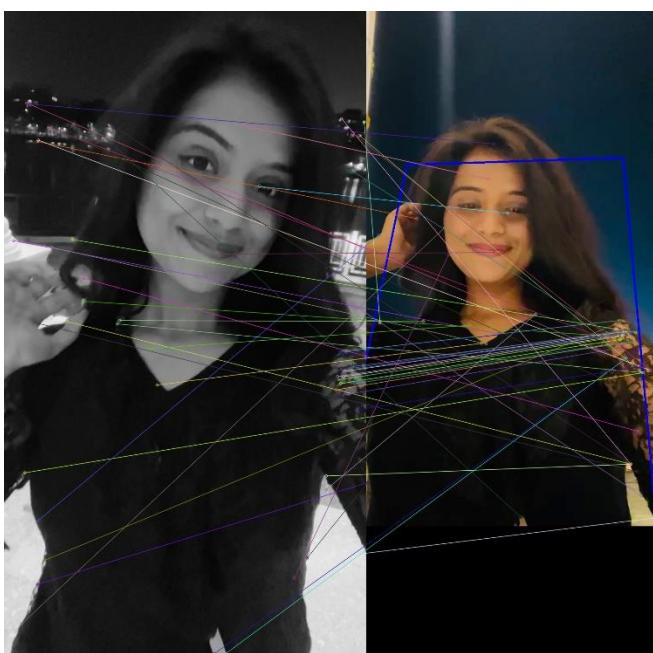
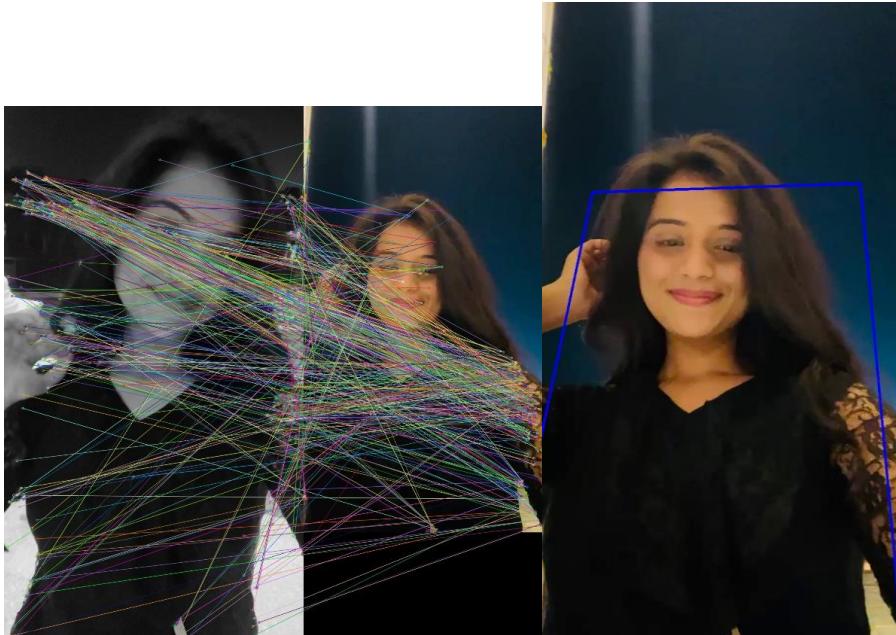
```

import cv2
import numpy as np
from google.colab.patches import cv2_imshow
img = cv2.imread("/content/fd_img.jpg", cv2.IMREAD_GRAYSCALE)
cap = cv2.VideoCapture("/content/fd.mp4")
sift = cv2.SIFT_create()
kp_image, desc_image = sift.detectAndCompute(img, None)
index_params = dict(algorithm=0, trees=5)
search_params = dict()
flann = cv2.FlannBasedMatcher(index_params, search_params)
ret, frame = cap.read()
if not ret:
    print("Failed to read video")
else:

```

```
grayframe = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
kp_grayframe, desc_grayframe = sift.detectAndCompute(grayframe, None)
matches = flann.knnMatch(desc_image, desc_grayframe, k=2)
good_points = []
for m, n in matches:
    if m.distance < 0.8 * n.distance: # Slightly higher ratio
        good_points.append(m)
all_matches = cv2.drawMatches(img, kp_image, frame, kp_grayframe, [m[0] for m in matches], None, flags=cv2.DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)
cv2_imshow(all_matches)
if len(good_points) > 10:
    # Extract location of good matches
    query_pts = np.float32([kp_image[m.queryIdx].pt for m in good_points]).reshape(-1, 1, 2)
    train_pts = np.float32([kp_grayframe[m.trainIdx].pt for m in good_points]).reshape(-1, 1, 2)
    matrix, mask = cv2.findHomography(query_pts, train_pts, cv2.RANSAC, 5.0)
    matches_mask = mask.ravel().tolist()
    h, w = img.shape
    pts = np.float32([[0, 0], [0, h], [w, h], [w, 0]]).reshape(-1, 1, 2)
    dst = cv2.perspectiveTransform(pts, matrix)
    homography = cv2.polylines(frame, [np.int32(dst)], True, (255, 0, 0), 3)
    cv2_imshow(homography)
else:
    print("Not enough good matches found.")
matching_output = cv2.drawMatches(img, kp_image, frame, kp_grayframe, good_points, None, flags=cv2.DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)
cv2_imshow(matching_output)
cap.release()
cv2.destroyAllWindows()
```

## Output



## Practical – 6

**Aim:** Perform Colorization.

### Theory:

We'll create a program to convert a black & white image i.e grayscale image to a colour image. We're going to use the Caffe colourization model for this program. And you should be familiar with basic OpenCV functions and uses like reading an image or how to load a pre-trained model using dnn module etc.

The procedure that we'll follow to implement the program:

Steps:

1. Load the model and the convolution/kernel points
2. Read and preprocess the image
3. Generate model predictions using the L channel from our input image
4. Use the output -> ab channel to create a resulting image

What is the L channel and ab channel? Basically like RGB colour space, there is something similar, known as Lab colour space. And this is the basis on which our program is based. Let's discuss what it is briefly:

What is Lab Colour Space?

Like RGB, lab colour has 3 channels L, a, and b. But here instead of pixel values, these have different significances i.e :

- L-channel: light intensity
- a channel: green-red encoding
- b channel: blue-red encoding

And In our program, we'll use the L channel of our image as input to our model to predict ab channel values and then rejoin it with the L channel to generate our final image.

Automatic colorization of photos using deep neural networks is a technology that can add color to black and white photos without the need for manual

coloring. This technology uses deep neural networks that have been trained on large datasets of color images to learn the relationship between luminance and color, which can then be used to predict the color channels of grayscale images. This technique has many applications, including restoring old photos and enhancing the visual appeal of images.

### Background

Before the advent of computer technology, adding color to a black-and-white photograph was a manual and time-consuming process that required skilled artists. With the introduction of automated colorization methods, the process became quicker but often inaccurate and still required manual intervention. Deep neural networks are a type of machine learning algorithm inspired by the human brain, and they have made it possible to automatically colorize black-and-white photographs with high accuracy and speed. These networks are trained on large datasets of color images and use what they learned to generate plausible colorizations for grayscale images. Using deep neural networks for automatic colorization has many practical applications, such as restoring old photographs, enhancing medical images, and creating realistic 3D models from 2D images.

To implement it, follow the below steps.

1. We need to import some libraries that will be used in our implementation.

```
import numpy as np
```

```
import cv2
```

2. These are variables used in the code for the automatic colorization of photos using a deep neural network. —PROTOTXT is a file that contains the network architecture, —POINTS is a file that stores color space information, and —MODEL is a pre-trained model file that contains the learned weights of the neural network.

```
PROTOTXT = "colorization_deploy_v2.prototxt"
```

```
POINTS = "pts_in_hull.npy"
```

```
MODEL = "colorization_release_v2.caffemodel"
```

3. These lines of code load the pre-trained colorization model and the corresponding points used for color mapping into the program. The model and

points are stored in files, which are read and loaded into the program using the OpenCV library and the numpy library, respectively.

```
net = cv2.dnn.readNetFromCaffe(PROTOTXT, MODEL)  
pts = np.load(POINTS)
```

4. This code sets up the pre-trained neural network by obtaining the layer IDs for class8 and conv8. It then reshapes the —pts|| variable to match the dimensions of the layers and sets the layer blobs to the reshaped —pts|| array and a  $313 \times 1$  array filled with the scalar value of 2.606.

```
class8 = net.getLayerId("class8_ab")  
conv8 = net.getLayerId("conv8_313_rh")  
pts = pts.transpose().reshape(2, 313, 1, 1)  
net.getLayer(class8).blobs = [pts.astype("float32")]  
net.getLayer(conv8).blobs = [np.full([1, 313], 2.606, dtype="float32")]
```

5. The code reads an image file named —flower.jpg|| using OpenCV and scales the pixel values to a range between 0 and 1. The image is then converted from BGR to LAB color space using cv2.cvtColor(). Finally, the image is resized to  $224 \times 224$  pixels using cv2.resize().

```
image = cv2.imread("flower.jpg")  
scaled = image.astype("float32") / 255.0  
lab = cv2.cvtColor(scaled, cv2.COLOR_BGR2LAB)  
resized = cv2.resize(lab, (224, 224))
```

6. In this code, the L channel of the resized LAB image is extracted using cv2.split(), and then its pixel values are subtracted by 50.

```
L = cv2.split(resized)[0]  
L-=50
```

7. The grayscale image —L|| is fed into the pre-trained neural network using the

setInput function to produce an output array —ab|. The output array is reshaped and resized to match the dimensions of the original image.

```
net.setInput(cv2.dnn.blobFromImage(L))
ab = net.forward()[0, :, :, :].transpose((1, 2, 0))
```

```
ab = cv2.resize(ab, (image.shape[1], image.shape[0]))
```

8. This code separates the L channel from a given LAB color image using OpenCV's —cv2.split| function.

```
L = cv2.split(lab)[0]
```

9. In this code, the colorized image is obtained by concatenating the L channel of the LAB color space with the predicted AB channels, converting the image back to the BGR color space, clipping the pixel values to lie between 0 and 1, and finally scaling the pixel values to the range of 0 to 255.

```
colorized = np.concatenate((L[:, :, np.newaxis], ab), axis=2)
colorized = cv2.cvtColor(colorized, cv2.COLOR_LAB2BGR)
colorized = np.clip(colorized, 0, 1)
colorized = (255 * colorized).astype("uint8")
```

10. This code creates a window with a name, resizes it to match the size of the original image, and displays the colorized image in the window.

```
window_name = "Colorized Image By PythonGeeks"
cv2.namedWindow(window_name, cv2.WINDOW_NORMAL)
cv2.resizeWindow(window_name, image.shape[1], image.shape[0])
cv2.imshow(window_name, colorized)
```

11. The code creates a window to display the original image with a specified name and size using OpenCV functions namedWindow and resizeWindow. Then, it shows the original image in the window using the imshow function.

```
cv2.namedWindow("Original Image By PythonGeeks", cv2.WINDOW_NORMAL)
cv2.resizeWindow("Original Image By PythonGeeks", image.shape[1], image.shape[0])
cv2.imshow("Original Image By PythonGeeks", image)
```

12. This program waits indefinitely until a key is pressed. Argument 0 specifies that the program should wait indefinitely.

```
cv2.waitKey(0)
```

Conclusion

Automated colorization using deep neural networks is a technique that uses pre-trained neural networks to automatically add color to grayscale images. This

method is useful for image and video editing, restoration of historical media, and improving accessibility for individuals with visual impairments. The process involves converting the grayscale image to the LAB color space, predicting the color channels using the neural network, and merging them with the grayscale image to produce the final colorized image. Advancements in deep learning algorithms and large-scale datasets will likely lead to further advancements in this field.

**Code -**

```
import numpy as np
import cv2
from cv2 import dnn

proto_file = 'colorization_deploy_v2.prototxt'
model_file = 'colorization_release_v2.caffemodel'
hull_pts = 'pts_in_hull.npy'
img_path = 'DBZ1.jpg'

net = dnn.readNetFromCaffe(proto_file, model_file)
kernel = np.load(hull_pts)

img = cv2.imread(img_path)
scaled = img.astype("float32") / 255.0
lab_img = cv2.cvtColor(scaled, cv2.COLOR_BGR2LAB)

class8 = net.getLayerId("class8_ab")
conv8 = net.getLayerId("conv8_313_rh")
pts = kernel.transpose().reshape(2, 313, 1, 1)
net.getLayer(class8).blobs = [pts.astype("float32")]
net.getLayer(conv8).blobs =[np.full([1, 313], 2.606, dtype = "float32")]

resized = cv2.resize(lab_img, (224, 224))

L = cv2.split(resized)[0]

L -= 50

net.setInput(cv2.dnn.blobFromImage(L))
ab_channel = net.forward()[0, :, :, :].transpose((1, 2, 0))
ab_channel = cv2.resize(ab_channel, (img.shape[1], img.shape[0]))
```

```
L = cv2.split(lab_img)[0]

colorized = np.concatenate((L[:, :, np.newaxis], ab_channel), axis=2)

colorized = cv2.cvtColor(colorized, cv2.COLOR_LAB2BGR)
colorized = np.clip(colorized, 0, 1)

colorized = (255 * colorized).astype("uint8")

img = cv2.resize(img, (250,500))
colorized = cv2.resize(colorized, (250,500))

result = cv2.hconcat([img,colorized])

cv2.imshow("Grayscale -> Colour", result)

cv2.waitKey(0)
```

### Output –



## Practical – 7

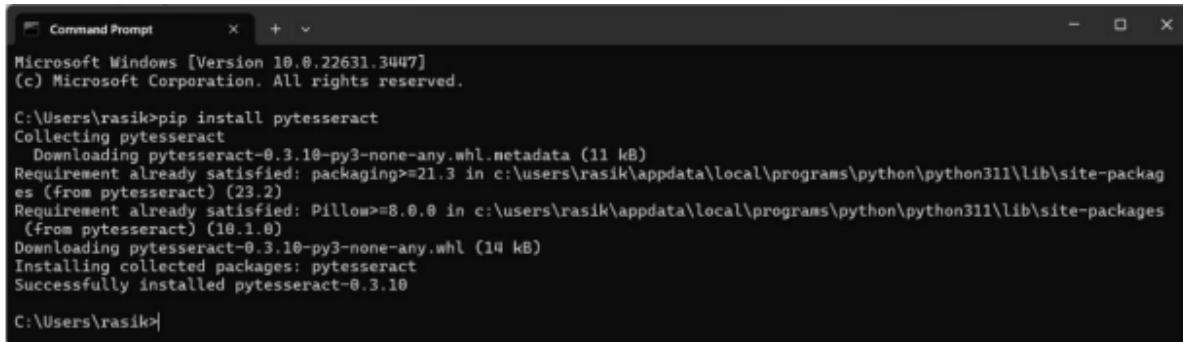
**Aim:** Perform Text Detection and Recognition.

### Theory:

Required Installations:

pip install opencv-python

pip install pytesseract



```
Microsoft Windows [Version 10.0.22631.3447]
(c) Microsoft Corporation. All rights reserved.

C:\Users\rasik>pip install pytesseract
Collecting pytesseract
  Downloading pytesseract-0.3.10-py3-none-any.whl.metadata (11 kB)
Requirement already satisfied: packaging>=21.3 in c:\users\rasik\appdata\local\programs\python\python311\lib\site-packages (from pytesseract) (23.2)
Requirement already satisfied: Pillow>=8.0.0 in c:\users\rasik\appdata\local\programs\python\python311\lib\site-packages (from pytesseract) (10.1.0)
  Downloading pytesseract-0.3.10-py3-none-any.whl (14 kB)
Installing collected packages: pytesseract
Successfully installed pytesseract-0.3.10

C:\Users\rasik>
```

Follow the instructions given in below Link:

<https://builtin.com/articles/python-tesseract>

OpenCV package is used to read an image and perform certain image

processing techniques. Python-tesseract is a wrapper for Google's Tesseract-OCR Engine which is used to recognize text from images.

Python-tesseract is an optical character recognition (OCR) tool for python. That is, it will recognize and —read— the text embedded in images.

Python-tesseract is a wrapper for Google's Tesseract-OCR Engine. It is also useful as a stand-alone invocation script to tesseract, as it can read all image types supported by the Pillow and Leptonica imaging libraries, including jpeg, png, gif, bmp, tiff, and others. Additionally, if used as a script, Python-tesseract will print the recognized text instead of writing it to a file.

Approach:

After the necessary imports, a sample image is read using the imread function of opencv.

Applying image processing for the image:

The colorspace of the image is first changed and stored in a variable. For color conversion we use the function cv2.cvtColor(input\_image, flag). The second parameter flag determines the type of conversion. We can chose among cv2.COLOR\_BGR2GRAY and cv2.COLOR\_BGR2HSV.

cv2.COLOR\_BGR2GRAY helps us to convert an RGB image to gray scale image and cv2.COLOR\_BGR2HSV is used to convert an RGB image to HSV

(Hue, Saturation, Value) color-space image. Here, we use cv2.COLOR\_BGR2GRAY. A threshold is applied to the converted image using cv2.threshold function.

There are 3 types of thresholding:

1. Simple Thresholding
2. Adaptive Thresholding
3. Otsu's Binarization

For more information on thresholding, refer Thresholding techniques using OpenCV.

cv2.threshold() has 4 parameters, first parameter being the color-space changed image, followed by the minimum threshold value, the maximum threshold value and the type of thresholding that needs to be applied.

To get a rectangular structure:

cv2.getStructuringElement() is used to define a structural element like elliptical, circular, rectangular etc. Here, we use the rectangular structural element (cv2.MORPH\_RECT). cv2.getStructuringElement takes an extra size of the kernel parameter. A bigger kernel would make group larger blocks of texts together. After choosing the correct kernel, dilation is applied to the image with cv2.dilate function. Dilation makes the groups of text to be detected more accurately since it dilates (expands) a text block.

Finding Contours:

`cv2.findContours()` is used to find contours in the dilated image. There are three arguments in `cv.findContours()`: the source image, the contour retrieval mode and the contour approximation method.

This function returns contours and hierarchy. Contours is a python list of all the contours in the image. Each contour is a Numpy array of (x, y) coordinates of boundary points in the object. Contours are typically used to find a white object from a black background. All the above image processing techniques are applied so that the Contours can detect the boundary edges of the blocks of text of the image. A text file is opened in write mode and flushed. This text file is opened to save the text from the output of the OCR.

Applying OCR:

Loop through each contour and take the x and y coordinates and the width and height using the function `cv2.boundingRect()`. Then draw a rectangle in the image using the function `cv2.rectangle()` with the help of obtained x and y coordinates and the width and height. There are 5 parameters in the `cv2.rectangle()`, the first parameter specifies the input image, followed by the x and y coordinates (starting coordinates of the rectangle), the ending coordinates of the rectangle which is  $(x+w, y+h)$ , the boundary color for the rectangle in RGB value and the size of the boundary. Now crop the rectangular region and then pass it to the tesseract toextract the text from the image. Then we open the created text file in append mode to append the obtained text and close the file.

**Code:**

```
import cv2
import pytesseract

pytesseract.pytesseract.tesseract_cmd = 'C:\\Program Files\\Tesseract-OCR\\tesseract.exe'

img = cv2.imread('sample.jpg')

gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

ret, threshl = cv2.threshold(gray, 0, 255, cv2.THRESH_OTSU | cv2.THRESH_BINARY_INV)

rect_kernel = cv2.getStructuringElement(cv2.MORPH_RECT, (18,18))

dilation= cv2.dilate(threshl, rect_kernel, iterations = 1)

contours, hierarchy = cv2.findContours(dilation, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_NONE)

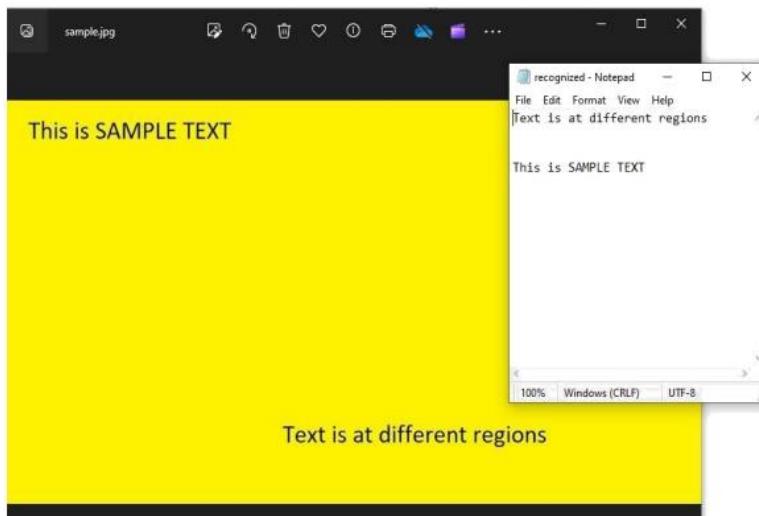
im2 = img.copy()

file = open('recognized.txt','w+')
file.write("")
file.close()

for cnt in contours:
    x,y,w,h = cv2.boundingRect(cnt)
    rect = cv2.rectangle(im2, (x,y), (x+w,y+h), (0,255,0),2)
    cropped = im2[y:y+h,x:x+w]

    file = open('recognized.txt','a')
    text = pytesseract.image_to_string(cropped)
    file.write(text)
    file.write("\n")

file.close()
```

**Output**

## Practical – 8

**Aim:** Construct 3D model from Images.

### Theory:

Transforming a 2D image into a 3D environment requires depth estimation, which can be a difficult operation depending on the amount of precision and information required. OpenCV supports a variety of depth estimation approaches, including stereo vision and depth from focus/defocus.

**In this practical we'll see a basic technique utilizing stereovision: Transform a 2D image into a 3D space using OpenCV**

Transforming a 2D image into a 3D space using OpenCV refers to the process of converting a two-dimensional image into a three-dimensional spatial representation using the Open Source Library (OpenCV). This transformation involves inferring the depth information from the 2D image, typically through techniques such as stereo vision, depth estimation, or other algorithms, to create a 3D model with depth perception. This process enables various applications such as 3D reconstruction, depth sensing, and augmented reality.

### Importance of transformations of a 2D image into a 3D space

Transforming 2D images into 3D space becomes crucial in various fields due to its numerous applications and benefits:

**Depth Perception:** We are able to detect depth by transforming 2D pictures into 3D space. This makes it possible to use augmented reality, object recognition, and scene understanding.

**3D Reconstruction:** Converting 2D photos into 3D space makes it easier to recreate 3D scenes, which is crucial in industries like robotics, , and the preservation of cultural assets.

**Stereo Vision:** Stereo vision depends on converting 2D images into 3D space. It entails taking pictures from various angles and calculating depth from the difference between matching spots. It is employed in 3D modeling, autonomous navigation, and depth sensing, among other applications.

**Medical Imaging:** Improved visualization, diagnosis, and treatment planning are possible in medical imaging when 2D medical scans—such as CT or MRI scans—are converted into 3D space.

**Virtual Reality and Simulation:** In virtual reality, simulation, and gaming, realistic 3D worlds must be constructed from 2D photos or video. This requires translating 2D visuals into 3D space.

### How you get a 3D image from a 2D?

In conventional photography, you can either utilize a mirror and attach a camera to it to create an immediate 3D effect, or you can take a shot, step to your right (or left), and then shoot another, ensuring that all components from the first photo are present in the second.

However, if you just move a 2D picture left by 10 pixels, nothing changes. This is because you are shifting the entire environment, and no 3D information is saved.

Instead, there must be a bigger shift distance between the foreground and backdrop. In other words, the farthest point distant from the lens remains motionless while the nearest point moves.

How is this related to using a 2D image to create a 3D image?

We need a method to move the pixels since a picture becomes three-dimensional when the foreground moves more than the background.

Fortunately, a technique known as depth detection exists that generates what is known as a depth map.

Now remember that this is only an estimate. Furthermore, it won't reach every nook and corner. All that depth detection does is use cues like shadows, haze, and depth of focus to determine if an object is in the forefront or background.

### We can instruct a program on how far to move pixels now that we have a depth map.

*Approach:*

- **Obtain Stereo Images:** Take or obtain two pictures of the same scene taken from two separate perspectives.
- **Stereo Calibration:** For precise depth estimation, ascertain each camera's intrinsic and extrinsic properties.
- **Rectification:** To make matching easier, make sure corresponding spots in stereo pictures line up on the same scanlines.
- **Stereo matching:** Use methods such as block matching or SGBM to find correspondences between corrected stereo pictures.
- **Disparity:** Calculate the disparity by dividing the horizontal locations of corresponding points by their pixels.

- **Depth Estimation:** Using camera settings and stereo geometry, convert disparity map to depth map.
- **3D Reconstruction:** Using the depth map as a guide, reconstruct the scene's three dimensions.

Implementations of a 2D image into a 3D space Transformations using OpenCV Input image:

- [cube 1](#)
- [cube 2](#)

*Code steps:*

- First we have imported the required libraries.  
 [PIL \(Python Imaging Library\)](#), which is used to work with images  
 [numpy](#) is used for numerical computations
- Then we have defined a function named `shift\_image` that takes three parameters: `img`, `depth\_img`, and `shift\_amount`. `img` is the base image that will be shifted, `depth\_img` is the depth map providing depth information for the shift, and `shift\_amount` specifies the maximum amount of horizontal shift allowed.
- After that we have converted the input images (`img` and `depth\_img`) into arrays using NumPy for further processing. The base image (**img**) is converted to RGBA format to ensure it has an alpha channel, while the depth image (**depth\_img**) is converted to grayscale (**L**) to ensure it contains only one channel representing depth values.
- Then calculates the shift amounts based on the depth values obtained from the depth image. It scales the depth values to the range [0, 1] and then multiplies it by the `shift\_amount`. The result is then converted to integers using `astype(int)`.
- Then initializes an array (`shifted\_data`) with the same shape as the base image (`data`) filled with zeros. This array will store the shifted image data.
- The nested loops iterate over the elements of the `deltas` array to perform the shift operation. For each pixel position (x, y) in the base image, the corresponding shift amount `dx` is applied horizontally. If the resulting position (x + dx, y) is within the bounds of the image, the pixel value at (x, y) in the base image is copied to (x + dx, y) in the shifted image.
- `Image.fromarray` converts the shifted image data

(`shifted\_data`) back to a PIL Image object. The data is converted to np.uint8 type before creating the image.

- Then we read the images from our local files and applied to the newly created functions.
- `shifted\_img.show` will plot the newly transformed image

### **Limitations of OpenCV to 2D image into a 3D space image transformations**

The process of simulating a 3D effect by shifting pixels based on depth information has several limitations also:

**Simplified Depth Representation:** Since the depth map is a grayscale picture, intricate depth fluctuations might not be properly captured.

It Only moves pixels horizontally, assuming that all depth fluctuations are horizontal. This assumption might not hold true in scenarios found in the actual world.

**Uniform Shift Amount:** Unrealistic effects may result from the linear connection between depth and pixel shift, which may not adequately reflect real-world depth perception.

**Limited Depth Cues:** It does not use other depth cues such as perspective distortion or occlusion, instead depending solely on the parallax effect.

**Boundary Artifacts:** Boundary artifacts may arise from the disregard of pixels that have been pushed outside the limits of the picture.

**Dependency on Parameter:** The shift amount parameter that is selected has a significant impact on the 3D effect's quality

**Limited Application Scope:** Although it works well for basic 3D effects, it might not be accurate enough for complex 3D activities like medical imaging or augmented reality.

#### **Code –**

```
from PIL import Image
import numpy as np
import matplotlib.pyplot as plt

def shift_image(img, depth_img, shift_amount=10):
    img = img.convert("RGBA")
    data = np.array(img)

    depth_img = depth_img.convert("L")
    depth_data = np.array(depth_img)
    deltas = ((depth_data / 255.0) * float(shift_amount)).astype(int)

    shifted_data = np.zeros_like(data)
    height, width, _ = data.shape

    for y, row in enumerate(deltas):
        for x, dx in enumerate(row):
            if 0 <= x + dx < width:
                shifted_data[y, x + dx] = data[y, x]

    shifted_image = Image.fromarray(shifted_data.astype(np.uint8))
```

```
return shifted_image

# Load images
img_path = "/content/sample_data/cube1.jpg" # Replace with your image path
depth_img_path = "/content/sample_data/cube2.jpg" # Replace with your depth image path

img = Image.open(img_path)
depth_img = Image.open(depth_img_path)

# Apply the shift

shifted_img = shift_image(img, depth_img, shift_amount=10)

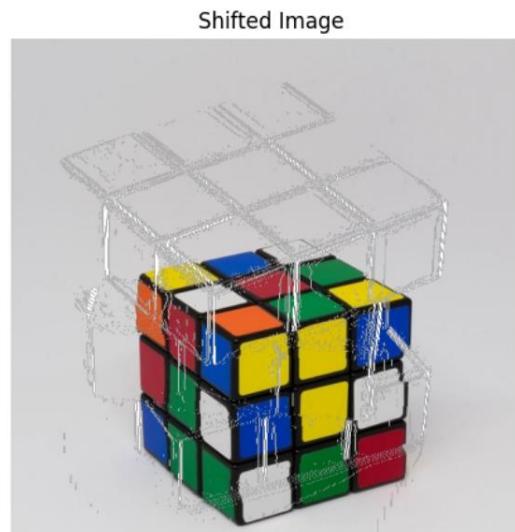
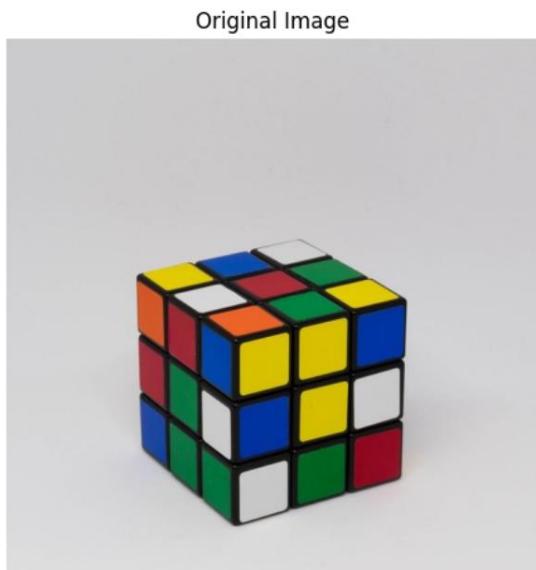
# Display the original and shifted images using matplotlib
fig, axs = plt.subplots(1, 2, figsize=(12, 6))

# Original Image
axs[0].imshow(img)
axs[0].set_title('Original Image')
axs[0].axis('off')

# Shifted Image
axs[1].imshow(shifted_img)
axs[1].set_title('Shifted Image')
axs[1].axis('off')

plt.show()
```

### Output -



## Practical – 9

**Aim:** Perform Feature extraction using RANSAC.

### Theory:

**Image registration** is a digital image processing technique that helps us align different images of the same scene. For instance, one may click the picture of a book from various angles. Below are a few instances that show the diversity of camera angles.

Now, we may want to —align| a particular image to the same angle as a reference image. In the images above, one may consider the first image to be an —ideal| cover photo, while the second and third images do not serve well for book cover photo purposes. The image registration algorithm helps us align the second and third pictures to the same plane as the first one.

### How does image registration work?

Alignment can be looked at as a simple coordinate transform. The algorithm works as follows:

- Convert both images to grayscale.
- Match features from the image to be aligned, to the reference image and store the coordinates of the corresponding key points. Keypoints are simply the selected few points that are used to compute the transform (generally points that stand out), and descriptors are histograms of the image gradients to characterize the appearance of a keypoint. In this post, we use ORB (Oriented FAST and Rotated BRIEF) implementation in the OpenCV library, which provides us with both key points as well as their associated descriptors.
- Match the key points between the two images. In this post, we use BFMatcher, which is a brute force matcher. BFMatcher.match() retrieves the best match, while BFMatcher.knnMatch() retrieves top K matches, where K is specified by the user.
- Pick the top matches, and remove the noisy matches.
- Find the homomorphy transform.
- Apply this transform to the original unaligned image to get the output image.

## Applications of Image Registration –

Some of the useful applications of image registration include:

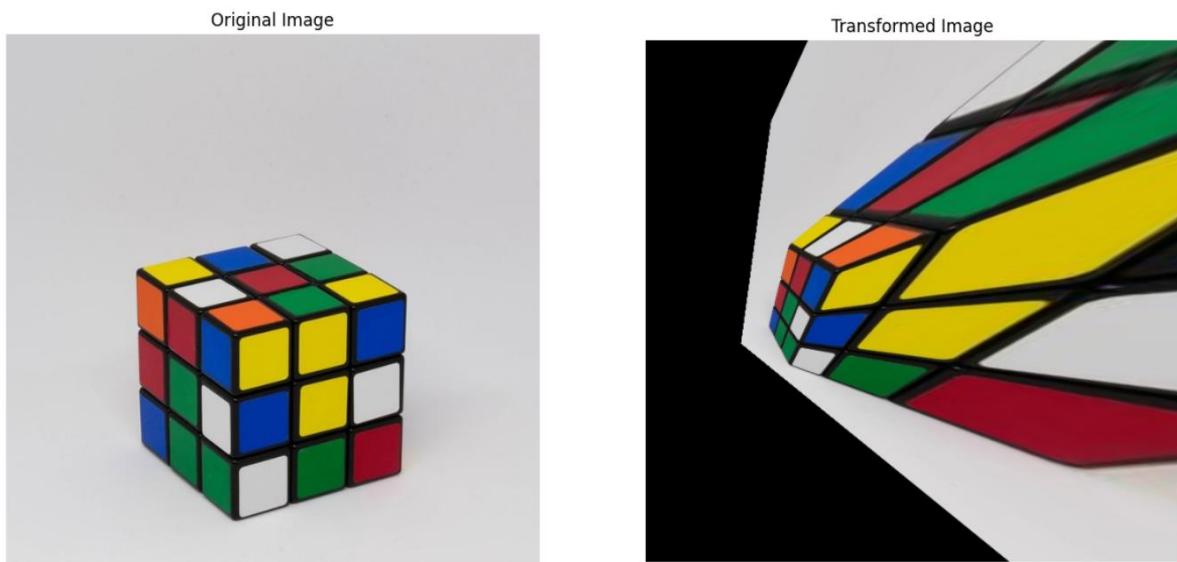
- Stitching various scenes (which may or may not have the same camera alignment) together to form a continuous panoramic shot.
- Aligning camera images of documents to a standard alignment to create realistic scanned documents.
- Aligning medical images for better observation and analysis.

### Code -

```
import cv2
import numpy as np
import matplotlib.pyplot as plt
img1_color = cv2.imread("/content/sample_data/cube1.jpg")
img2_color = cv2.imread("/content/sample_data/cube2.jpg")
img1 = cv2.cvtColor(img1_color, cv2.COLOR_BGR2GRAY)
img2 = cv2.cvtColor(img2_color, cv2.COLOR_BGR2GRAY)
height, width = img2.shape
orb_detector = cv2.ORB_create(5000)
kp1, d1 = orb_detector.detectAndCompute(img1, None)
kp2, d2 = orb_detector.detectAndCompute(img2, None)
matcher = cv2.BFMatcher(cv2.NORM_HAMMING, crossCheck=True)
matches = matcher.match(d1, d2)
matches = sorted(matches, key=lambda x: x.distance)
matches = matches[:int(len(matches) * 0.9)]
no_of_matches = len(matches)
p1 = np.zeros((no_of_matches, 2))
p2 = np.zeros((no_of_matches, 2))
```

```
for i in range(len(matches)):  
    p1[i, :] = kp1[matches[i].queryIdx].pt  
    p2[i, :] = kp2[matches[i].trainIdx].pt  
  
homography, mask = cv2.findHomography(p1, p2, cv2.RANSAC)  
transformed_img = cv2.warpPerspective(img1_color, homography, (width,  
height))  
  
cv2.imwrite('output.jpg', transformed_img)  
  
fig, axs = plt.subplots(1, 2, figsize=(15, 10))  
axs[0].imshow(cv2.cvtColor(img1_color, cv2.COLOR_BGR2RGB))  
axs[0].set_title('Original Image')  
axs[0].axis('off')  
  
axs[1].imshow(cv2.cvtColor(transformed_img, cv2.COLOR_BGR2RGB))  
axs[1].set_title('Transformed Image')  
axs[1].axis('off')  
  
plt.show()
```

### Output –



MSC IT Part 1 Sem 2

Roll No.: 00